

High Performance Natural Language Processing

EMNLP 2020



Presenters



Gabriel Ilharco

University of Washington



Cesar Ilharco

Google



Iulia Turc

Google



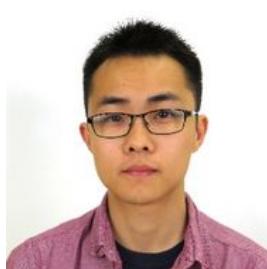
Tim Dettmers

University of Washington



Felipe Ferreira

Google



Kenton Lee

Google

High Performance NLP

Slides available at: bit.ly/2SmhKY7



Google Research

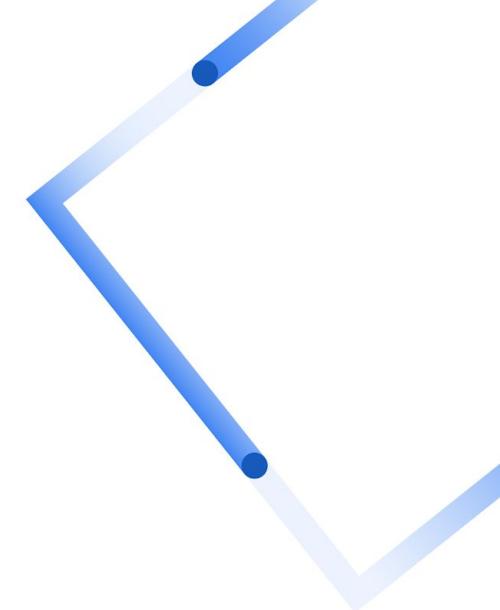


Agenda

- 01 Introduction
- 02 Fundamentals
- 03 Core Techniques
- 04 Efficient Attention
- 05 Case Studies
- 06 Scaling in Practice
- 07 Closing Notes

01

Introduction

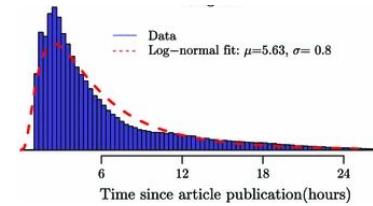


Motivation & Applications

Why do we need it ?

SCALE

- NEWS
 - Realtime: Majority of content is consumed within a few hours after publication [1]
 - Thousands of news articles per second
 - 40-80 sentences per article
- SOCIAL NETWORKS: ~6 Thousand tweets per second [2]
- THE WEB: Orders of magnitude bigger



What could we do, if we had it ?

[1] Tatar, A., Antoniadis, P., Amorim, M.D.d. et al. From popularity prediction to ranking online news. Soc. Netw. Anal. Min. 4, 174 (2014). <https://doi.org/10.1007/s13278-014-0174-8>

[2] https://blog.twitter.com/engineering/en_us/a/2013/new-tweets-per-second-record-and-how.html

Summarization

5:18 • 51%

Full Coverage

Eddie Van Halen has died

Top coverage



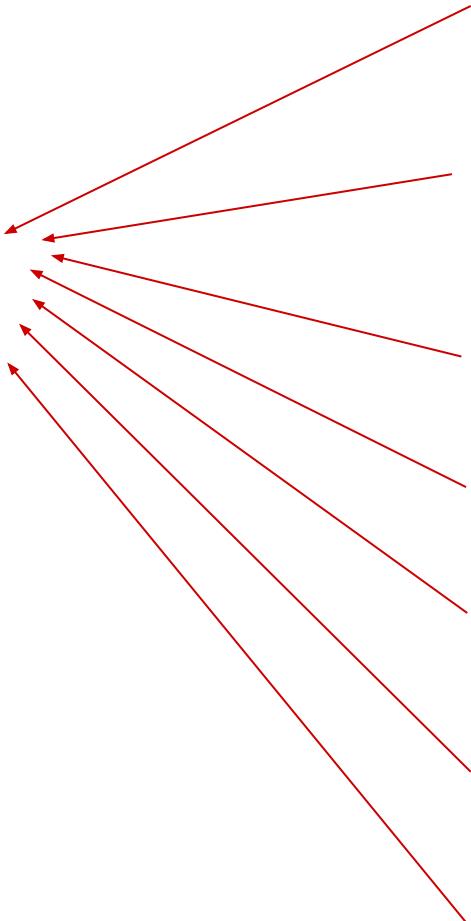
The New York Times

Eddie Van Halen, Virtuoso of the Rock Guitar, Dies at 65

11 minutes ago

TMZ

Eddie Van Halen Dead at 65 from Cancer



PEOPLE.com

Eddie Van Halen's Son Wolf Posts Adorable Childhood Photo with Dad Days After His Death

Eddie Van Halen died on Tuesday after battling cancer for years.
2 days ago



Billboard

Van Halen's Songs Were Streamed More Than 30 Million Times After Eddie Van Halen's Death

Following the death of Eddie Van Halen on Oct. 6, his namesake band's catalog of songs were streamed over 31 million times in the U.S., ...
2 days ago



NME.com

Van Halen streams increase by over 1,300% following Eddie Van Halen's death

Van Halen's US music streams have increased by over 1300% since the death of the band's guitarist Eddie Van Halen, figures have revealed.
23 hours ago



TMZ

Eddie Van Halen Dead at 65 from Cancer

Eddie Van Halen has died at age 65.
5 days ago



Yahoo Entertainment

Eddie Van Halen Dies at 65

Eddie Van Halen, whose innovative and explosive guitar playing kept the hard rock band that bore his family name cemented to the top of the ...
5 days ago



The New York Times

Eddie Van Halen, Virtuoso of the Rock Guitar, Dies at 65

Eddie Van Halen, whose razzle-dazzle guitar-playing — combining complex harmonics, innovative fingerings and ingenious devices he ...
4 days ago



Fox News

Eddie Van Halen, legendary rock guitarist, dead at 65

Eddie Van Halen died after a battle with cancer. His son Wolfgang confirmed on ...



Summarization

5:18 51%

Full Coverage

Eddie Van Halen has died

Top coverage



Associated Press

The New York Times

Eddie Van Halen, Virtuoso of the Rock Guitar, Dies at 65

11 minutes ago

TMZ

Eddie Van Halen Dead at 65 from Cancer

Eddie Van Halen

Musician



Available on

- YouTube
- Spotify
- Apple Music
- More music services

Edward Lodewijk "Eddie" Van Halen was an American musician, songwriter, producer, and inventor. He was the main songwriter and lead guitarist of the American rock band Van Halen, which he co-founded in 1972 with his brother, drummer Alex Van Halen, bassist Mark Stone, and singer David Lee Roth. [Wikipedia](#)

Died: October 6, 2020, Santa Monica, CA Trending

Born: January 26, 1955, Amsterdam, Netherlands

Spouse: Janie Liszewski (m. 2009–2020), Valerie Bertinelli (m. 1981–2007)

Children: Wolfgang Van Halen

Facts Extraction

Quinton Byfield X | ⚡ | 🔎

All News Videos Images Shopping More Settings Tools

About 361,000 results (0.54 seconds)

Top stories



yahoosports
Quinton Byfield becomes highest-ever drafted Black player in history after Kings take him No...

15 hours ago



SPORTINGNEWS
NHL Draft 2020 Round 1 winners, losers: Alex Trebek, Quinton Byfield make their mark, Maple...

3 hours ago



The Athletic
The winners, losers and late maneuvers of the NHL Draft's first round

11 hours ago



Quinton Byfield
Ice hockey centre

Quinton Byfield is a Canadian junior ice hockey centre. He is currently playing for the Sudbury Wolves of the Ontario Hockey League. Byfield was selected second overall by the Los Angeles Kings in the 2020 NHL Entry Draft, becoming the highest drafted black player in NHL history. [Wikipedia](#)

Born: August 19, 2002 (age 18 years), Newmarket, Canada
Height: 6' 4"
Nationality: Canadian
Playing career: TBD–present
Current team: Sudbury Wolves (Centerman)
NHL Draft: Eligible 2020

www.eliteprospects.com > player > quinton-byfield ▾

Quinton Byfield - Elite Prospects

Eliteprospects.com hockey player profile of Quinton Byfield, 2002-08-19 Newmarket, ON, CAN. Most recently in the OHL with Sudbury Wolves.

Team: Sudbury Wolves Place of birth: Newmarket

Quinton Byfield (F) · Quinton Byfield - Elite Prospects · Quinton Byfield Embed Stats

Highest-ever drafted Black player in history.

Facts Extraction

- Noteworthy Facts
- Trendiness



yahoo/sports

Quinton Byfield becomes highest-ever drafted Black player in history after Kings take him No...

15 hours ago

VS

Rink Royalty

LA Kings: Scout believes Quinton Byfield doesn't need extra time

Winners in the lottery and coming away with the second overall pick, the pick is likely going to be either Quinton Byfield or Tim Stützle. The debate ...
1 week ago



dobberprospects.com

How Quinton Byfield Stacks Up to Previous Second Overall

...
@Hockey_Robinson explains why he stuck to his guns and ranked Quinton Byfield ahead of Alexis Lafreniere in his latest draft rankings.
2 weeks ago



Sportsnet.ca

Quinton Byfield selected second overall by Kings in 2020 NHL ...

The Los Angeles Kings selected Quinton Byfield of the Sudbury Wolves with the second-overall pic in the 2020 NHL Draft. Top Videos ...
6 days ago



Detroit Jock City

Red Wings: Quinton Byfield or Tim Stützle falling to four is dream come true

The Detroit Red Wings have the fourth overall pick, and it could be a dream come true if Quinton Byfield or Tim Stützle falls into Steve ...
4 weeks ago



Sentence Entailment

 Snopes Search Snopes.com

Submit a Topic | Shop Snopes | What's New | Hot 50 | Fact Checks

In September 2020, a **false rumor** saying that Netflix CEO Reed Hastings had been arrested on child pornography charges was widely circulated on social media. While that rumor was entirely made up out of whole cloth, the pushers of that misinformation at least correctly identified the CEO of the streaming giant.

A few days after that rumor went viral, an even more incorrect version started to make its way around the internet. This time, **social media users claimed** that it was actually Netflix CEO "Kim Martin Morrow" — who is nonexistent — had been arrested for child pornography charges:



False

About this rating ↗

www.wbtw.com > news > national > sen-cruz-calls-for-criminal-invest... ▾

Sen. Cruz calls for criminal investigation into Netflix's 'Cuties' — WBTW

Sep 12, 2020 — Unedited press release from the press office of Sen. Ted Cruz (R-Texas) sent a letter calling on the Department of Justice to investigate whether Netflix, its executives or the makers of the film violated any federal laws against the production and distribution of child pornography.



Correct Attribution

About this rating ↗

7:58 ⓘ
◀ Search

Benjamin Schoch Carman Castle ...

Kim Martin Morrow, the CEO of Netflix has just been charged with 15 charges for child pornography and 31,000 files have been found on his personal computers for child porn from ages 8 and as young as toddlers. So, I think the investigation was needed. #SaveOurChildren



i

WBTW.COM

Sen. Cruz calls for criminal investigation into Netflix's 'Cuties'

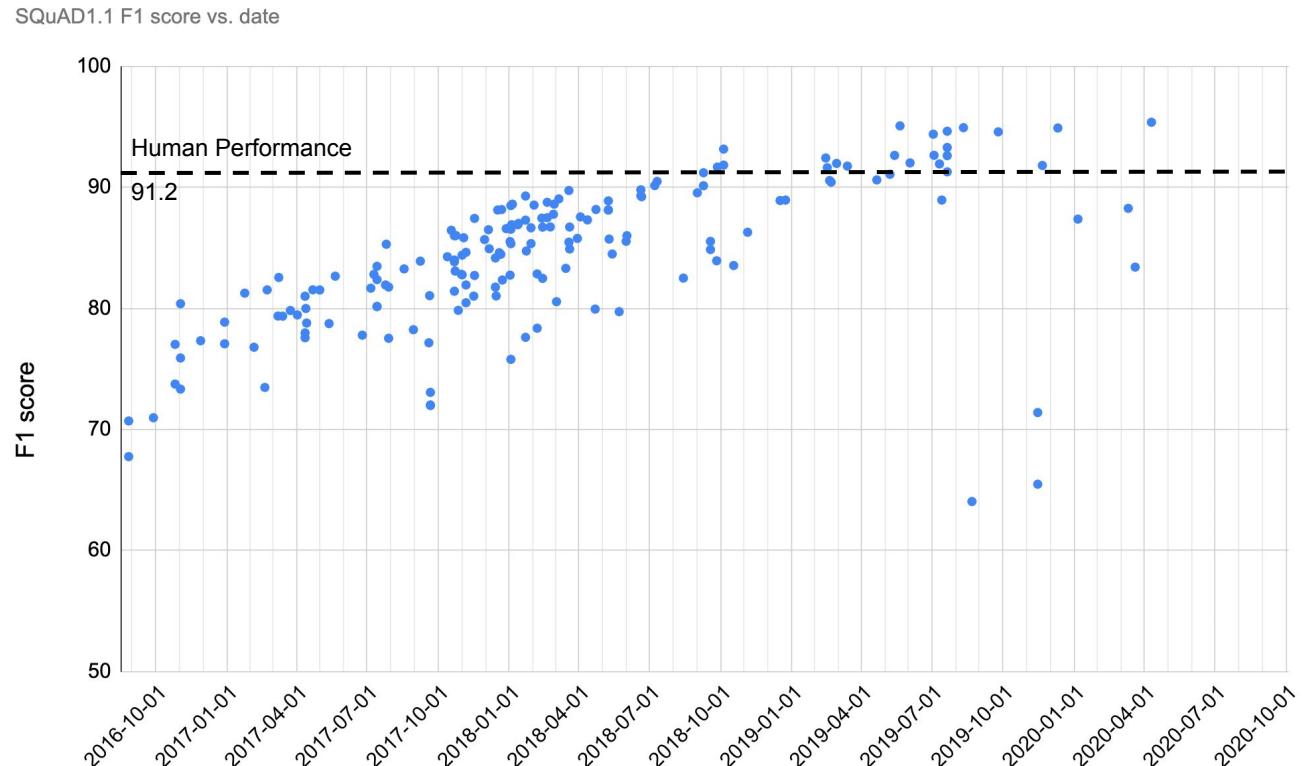
Share

10

11,937 Shares

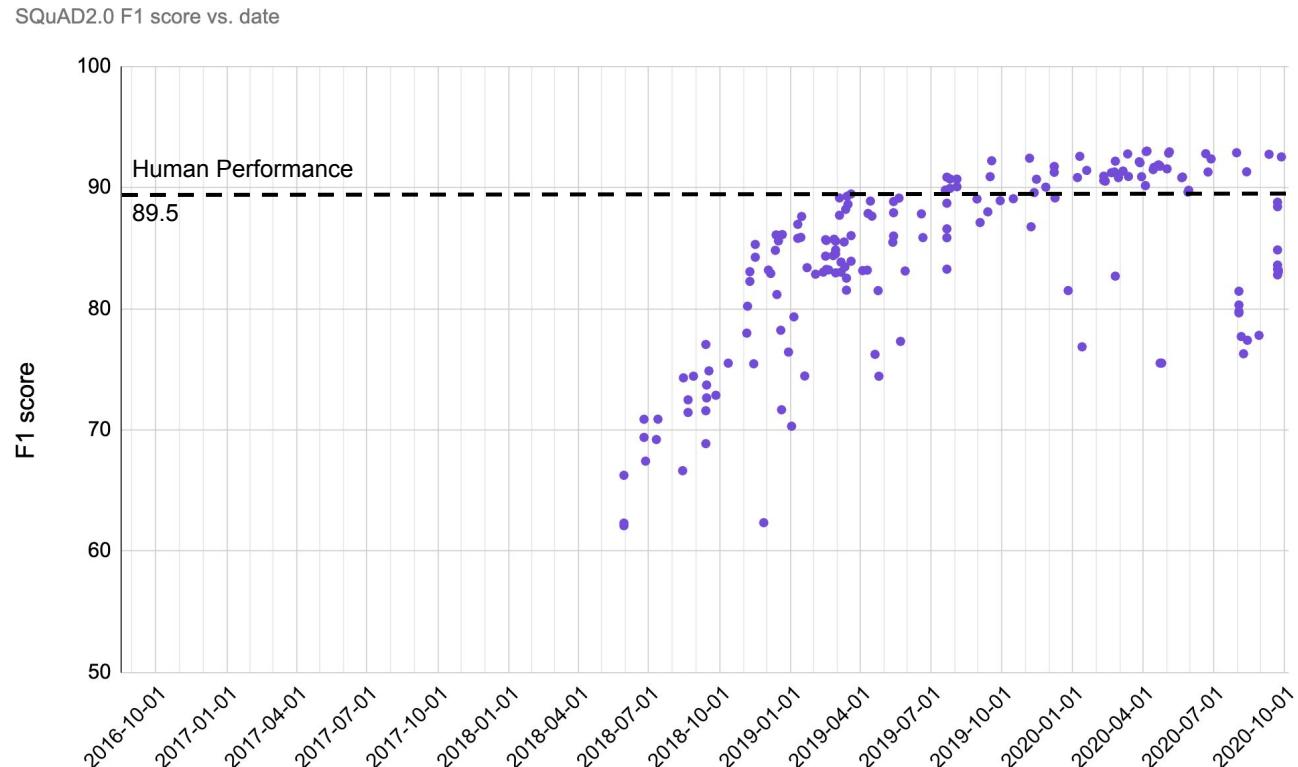
Recent years in Natural Language Processing

Benchmarks through the years - SQuAD 1.1



The Stanford Question Answering Dataset, <https://rajpurkar.github.io/SQuAD-explorer/>

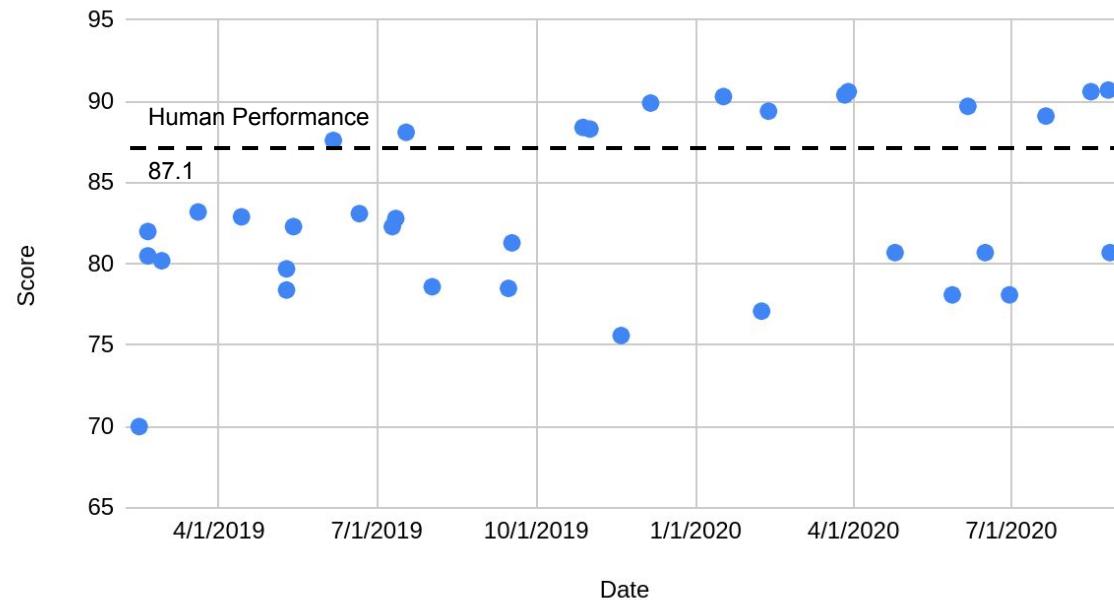
Benchmarks through the years - SQuAD 2.0



The Stanford Question Answering Dataset, <https://rajpurkar.github.io/SQuAD-explorer/>

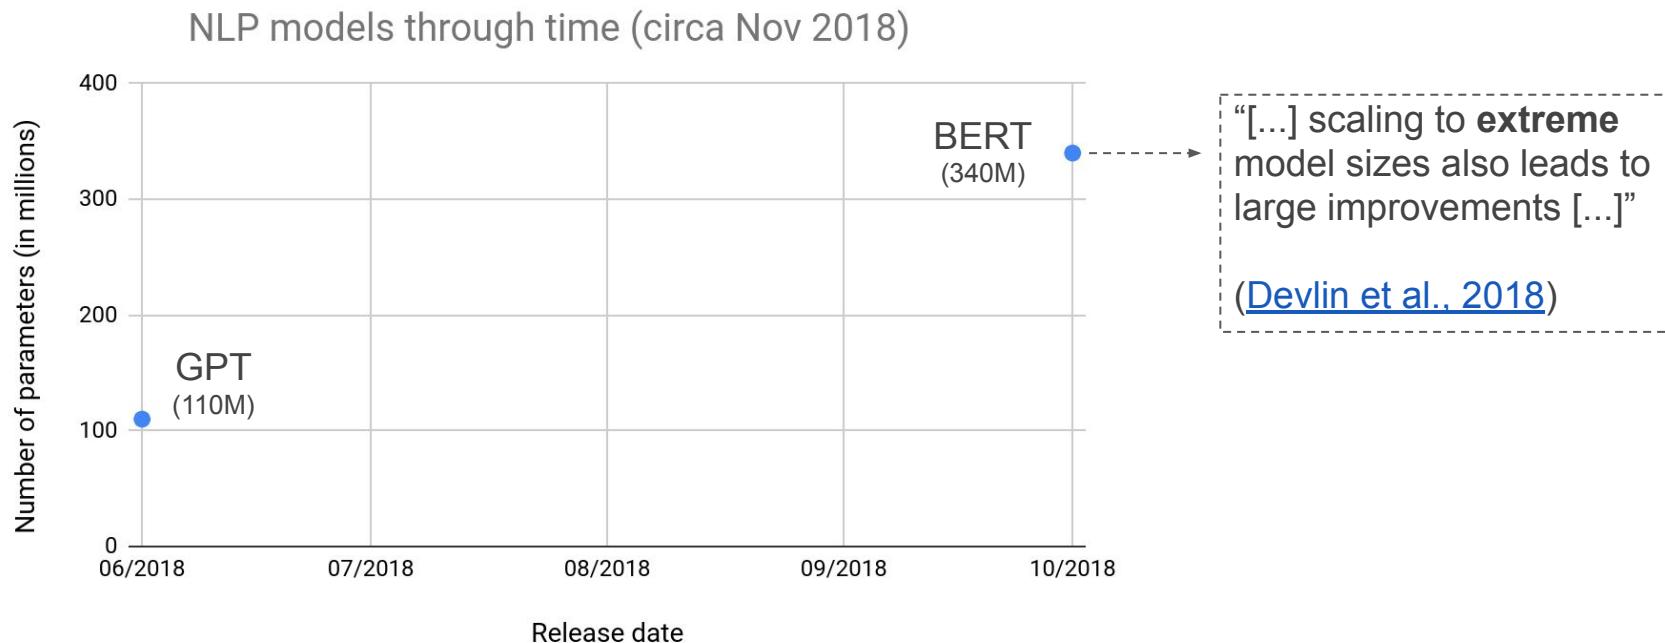
Benchmarks through the years - GLUE

GLUE aggregated score vs. Date

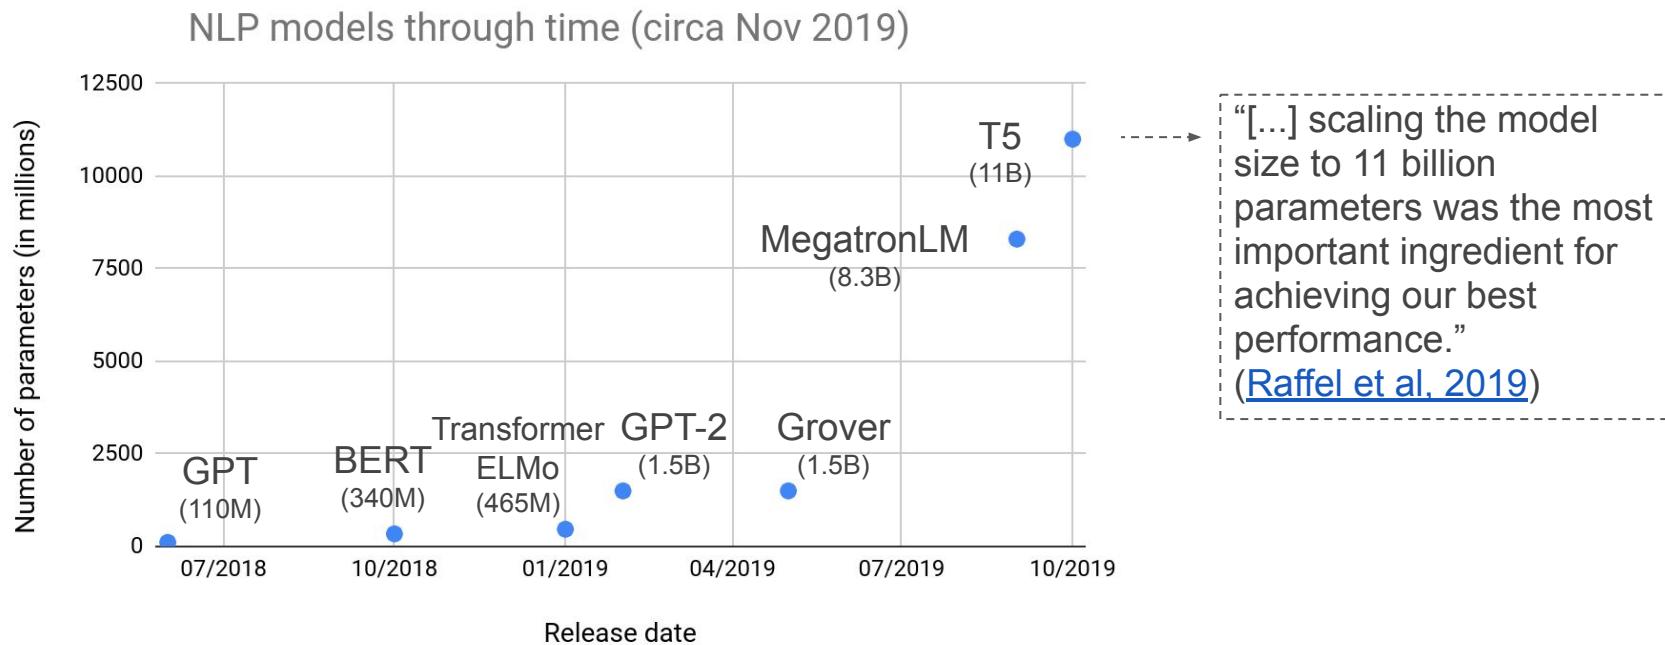


The GLUE Benchmark ([Wang et al., 2018](#))

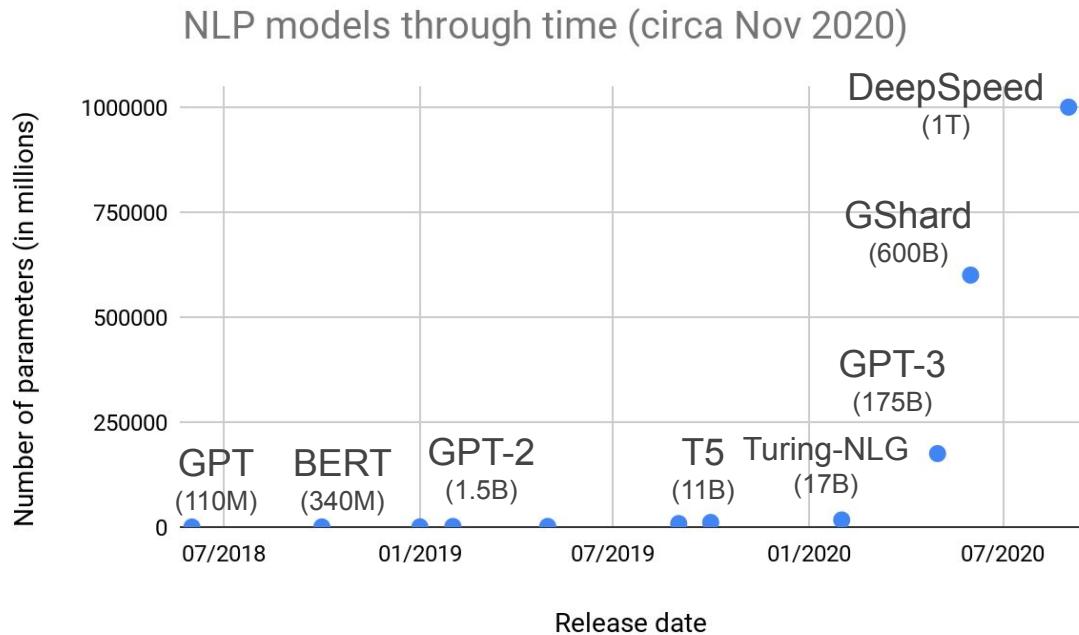
A brief recent history of scale in NLP



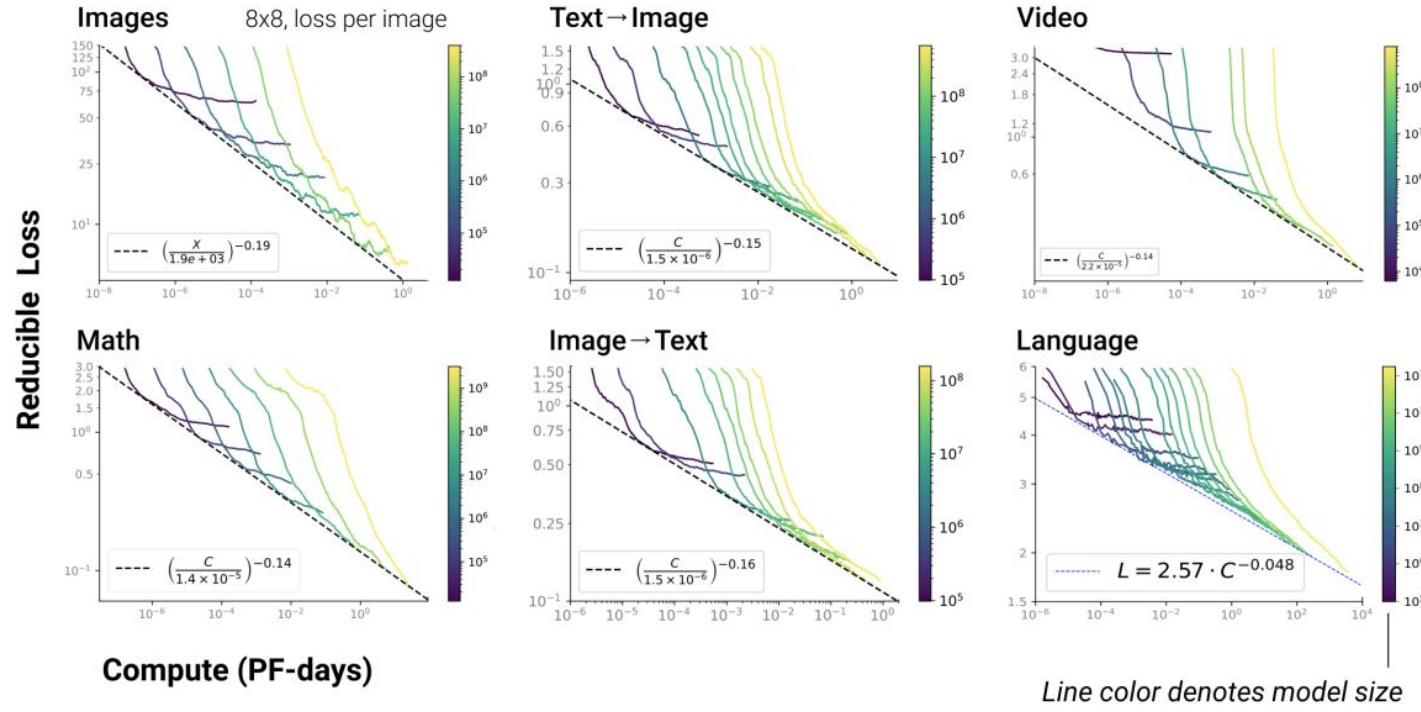
A brief recent history of scale in NLP



A brief recent history of scale in NLP



Scaling Laws

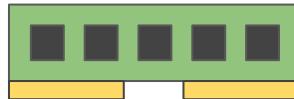


[Henighan et al., 2020](#)

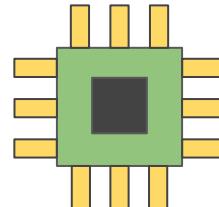
The drawbacks of naive scaling

1) Disconnect with production systems

- Latency
- Hardware constraints
- Energy costs



Memory



CPU/GPU/TPU



Storage



Battery

The drawbacks of naive scaling

1) Disconnect with production systems

2) Costs

- Hardware

- 2048 TPU v3 accelerators (*GShard*, [Lepikhin et al., 2020](#))

- 285,000 CPU cores, 10,000 GPUs (*GPT-3*, [Brown et al., 2020](#))

- Financial

- GPT-3 training cost is [estimated](#) at 4.6 million dollars.

The drawbacks of naive scaling

- 1) Disconnect with production systems
- 2) Costs

3) Accessibility

- Ever-larger hardware and financial requirements impose great barriers to many researchers and institutions
- This can have a serious impact in our research community
 - For instance, 62% of PhD students have access to 4 or less GPUs, according to a recent [poll](#).

The drawbacks of naive scaling

- 1) Disconnect with production systems
- 2) Costs
- 3) Accessibility

Altogether, this is especially relevant to a field that scaled by
3 orders of magnitude in 2 years.

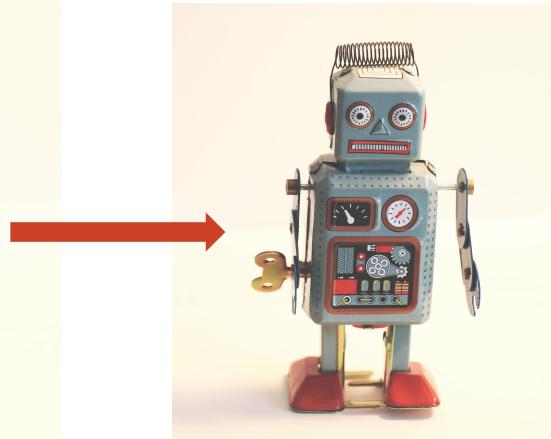
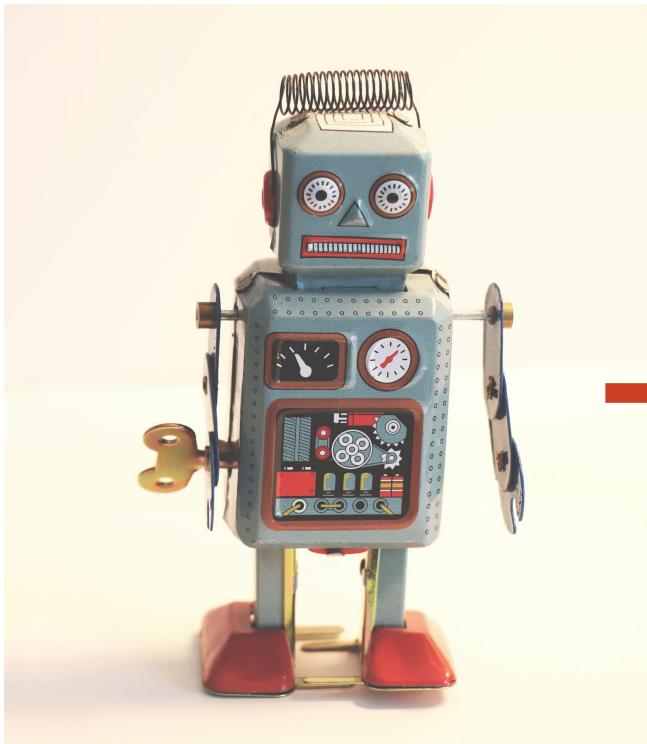
We should **strive for efficiency**



Towards more efficient NLP

1) Core techniques

- Knowledge Distillation

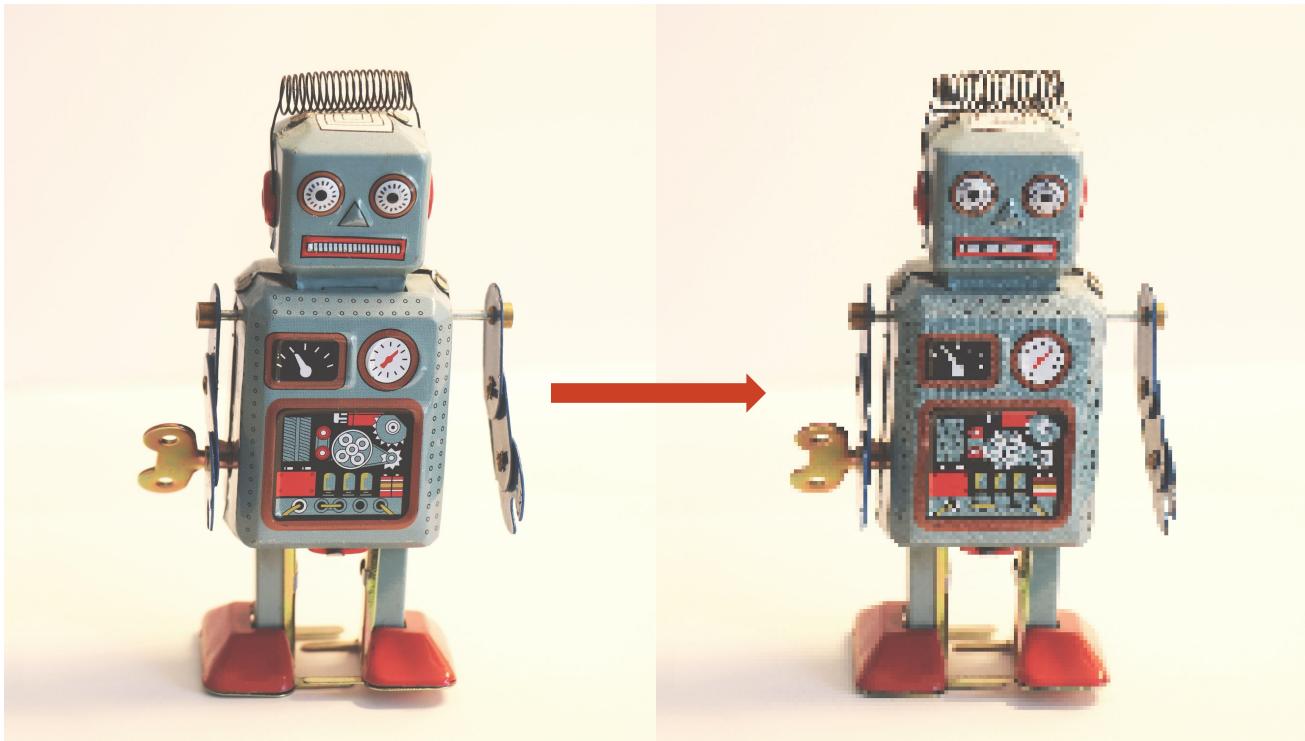


Source: unsplash.com

Towards more efficient NLP

1) Core techniques

- Knowledge Distillation
- Quantization

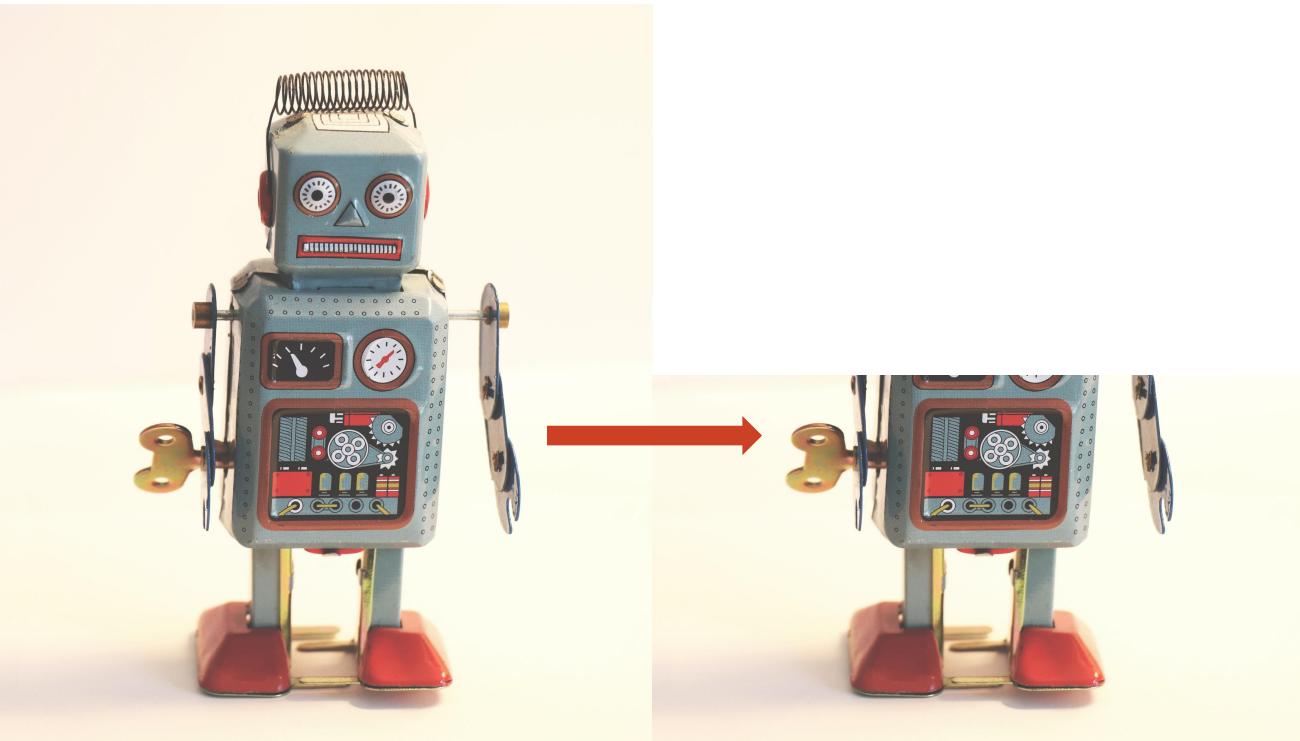


Source: unsplash.com

Towards more efficient NLP

1) Core techniques

- Knowledge Distillation
- Quantization
- Pruning



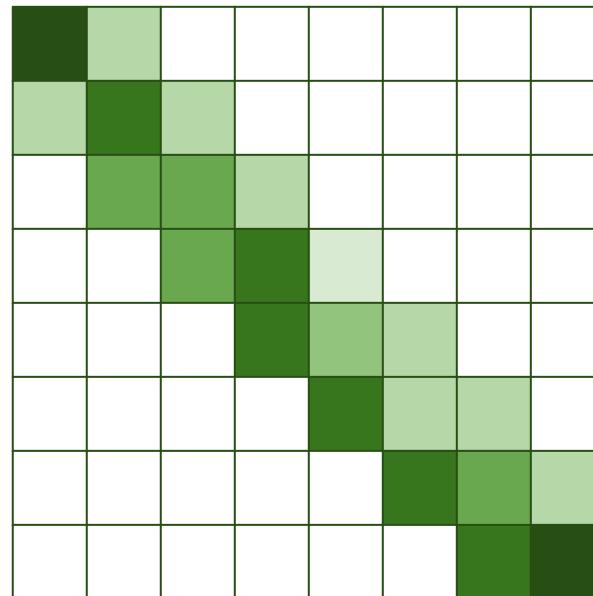
Source: unsplash.com

Towards more efficient NLP

1) Core techniques

2) Efficient attention

- Data-Independent Patterns

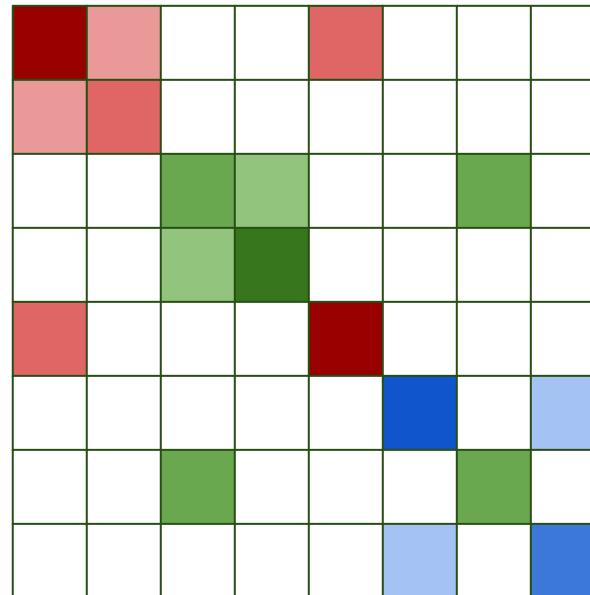


Towards more efficient NLP

1) Core techniques

2) Efficient attention

- Data-Independent Patterns
- Data-Dependent Patterns

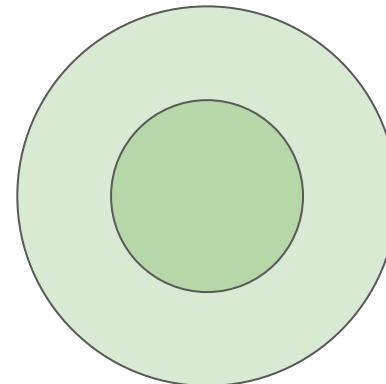


Towards more efficient NLP

1) Core techniques

2) Efficient attention

- Data-Independent Patterns
- Data-Dependent Patterns
- Kernels and Alternative Attention Mechanisms

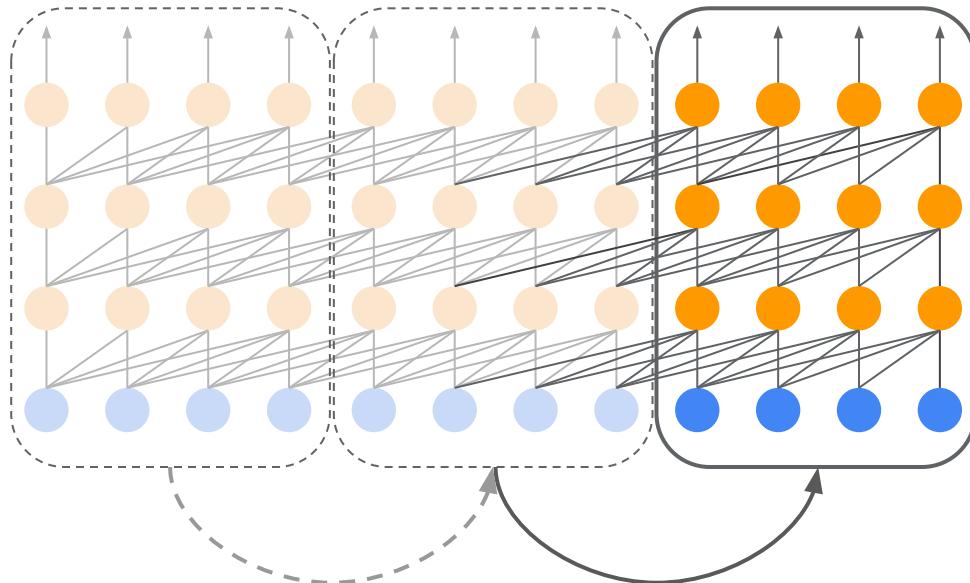


Towards more efficient NLP

1) Core techniques

2) Efficient attention

- Data-Independent Patterns
- Data-Dependent Patterns
- Alternative Attention Mechanisms
- Recurrence

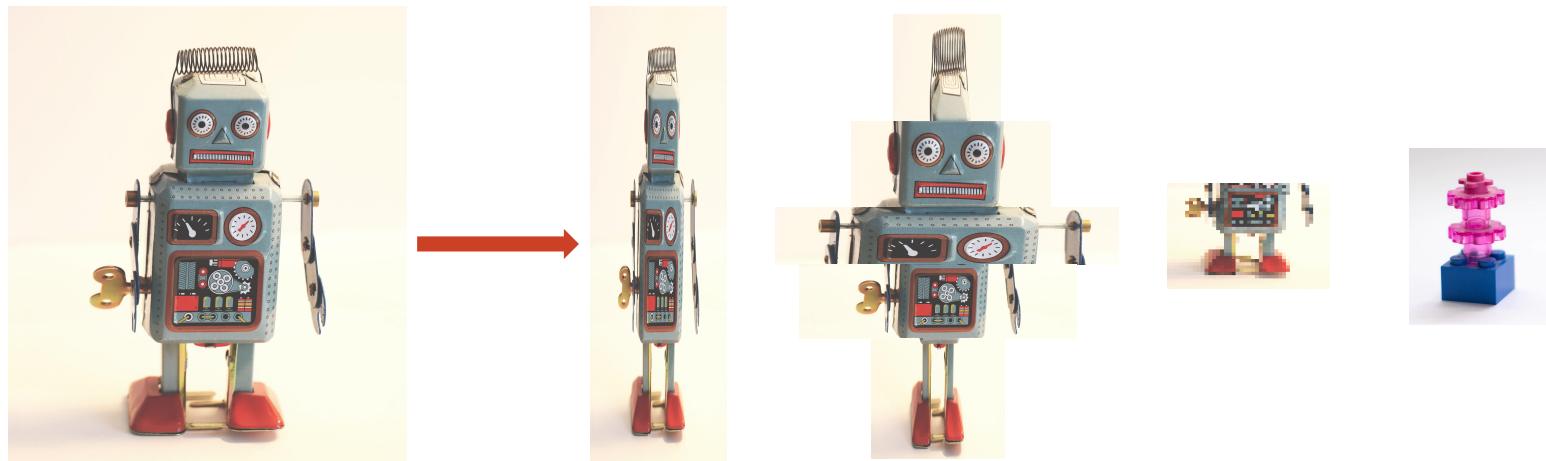


Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention

3) Case studies

- Efficient Language Models



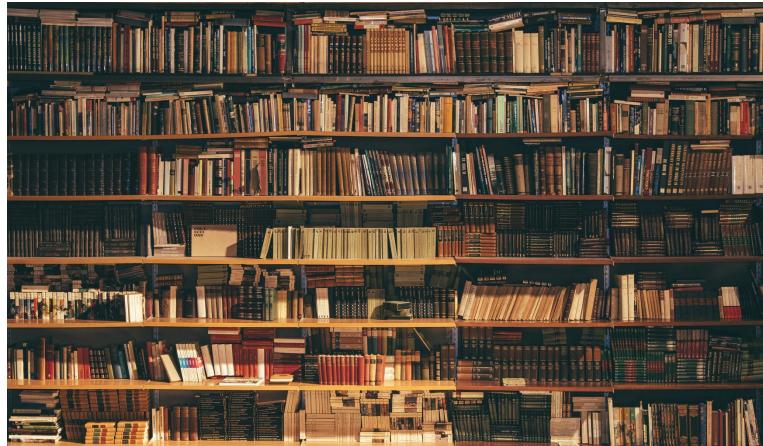
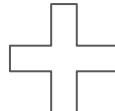
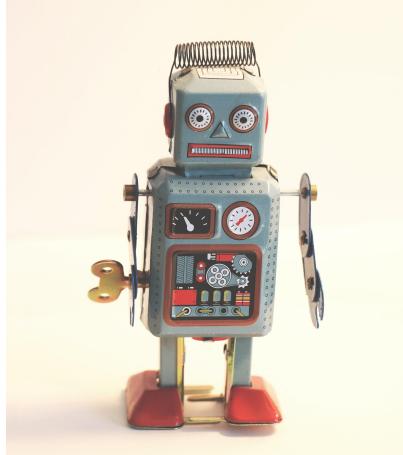
Source: unsplash.com

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention

3) Case studies

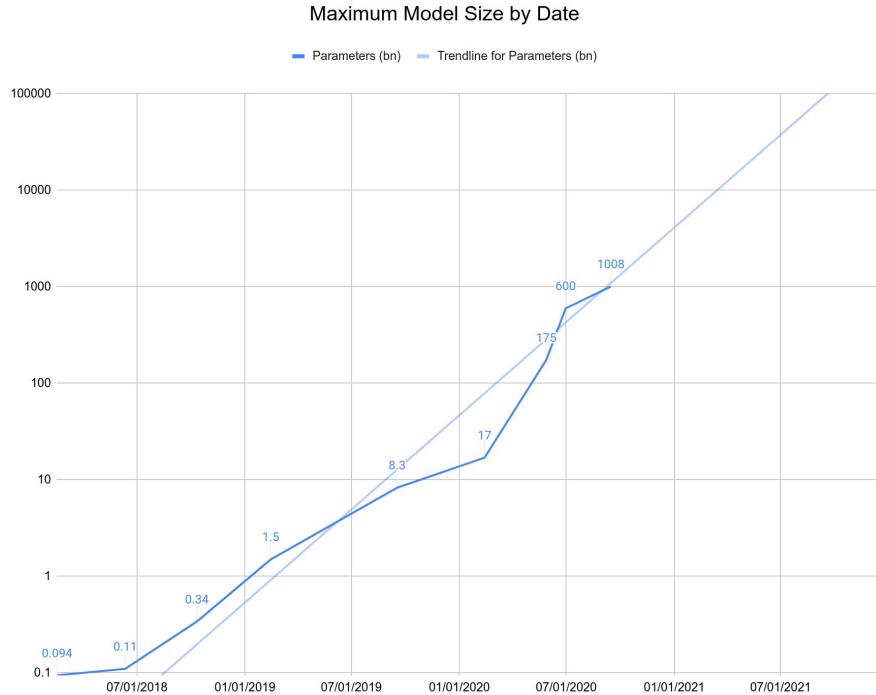
- Efficient Language Models
- Retrieval



Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies
- 4) Scaling in Practice**

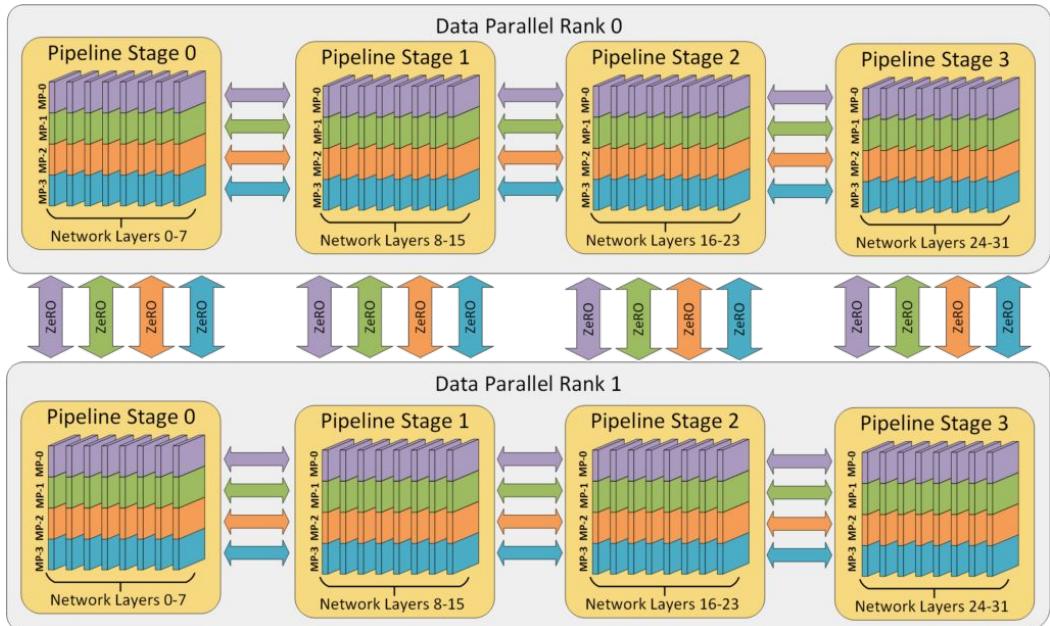
- Scaling Laws of Neural Language Models



Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies
- 4) Scaling in Practice**

- Scaling Laws of Neural Language Models
- Parallelism Techniques



Source: Microsoft [Blog Post](#)

Towards more efficient NLP

1) Core techniques

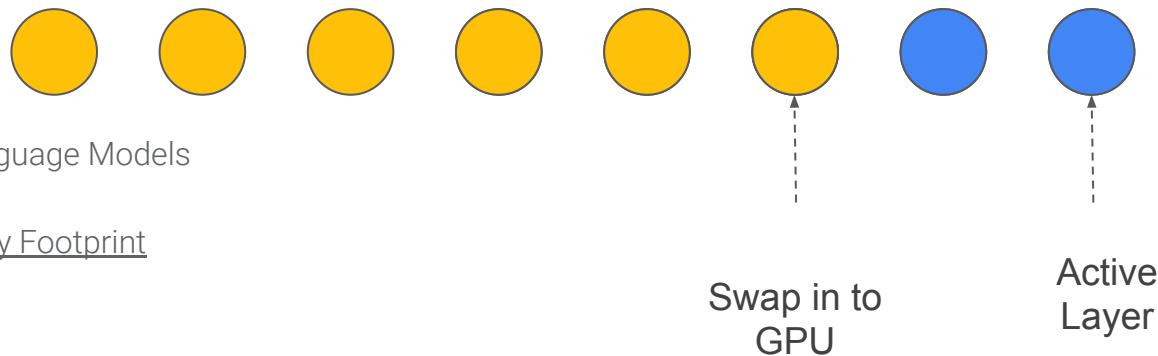
■ CPU ■ GPU

2) Efficient attention

3) Case studies

4) Scaling in Practice

- Scaling Laws of Neural Language Models
- Parallelism Techniques
- Methods to Reduce Memory Footprint



Swap in to
GPU

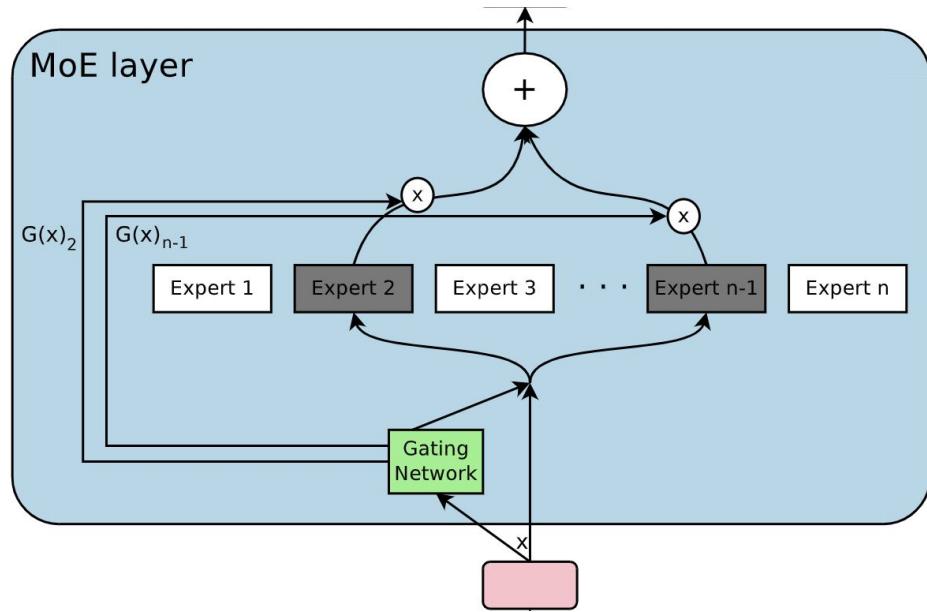
Active
Layer

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies

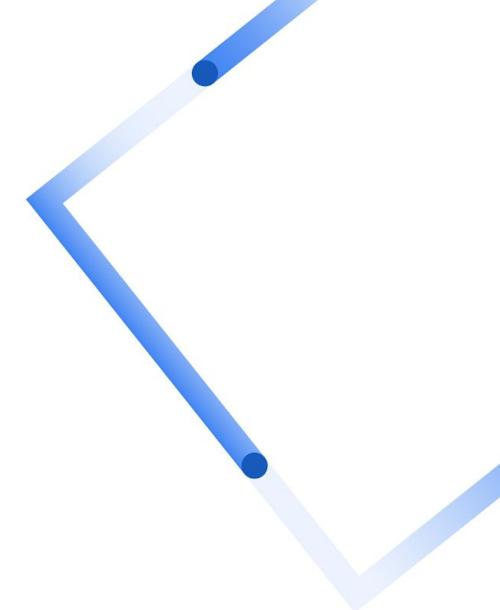
4) Scaling in Practice

- Scaling Laws of Neural Language Models
- Parallelism Techniques
- Methods to Reduce Memory Footprint
- Mixture of Experts

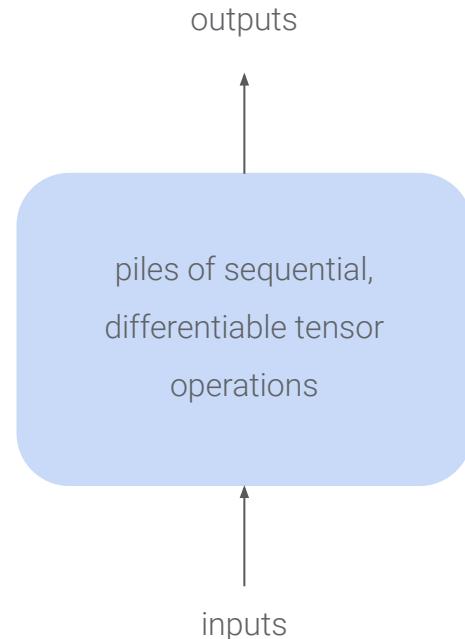


02

Fundamentals

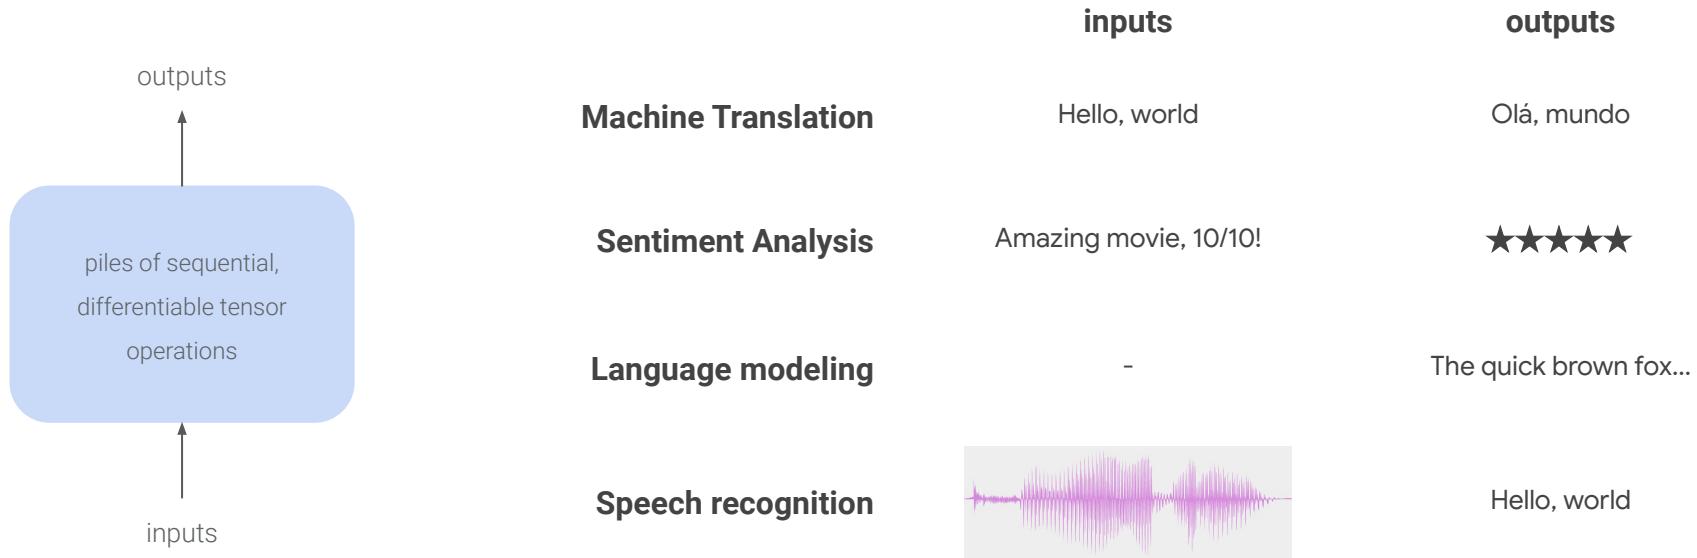


Sequence-to-sequence models



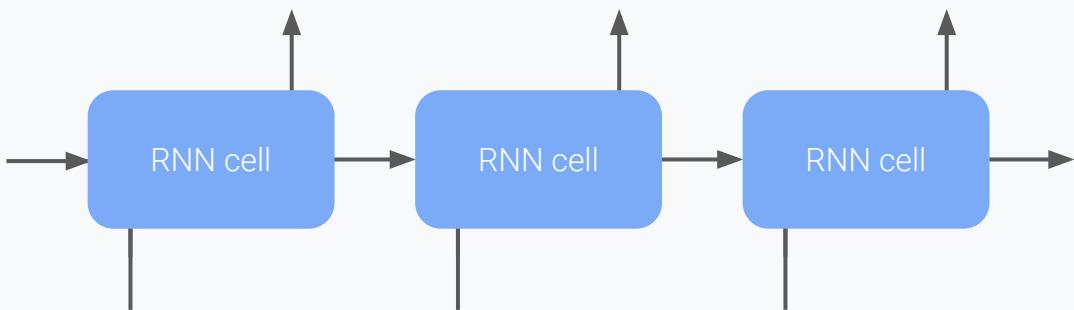
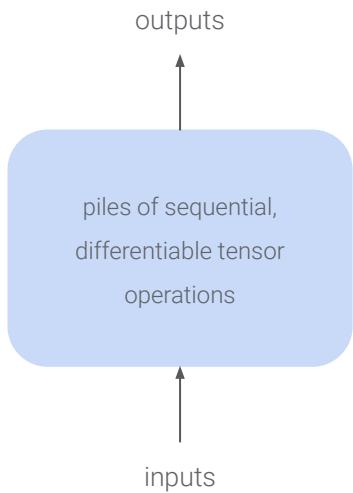
<https://xkcd.com/1838/>

Sequence-to-sequence models



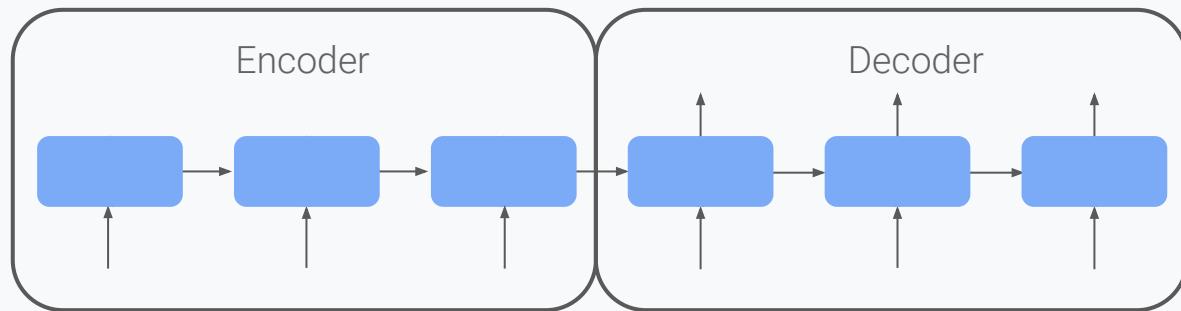
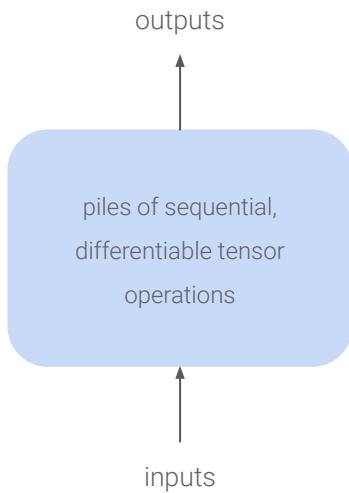
RNNs

RNNs allow computations over sequences of arbitrary length

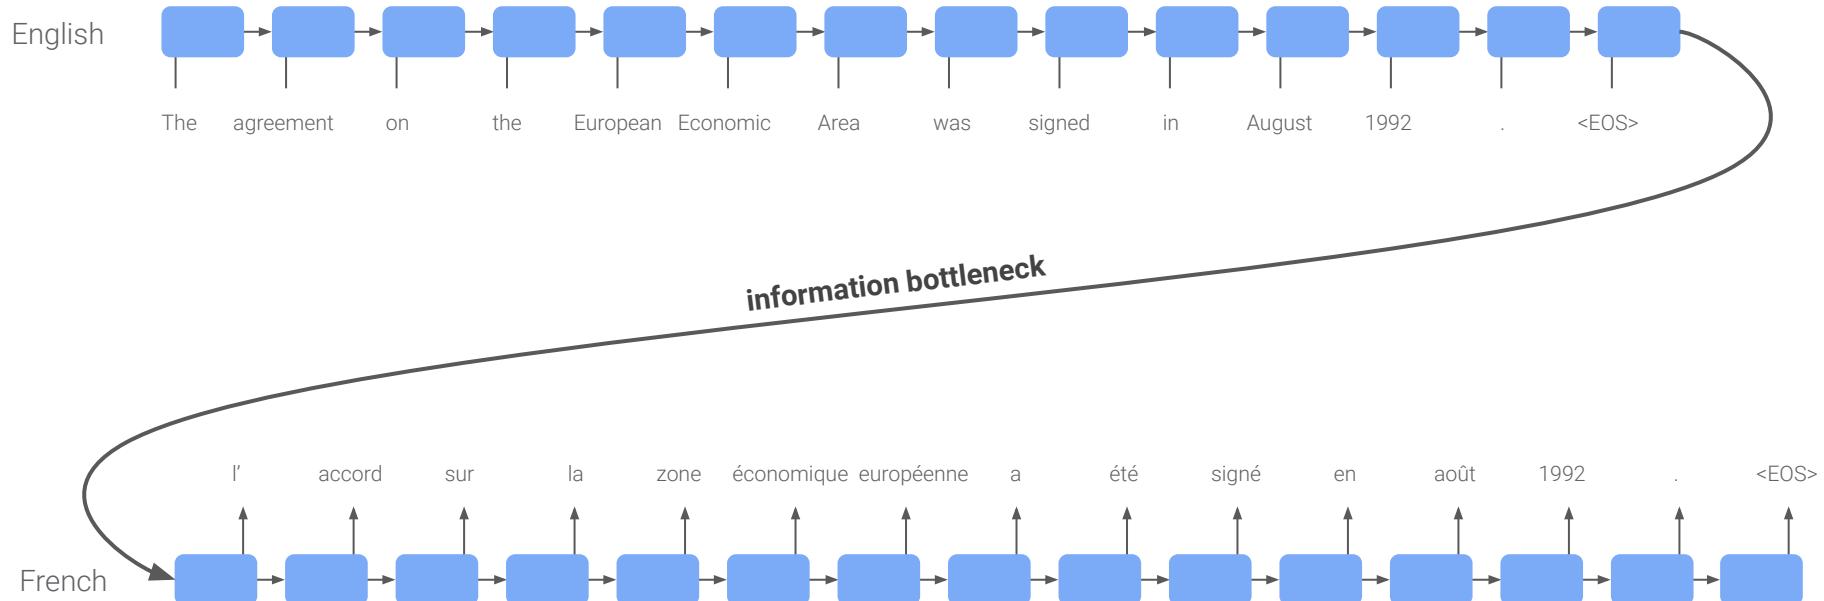


Encoders and Decoders

RNNs allow computations over sequences of arbitrary length

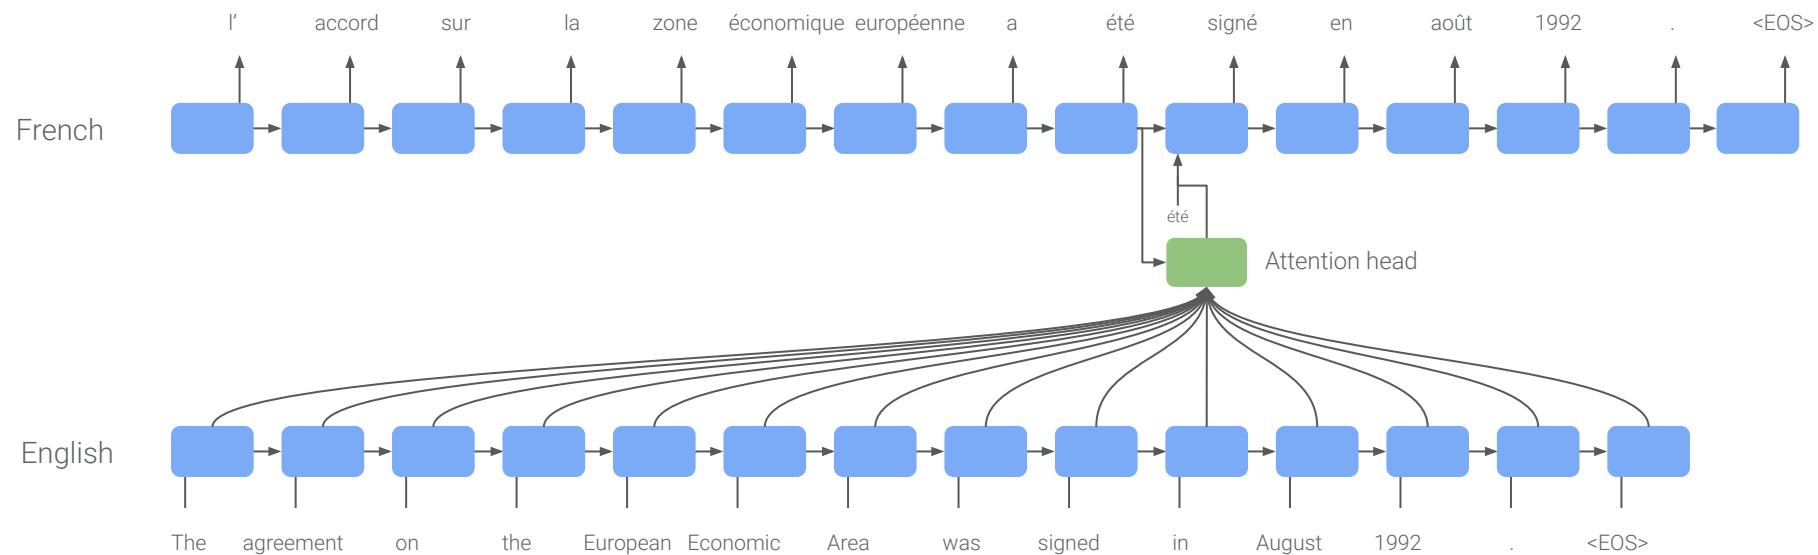


The encoder-decoder bottleneck



Example derived from Bahdanau, et al. 2014 (<https://arxiv.org/pdf/1409.0473.pdf>)

Attention



Example derived from Bahdanau, et al. 2014 (<https://arxiv.org/pdf/1409.0473.pdf>)

Attention

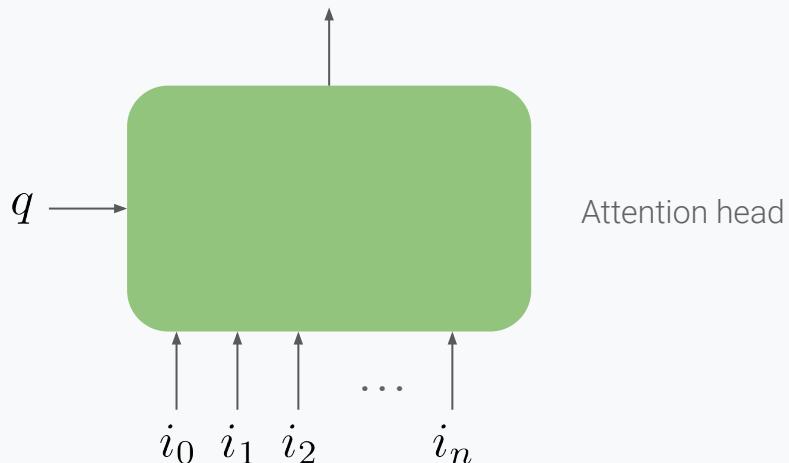
A summary of L
based on how similar their
are with the query

Bahdanau et al.

Neural machine translation by
jointly learning to align and
translate. 2014

Thang Luong et al.

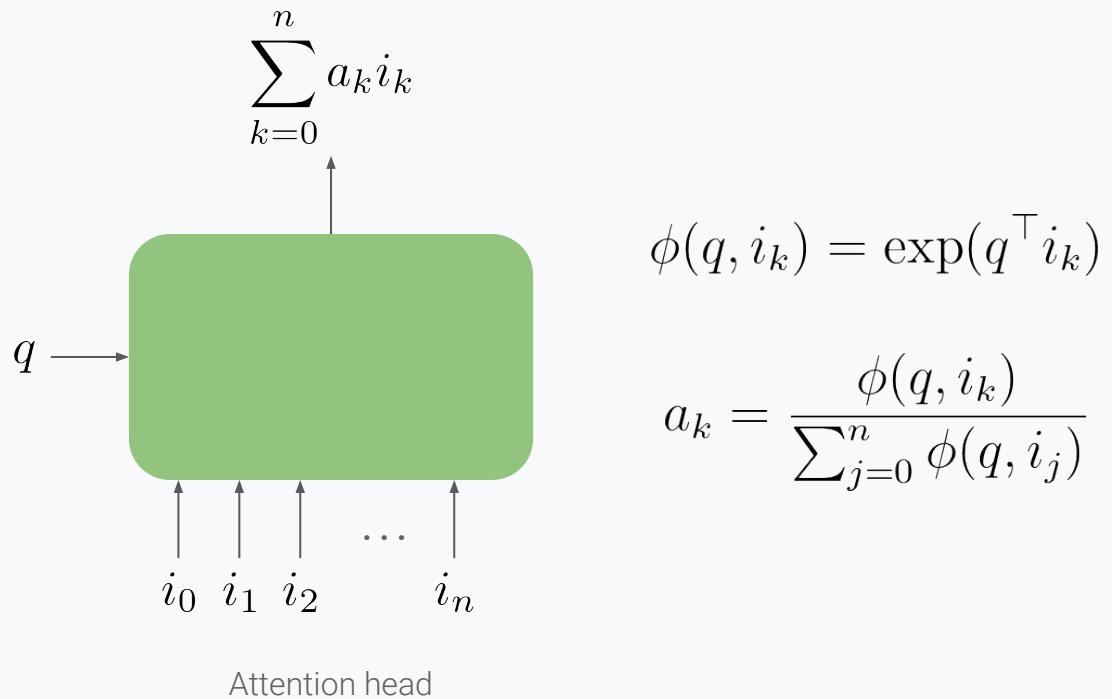
Effective approaches to
attention-based neural machine
translation. 2015



Dot product attention

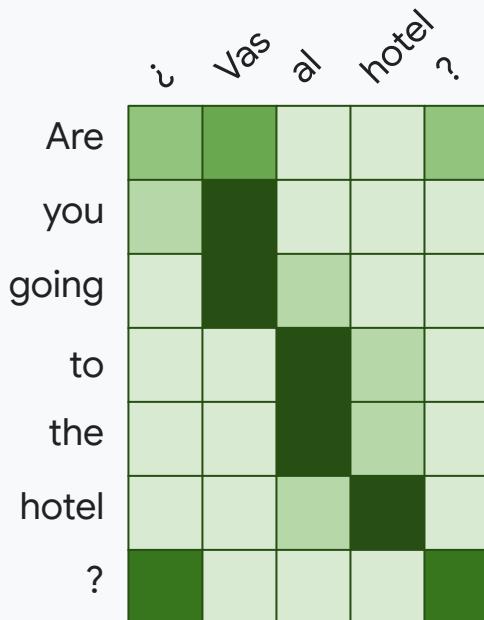
Thang Luong et al.

Effective approaches to
attention-based neural machine
translation. 2015



Attention mechanisms

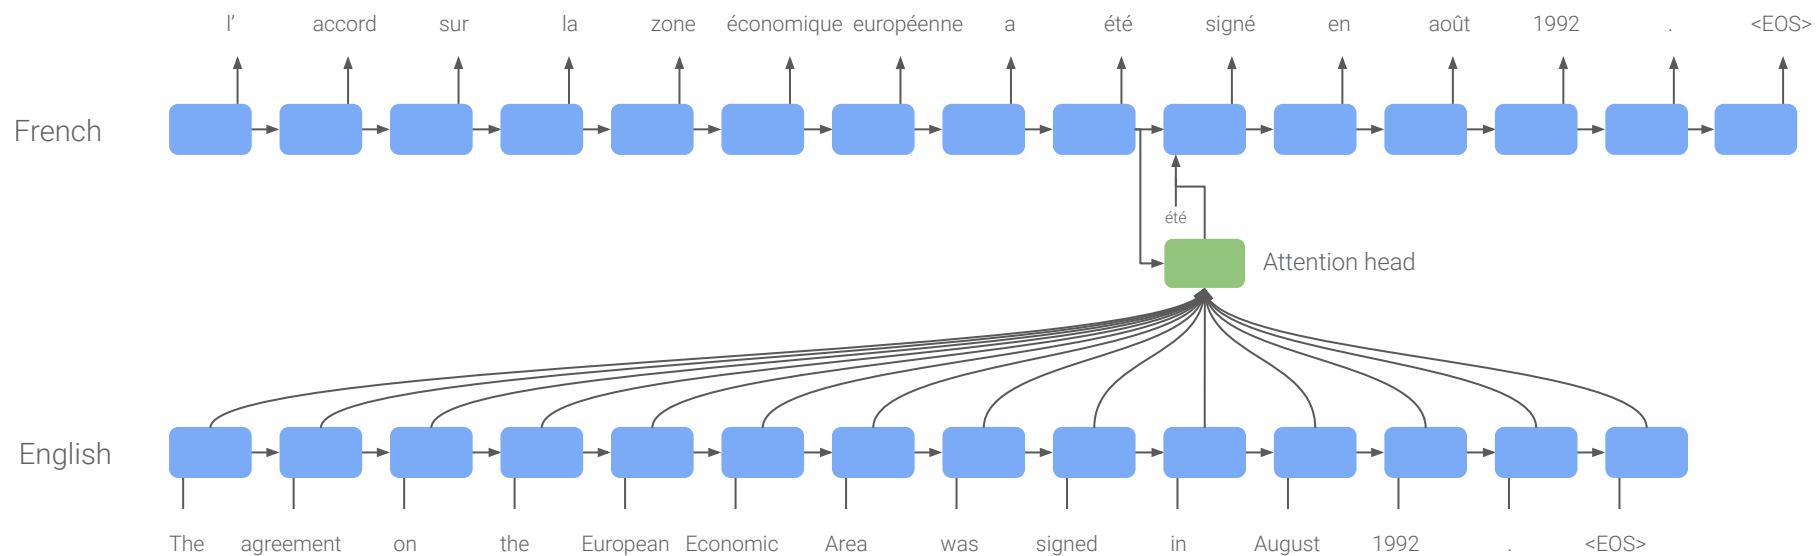
The attention matrix



Transformers



Motivation

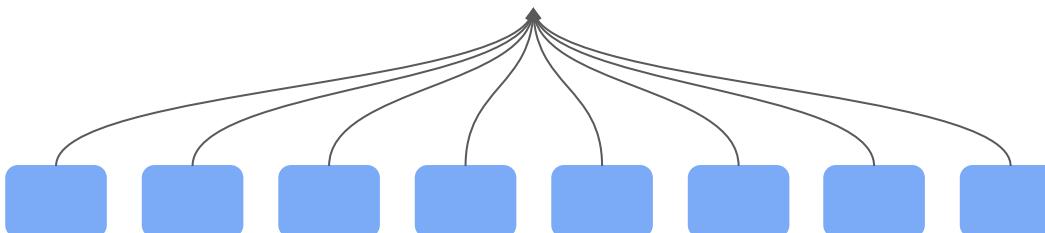


Example derived from Bahdanau, et al. 2014 (<https://arxiv.org/pdf/1409.0473.pdf>)

Motivation



sequential

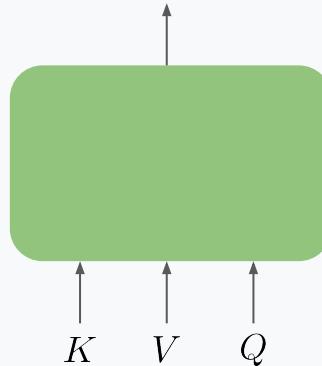


parallel

Scaled Dot-Product Attention

Queries, keys and values

A summary of values,
based on how similar their
corresponding keys are
with the query



Scaled

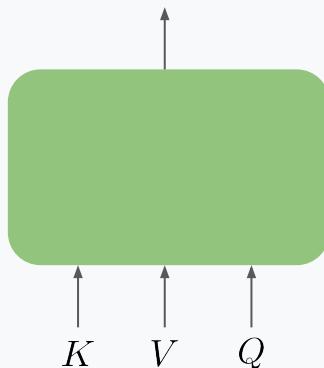
Dot-Product

Attention

Queries, keys and values

For some similarity
function $\underline{\phi}$

$$O_i = \sum_{j=0}^l a_{ij} V_j$$



$$a_{ij} = \frac{\phi(Q_i, K_j)}{\sum_{p=0}^l \phi(Q_i, K_p)}$$

Scaled

Dot-Product Attention

Attention

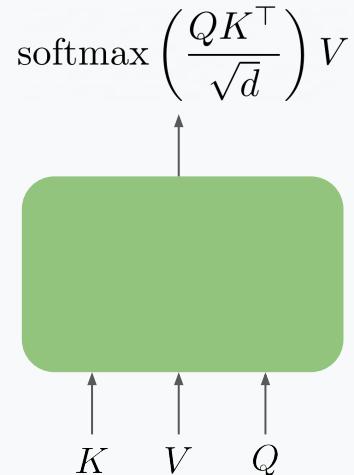
Using dot-product similarity,
we can vectorize nicely

$$\phi(Q_i, K_j) = \exp\left(\frac{Q_i K_j^\top}{\sqrt{d}}\right)$$

d = feature dim

$$\text{softmax}(x)_i = \frac{\exp x_i}{\sum_j \exp x_j}$$

Normalization factor for
numerical stability



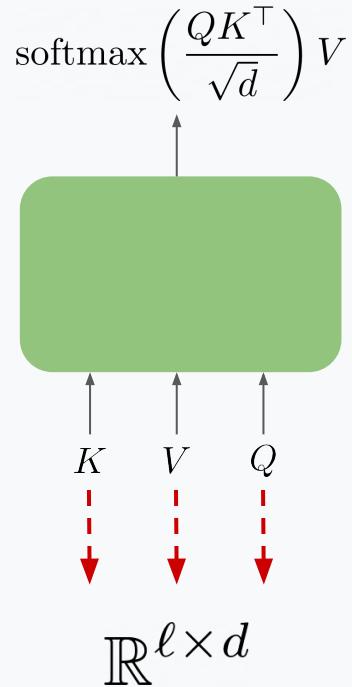
Scaled

Dot-Product Attention

Let's dive into the dimensions
(batch omitted for simplicity)

ℓ = sequence length

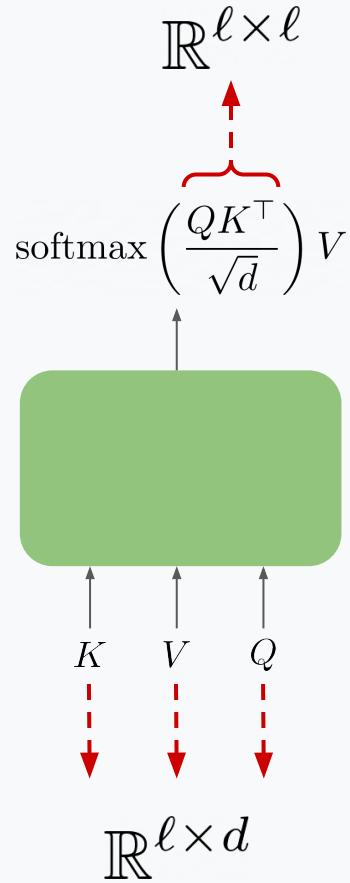
d = feature dim



Scaled Dot-Product Attention

Let's dive into the dimensions
(batch omitted for simplicity)

ℓ = sequence length
 d = feature dim

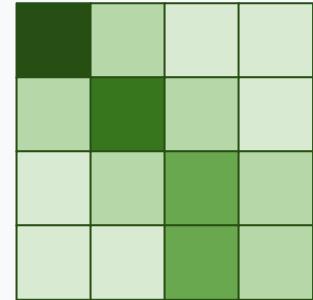
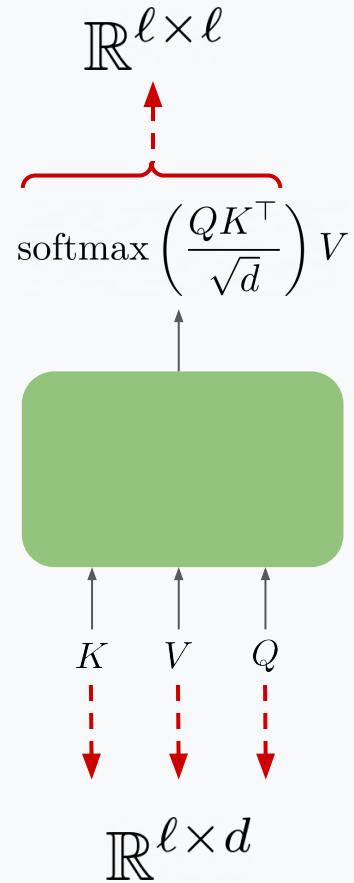


Scaled Dot-Product Attention

Let's dive into the dimensions
(batch omitted for simplicity)

ℓ = sequence length

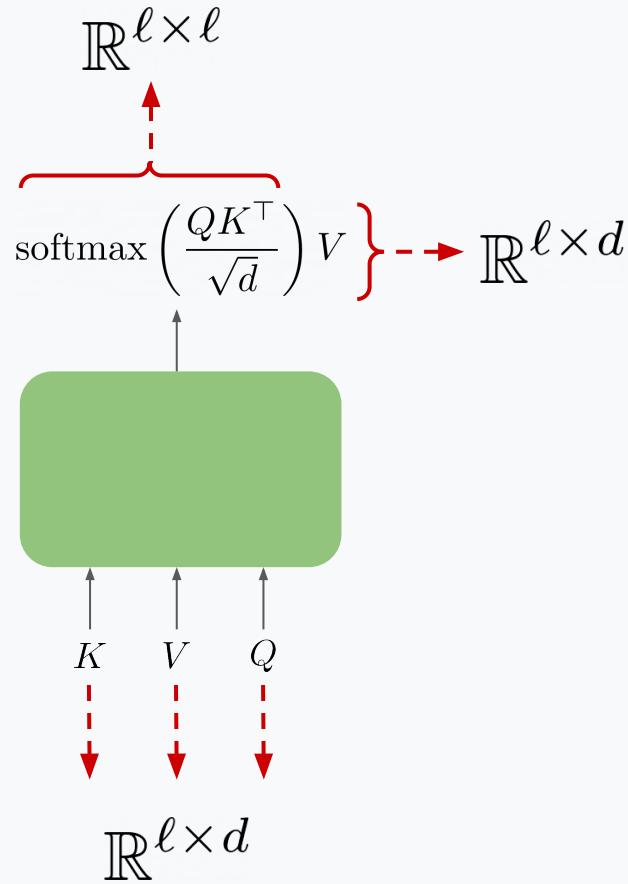
d = feature dim



Scaled Dot-Product Attention

Let's dive into the dimensions
(batch omitted for simplicity)

ℓ = sequence length
 d = feature dim

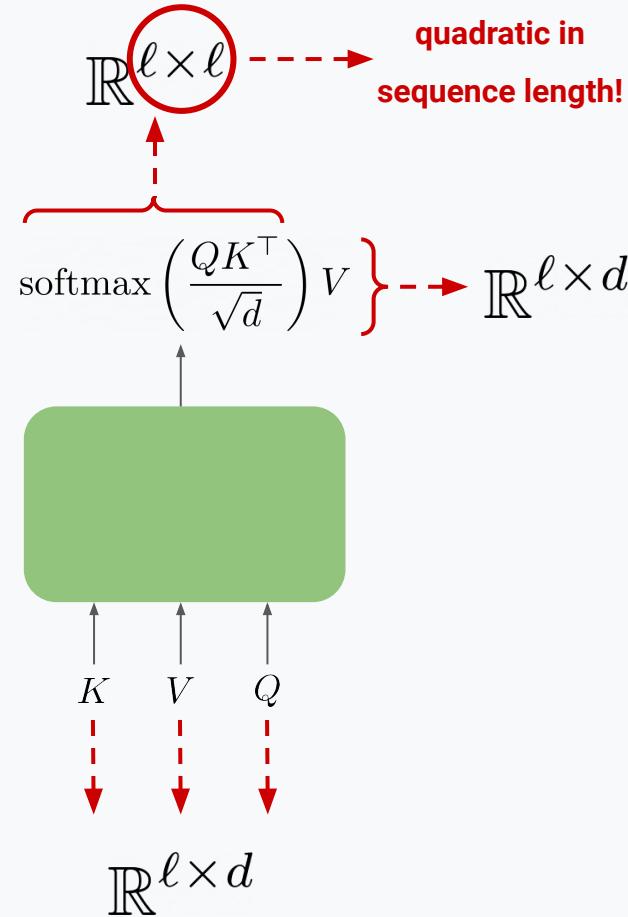


Scaled Dot-Product Attention

Let's dive into the dimensions
(batch omitted for simplicity)

ℓ = sequence length

d = feature dim

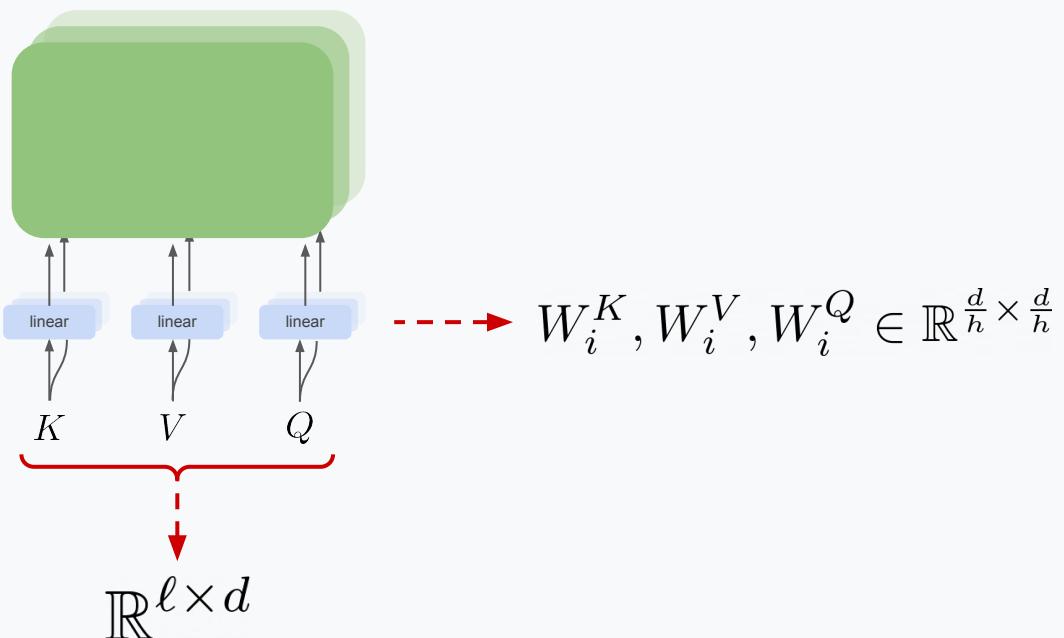


Multi-head attention

ℓ = sequence length

d = feature dim

h = # of attention heads

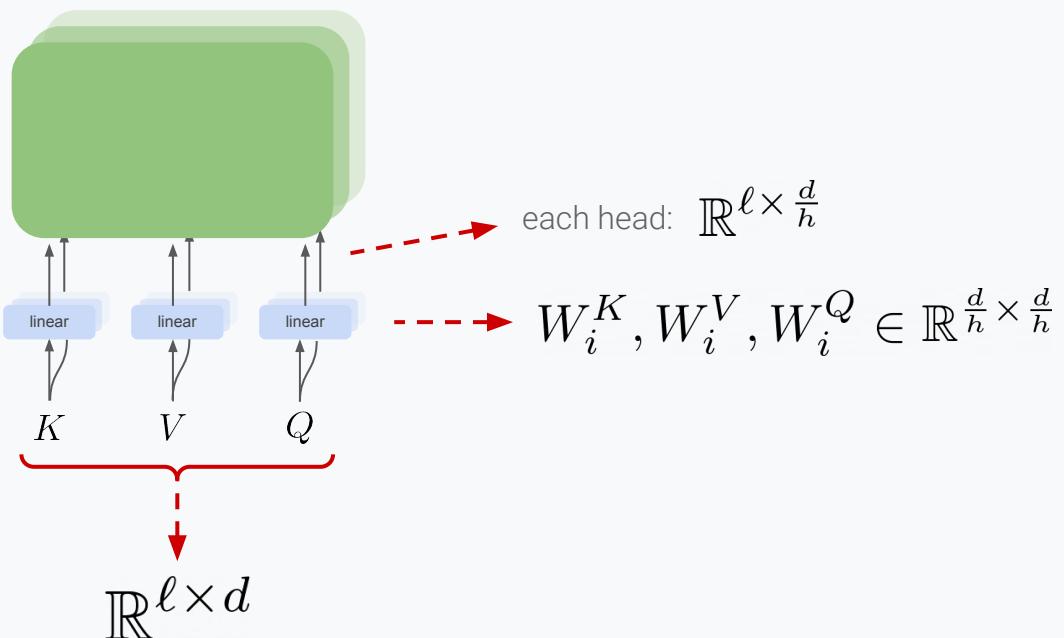


Multi-head attention

ℓ = sequence length

d = feature dim

h = # of attention heads

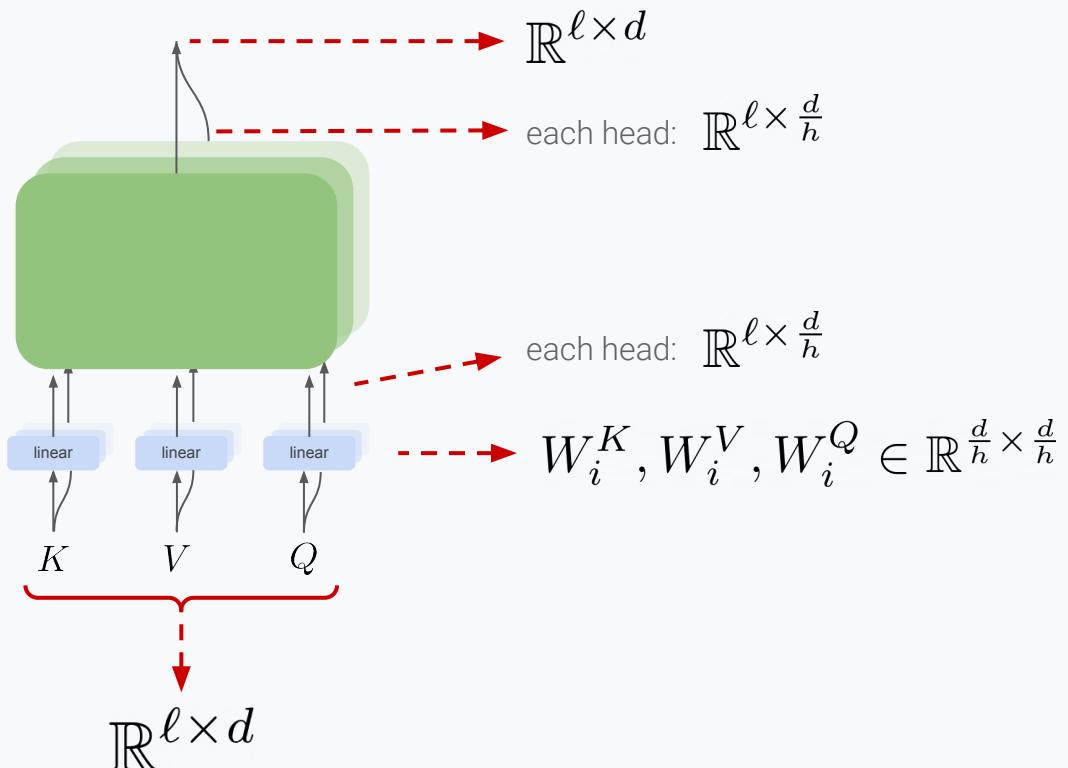


Multi-head attention

ℓ = sequence length

d = feature dim

h = # of attention heads

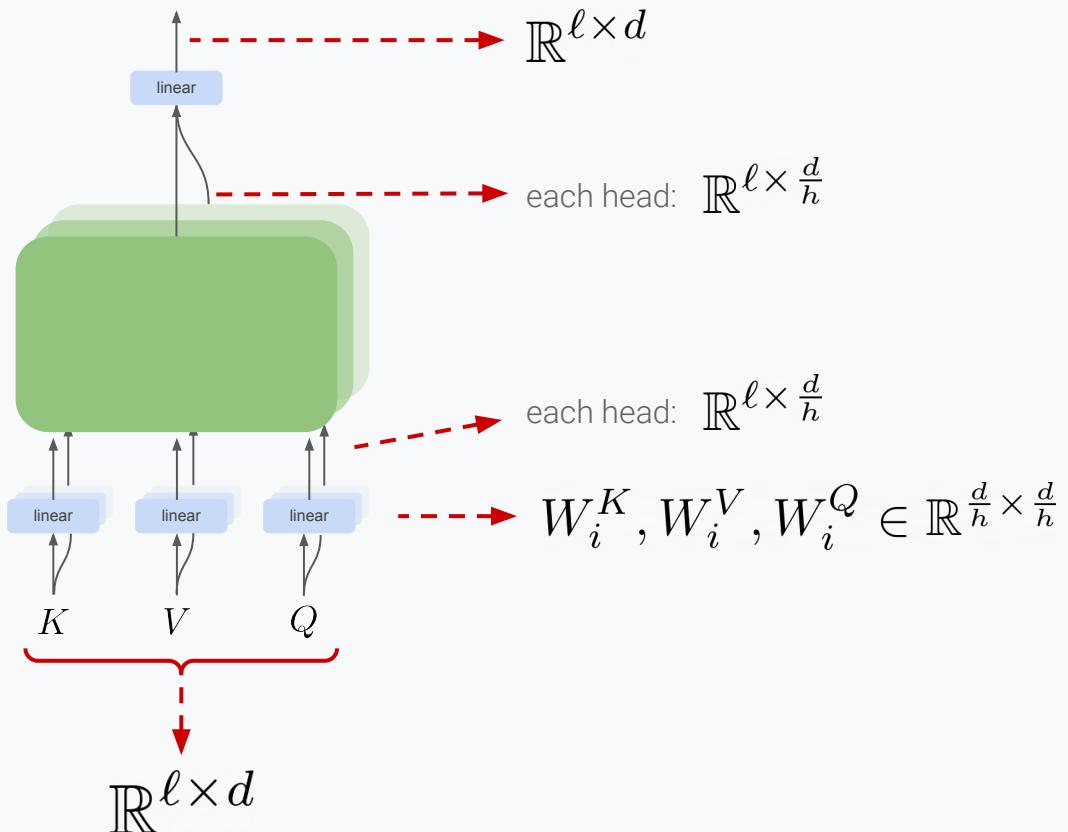


Multi-head attention

ℓ = sequence length

d = feature dim

h = # of attention heads

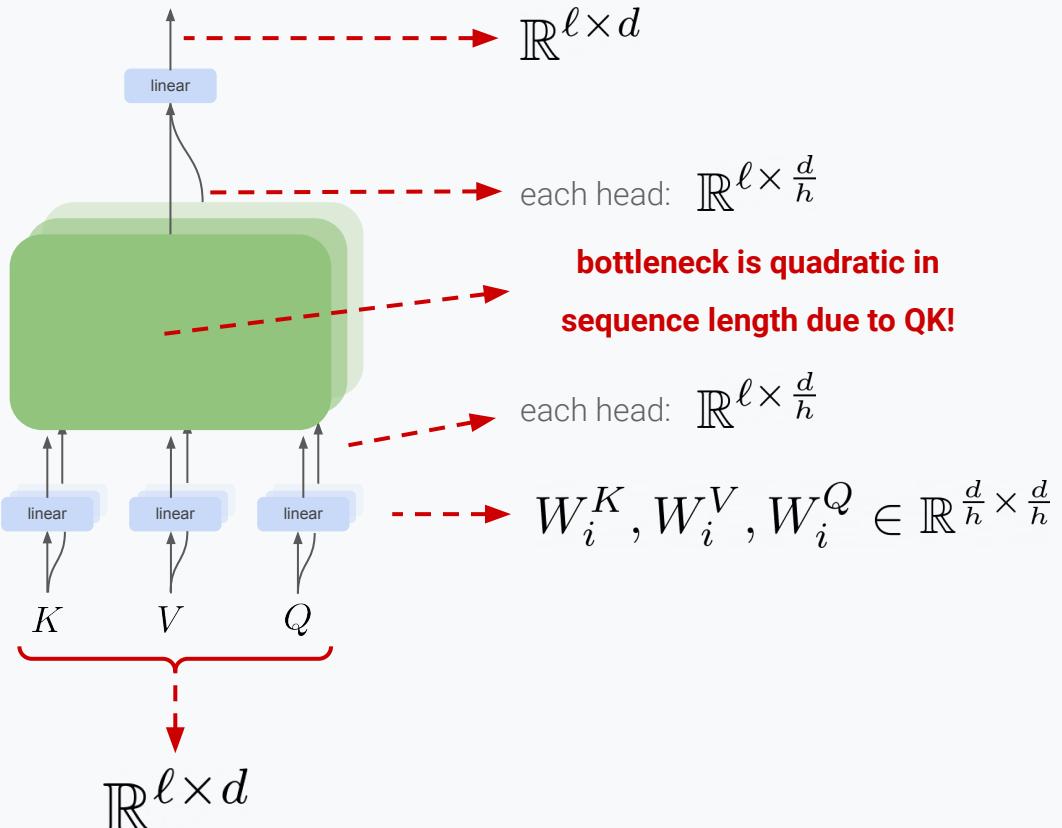


Multi-head attention

ℓ = sequence length

d = feature dim

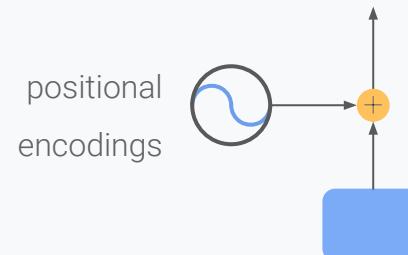
h = # of attention heads



Positional encodings

So far, attention has been a set operation.

Let's add positional information!



Positional encodings

So far, attention has been a set operation.

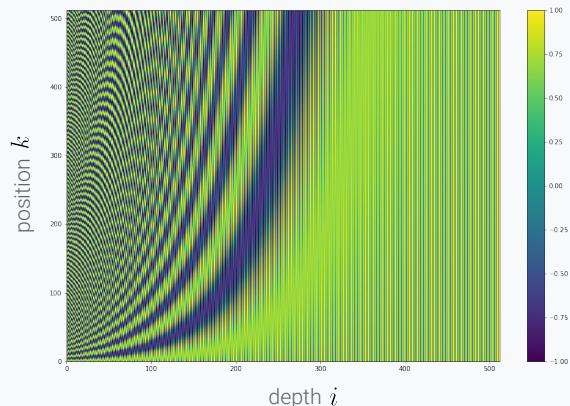
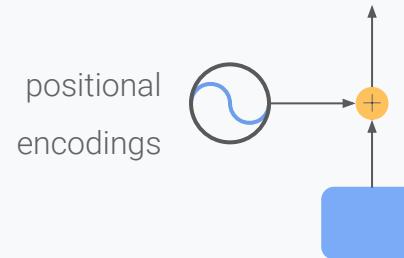
Let's add positional information!

These can be either **learned** or **fixed**.

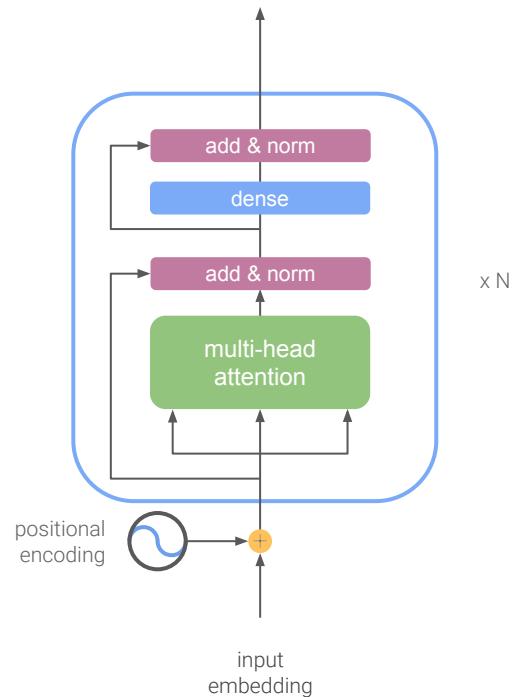
Fixed:

for a position k in the sequence and i in the feature space

$$E_{ki} = \begin{cases} \sin\left(k/10000^{\frac{i}{N}}\right) & \text{if } i \text{ is even} \\ \cos\left(k/10000^{\frac{i-1}{N}}\right) & \text{if } i \text{ is odd} \end{cases}$$

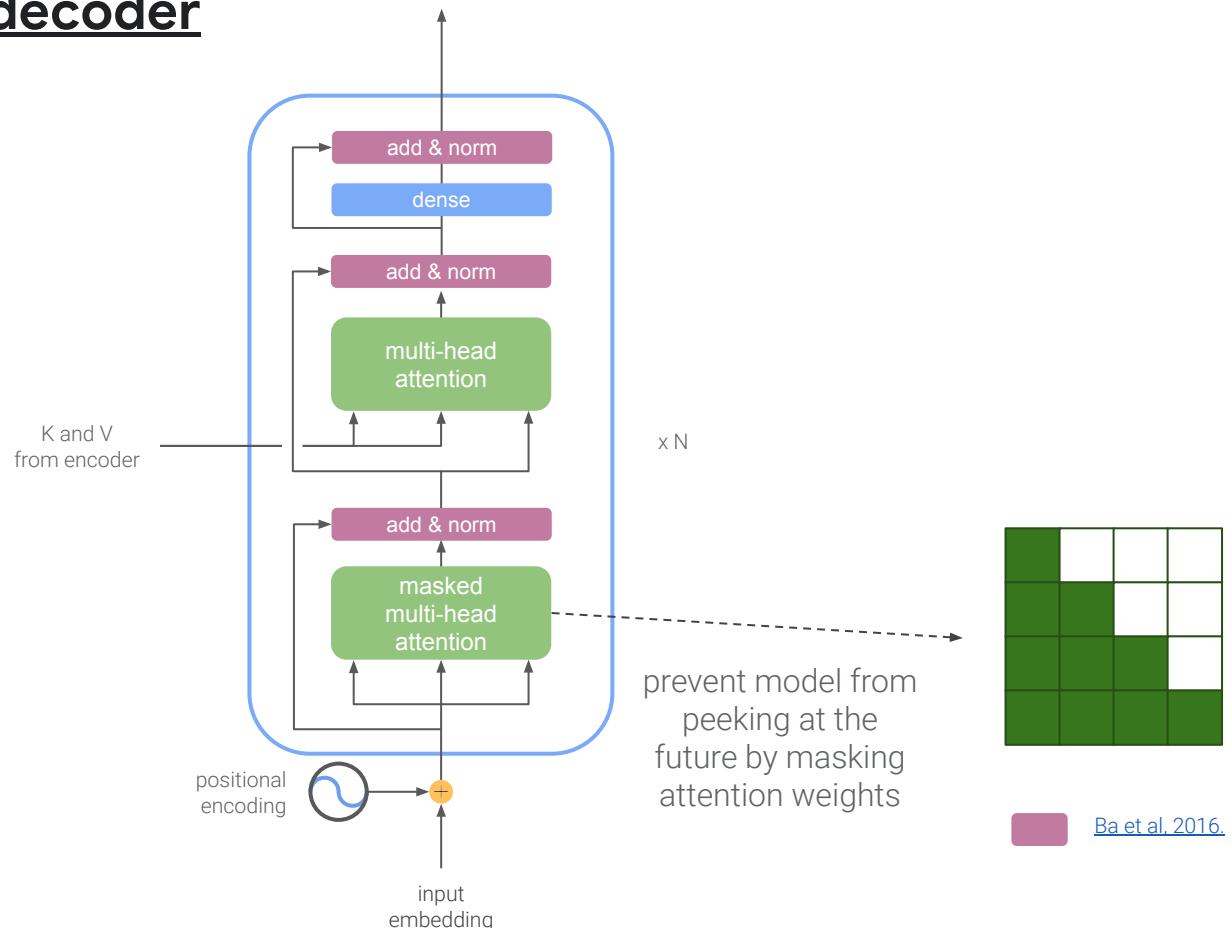


The transformer encoder

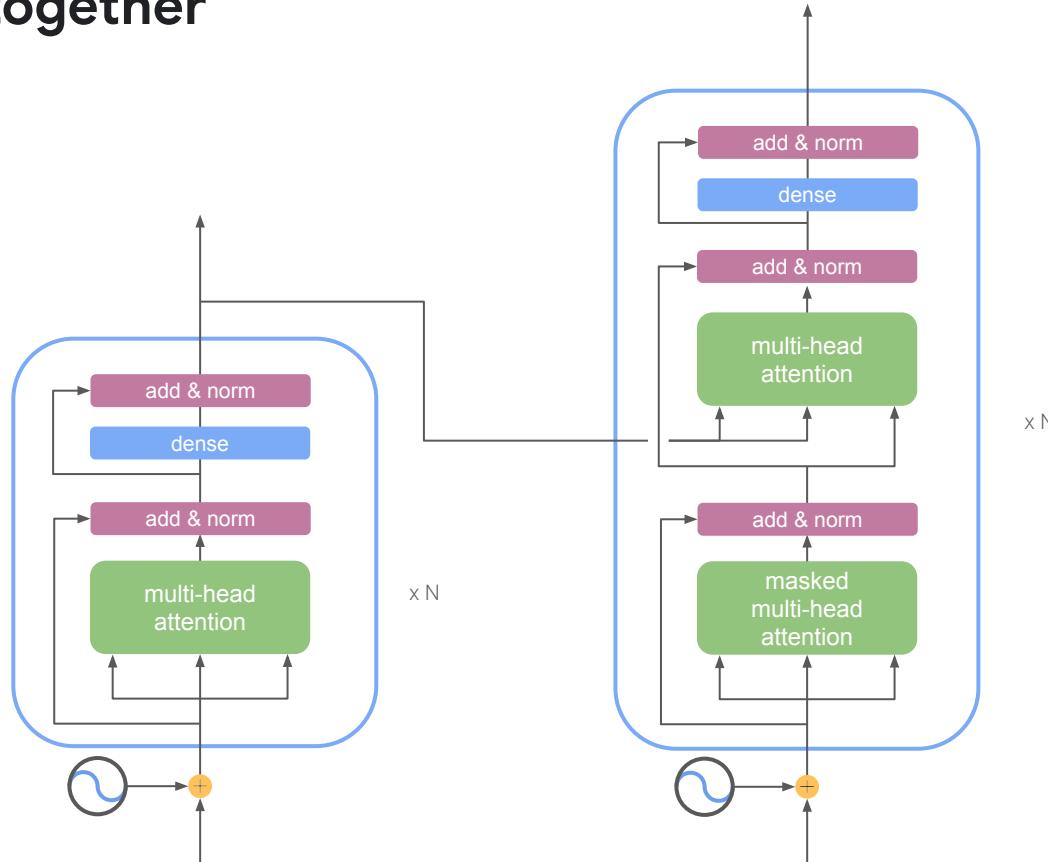


Ba et al. 2016.

The transformer decoder



Putting it all together



Transformers in recent literature

Transformers have become successful in a wide range of domains and applications, including:

- Mathematics and theorem proving (e.g. [Lample et al., 2019](#), [Clark et al., 2020](#))

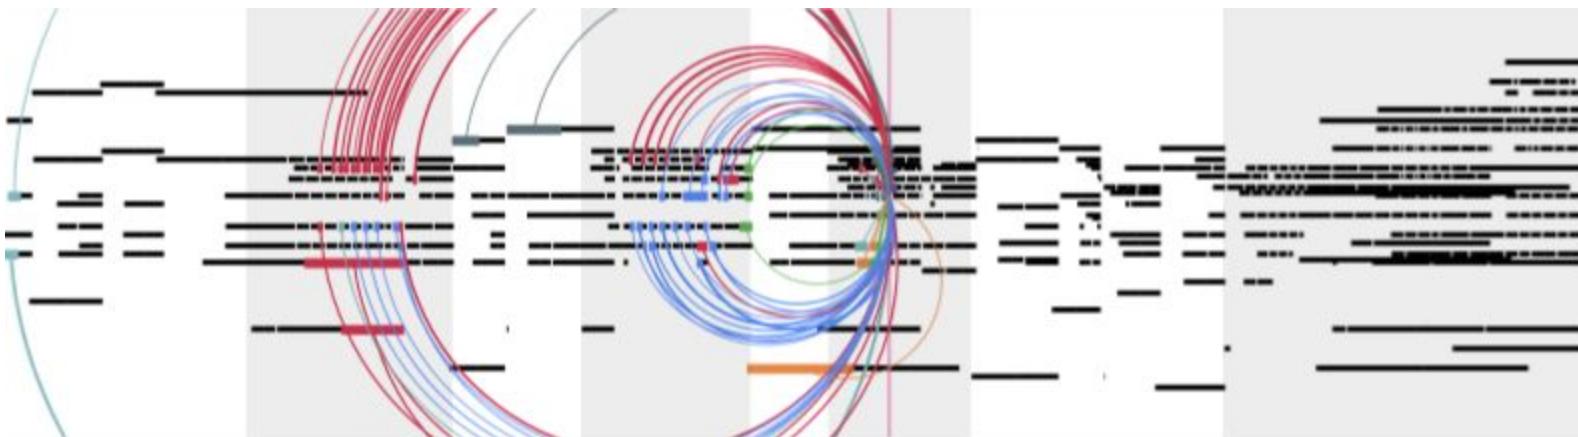
EQUATION	SOLUTION
$y' = \frac{16x^3 - 42x^2 + 2x}{(-16x^8 + 112x^7 - 204x^6 + 28x^5 - x^4 + 1)^{1/2}}$	$y = \sin^{-1}(4x^4 - 14x^3 + x^2)$
$3xy \cos(x) - \sqrt{9x^2 \sin(x)^2 + 1}y' + 3y \sin(x) = 0$	$y = c \exp(\sinh^{-1}(3x \sin(x)))$
$4x^4yy'' - 8x^4y'^2 - 8x^3yy' - 3x^3y'' - 8x^2y^2 - 6x^2y' - 3x^2y'' - 9xy' - 3y = 0$	$y = \frac{c_1 + 3x + 3\log(x)}{x(c_2 + 4x)}$

<https://ai.facebook.com/blog/using-neural-networks-to-solve-advanced-mathematics-equations/>

Transformers in recent literature

Transformers have become successful in a wide range of domains and applications, including:

- Mathematics and theorem proving (e.g. [Lample et al., 2019](#), [Clark et al., 2020](#))
- Music generation (e.g. [Anna Huang et al., 2019](#))

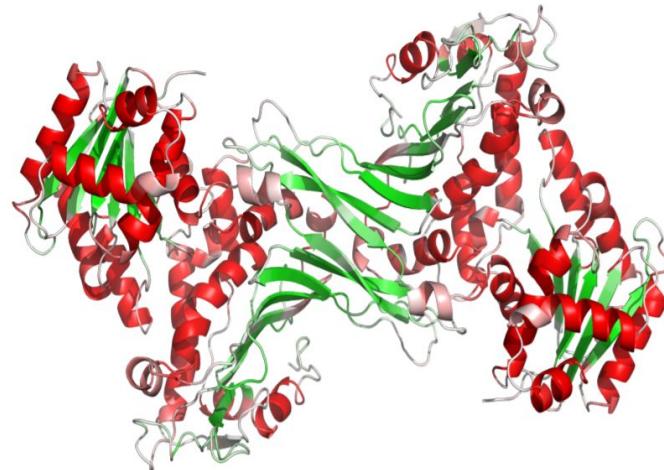


<https://magenta.tensorflow.org/music-transformer>

Transformers in recent literature

Transformers have become successful in a wide range of domains and applications, including:

- Mathematics and theorem proving (e.g. [Lample et al., 2019](#), [Clark et al., 2020](#))
- Music generation (e.g. [Anna Huang et al., 2019](#))
- Biology (e.g. [Rives et al., 2019](#), [Madani et al., 2020](#))



Transformers in recent literature

Transformers have become successful in a wide range of domains and applications, including:

- Mathematics and theorem proving (e.g. [Lample et al., 2019](#), [Clark et al., 2020](#))
- Music generation (e.g. [Anna Huang et al., 2019](#))
- Biology (e.g. [Rives et al., 2019](#), [Madani et al., 2020](#))
- Vision and Language (e.g. [Tan et al., 2019](#), [Lu et al., 2019](#), [Chen et al., 2020](#))



What color are her eyes?
What is the mustache made of?



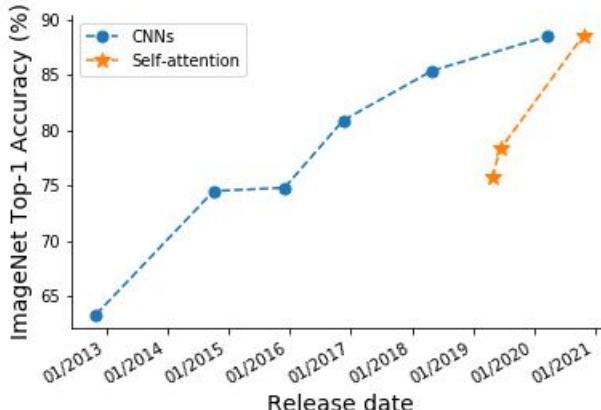
How many slices of pizza are there?
Is this a vegetarian pizza?

Visual Question Answering
([Agrawal et al., 2015](#))

Transformers in recent literature

Transformers have become successful in a wide range of domains and applications, including:

- Mathematics and theorem proving (e.g. [Lample et al., 2019](#), [Clark et al., 2020](#))
- Music generation (e.g. [Anna Huang et al., 2019](#))
- Biology (e.g. [Rives et al., 2019](#), [Madani et al., 2020](#))
- Vision and Language (e.g. [Tan et al., 2019](#), [Lu et al., 2019](#), [Chen et al., 2020](#))
- Computer Vision (e.g. [Ramachandran et al., 2019](#), [Dosovitskiy et al., 2020](#))



Transformers in NLP

Transformers are ubiquitous in NLP.

Large-scale pre-training has been enormously successful (e.g. [BERT](#), [ALBERT](#), [T5](#), [GPT-3](#)).

Models are typically used in 3 scenarios:

Transformers in NLP

Transformers are ubiquitous in NLP.

Large-scale pre-training has been enormously successful (e.g. [BERT](#), [ALBERT](#), [T5](#), [GPT-3](#)).

Models are typically used in 3 scenarios:

Pre-training

- Large corpus
(e.g. web crawled data)
- Typically unsupervised
(e.g. masked language modeling)
- Usually runs in GPUs or TPUs

Transformers in NLP

Transformers are ubiquitous in NLP.

Large-scale pre-training has been enormously successful (e.g. [BERT](#), [ALBERT](#), [T5](#), [GPT-3](#)).

Models are typically used in 3 scenarios:

Pre-training

- Large corpus
(e.g. web crawled data)
- Typically unsupervised
(e.g. masked language modeling)
- Usually runs in GPUs or TPUs

Fine-tuning

- Smaller corpus
- Typically supervised
(e.g. question answering, natural language inference)
- Usually runs in GPUs or TPUs

Transformers in NLP

Transformers are ubiquitous in NLP.

Large-scale pre-training has been enormously successful (e.g. [BERT](#), [ALBERT](#), [T5](#), [GPT-3](#)).

Models are typically used in 3 scenarios:

Pre-training

- Large corpus
(e.g. web crawled data)
- Typically unsupervised
(e.g. masked language modeling)
- Usually runs in GPUs or TPUs

Fine-tuning

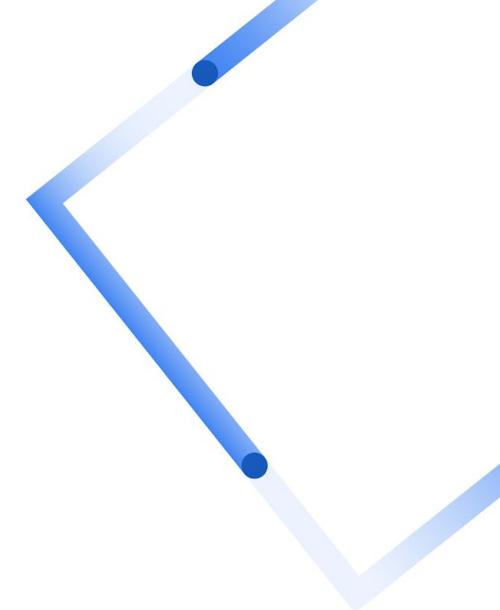
- Smaller corpus
- Typically supervised
(e.g. question answering, natural language inference)
- Usually runs in GPUs or TPUs

Production

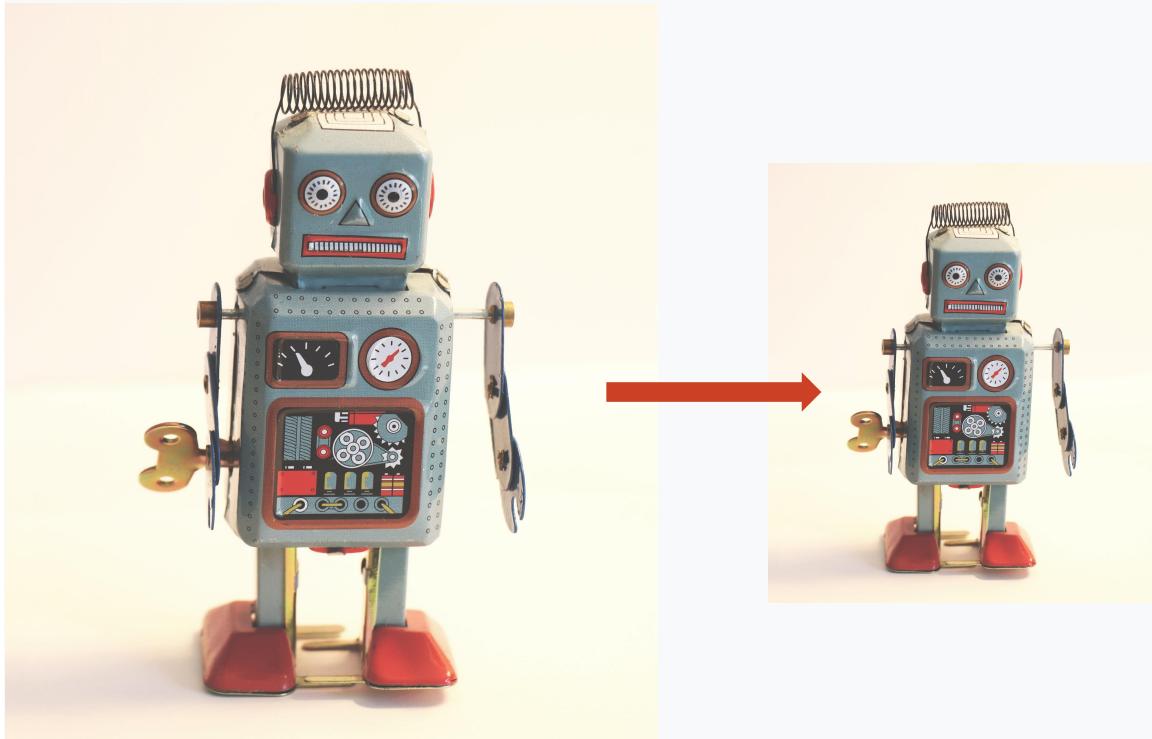
- Inference
- Usually runs in CPUs, sometimes in mobile devices

03

Core Techniques



Knowledge Distillation

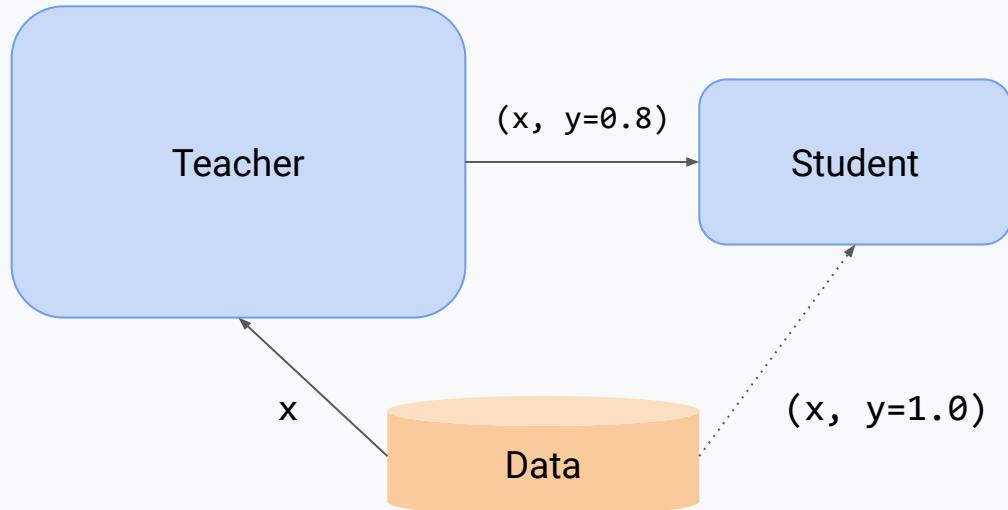


Source: unsplash.com

Knowledge Distillation

Hinton et al., 2015

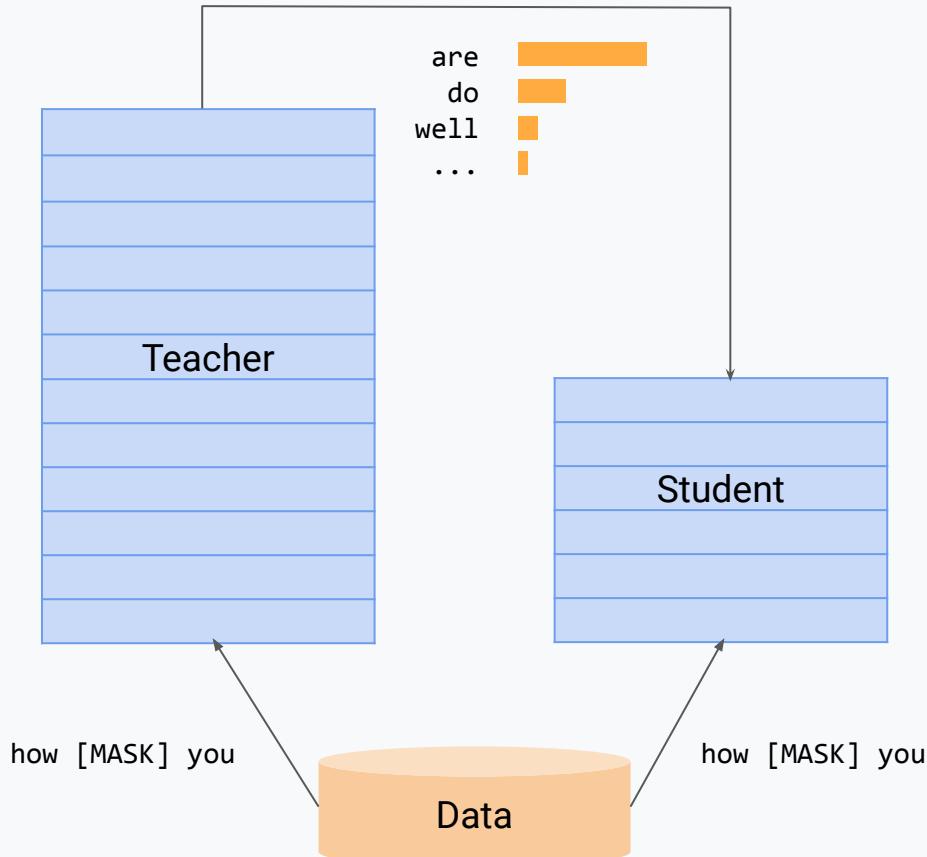
Distilling the Knowledge in a Neural Network



Knowledge Distillation for Pre-training

Sanh et al., 2019

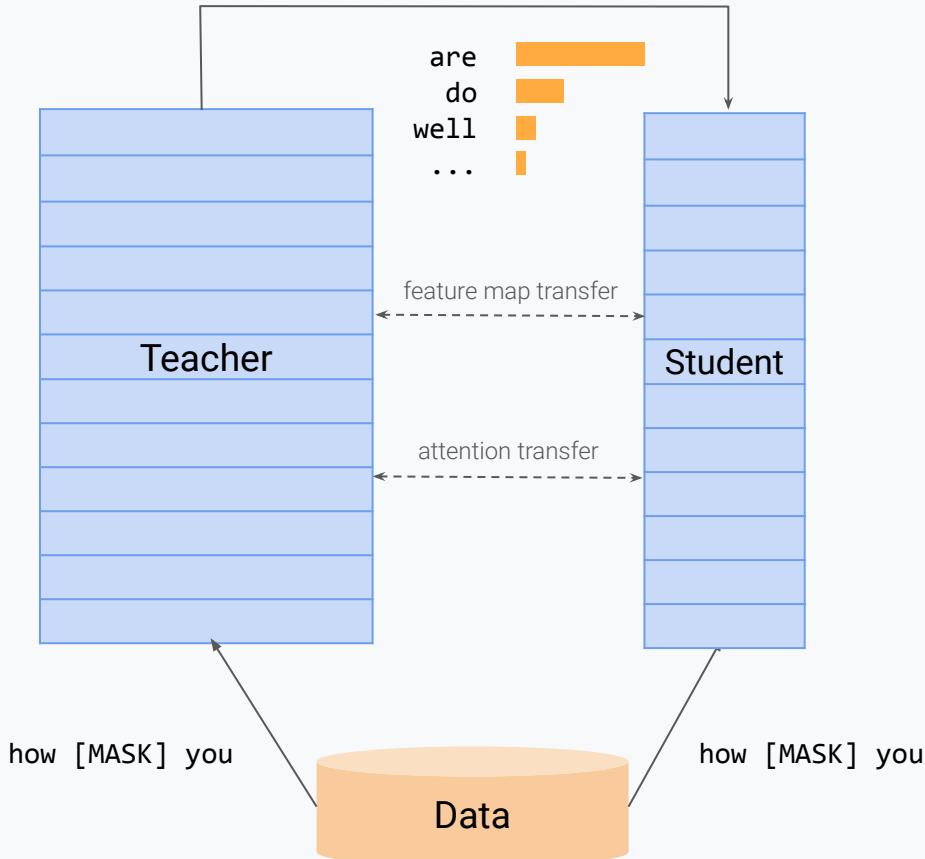
DistilBERT, a distilled version of BERT: smaller,
faster, cheaper and lighter



Knowledge Distillation for Pre-training

Sun et al., 2019

MobileBERT: a Compact Task-Agnostic BERT
for Resource-Limited Devices

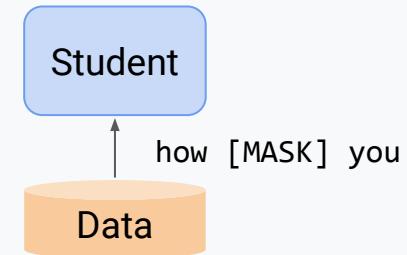


Knowledge Distillation for Fine-Tuning

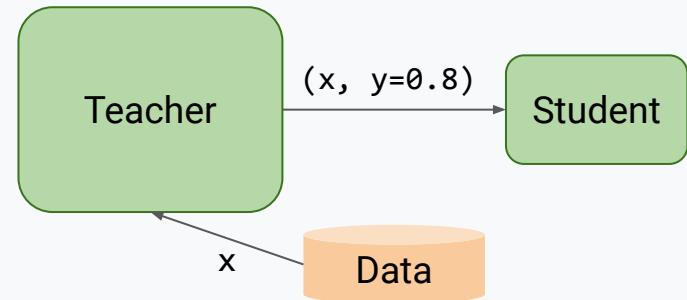
Turc et al., 2019

Well-Read Students Learn Better: On the
Importance of Pre-training Compact Models

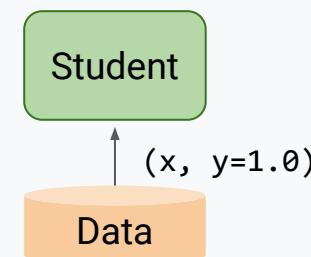
1 Regular Pre-training



2 Fine-tuning via distillation



3 (Optional) regular fine-tuning

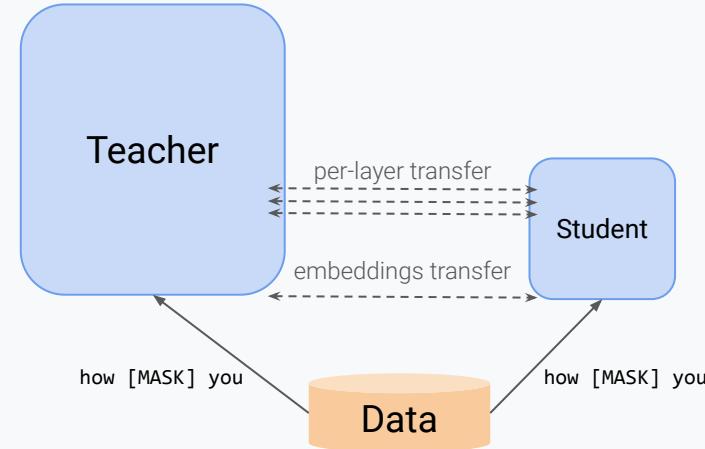


Knowledge Distillation for Pre-training and Fine-tuning

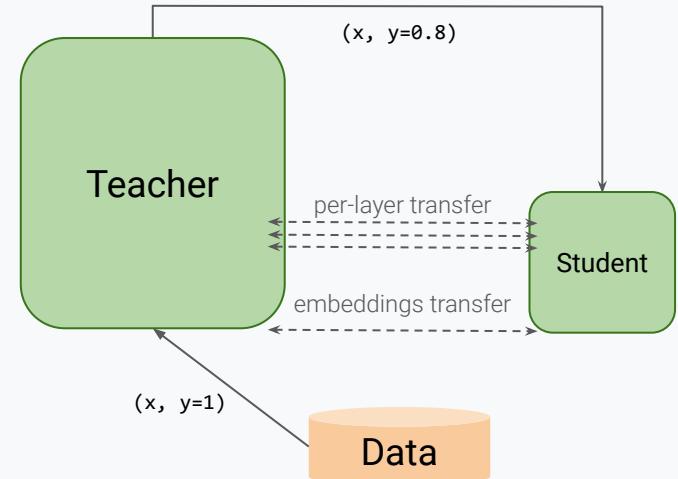
Jiao et al., 2019

TinyBERT: Distilling BERT for Natural Language Understanding

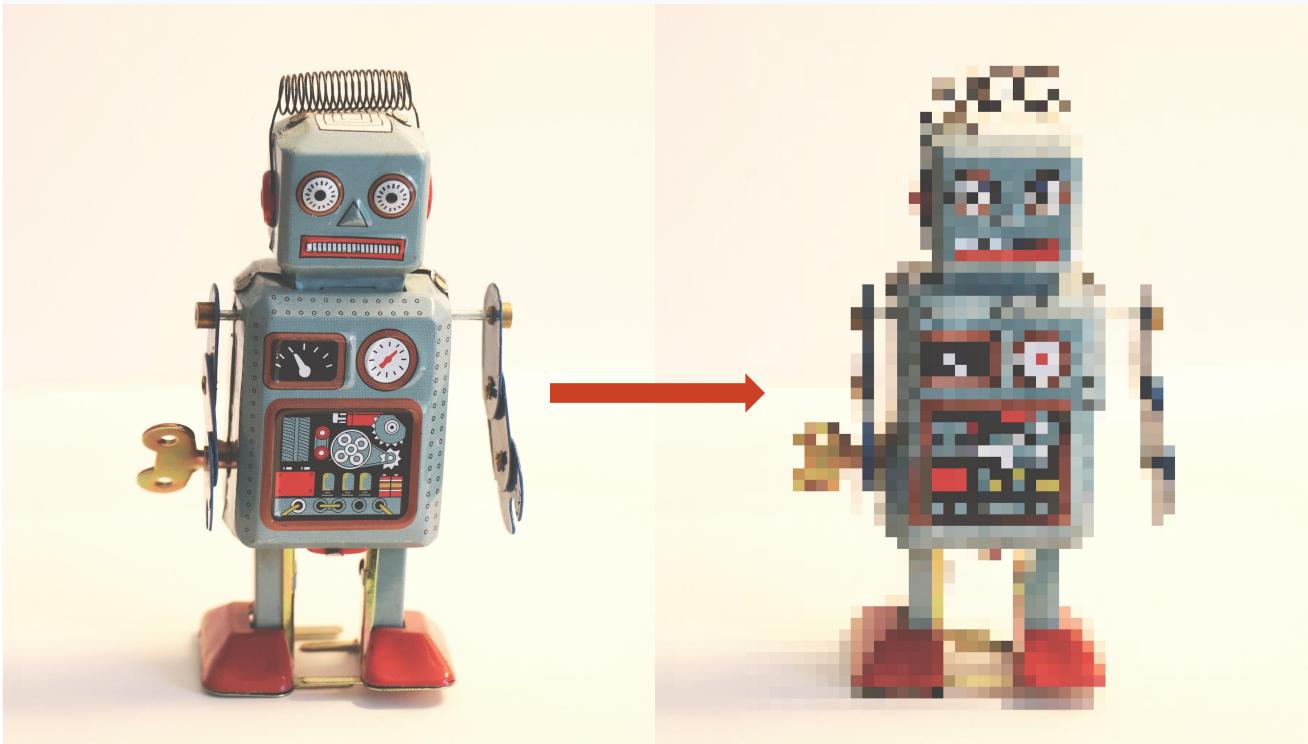
1 Pre-training via distillation



2 Fine-training via distillation



Quantization



Source: unsplash.com

Quantization

Definition

$$Q(z) = q_j$$

$$z \in (t_j, t_{j+1}]$$

$$j=0, \dots, 2^k-1$$

quantization operator

real-valued tensor (activation or weight)

quantization precision

Quantization

Definition

$$Q(z) = q_j \quad z \in (t_j, t_{j+1}] \quad j=0, \dots, 2^k-1$$

real-valued tensor (activation or weight)

quantization precision

quantization operator

```
graph TD; A[Q(z) = qj] --> C[j = 0, ..., 2k-1]; B[z ∈ (tj, tj+1)] --> C; C -- "quantization precision" --> D[real-valued tensor activation or weight]; E[quantization operator] --> F[j = 0, ..., 2k-1]
```

Linear Quantization

$$z = S (q_j - Z)$$

zero point

scaling factor

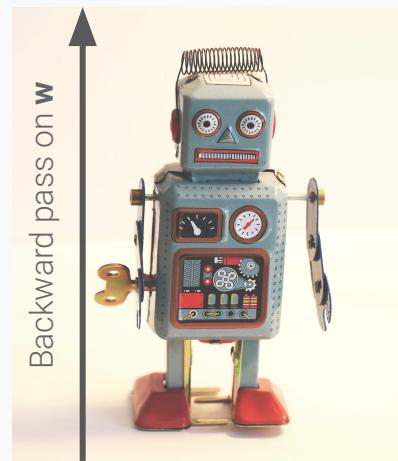
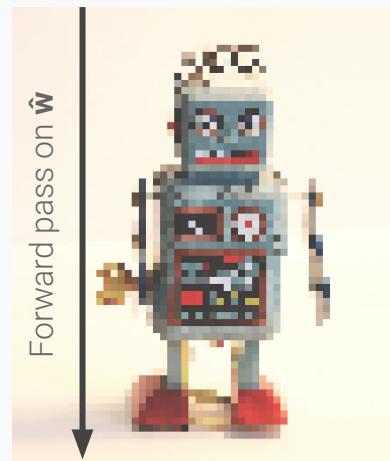
```
graph TD; A[z = S (qj - Z)] --> B[zero point]; A --> C[scaling factor]
```

Quantization

Jacob et al., 2017

Quantization and Training of Neural Networks
for Efficient
Integer-Arithmetic-Only Inference

Quantization-Aware Training



$$w^{t+1} = \text{UpdateParameter}(w^t, \frac{\partial L}{\partial \hat{w}^t}, \eta^t)$$

Quantization

Zafrir et al., 2019

[Q8BERT: Quantized 8Bit BERT](#)

Shen et al., 2019

[Q-BERT: Hessian Based Ultra Low Precision Quantization of BERT](#)

- **Q8BERT**: symmetric linear quantization:

$Q(z) = \text{clamp}(L_z \times S^z, -127, +127)$, where S^z is a statistic computed during or post-training.

- **Q-BERT**: uniform quantization to $\{0, \dots, 2^k-1\}$ with:
 - mixed precision (higher Hessian spectrum => higher precision for layer)
 - group precision (each matrix $W_k W_q W_v W_o$ is its own group)

Quantization with Distillation

Zhang et al., 2020

TernaryBERT: Distillation-aware Ultra-low Bit
BERT

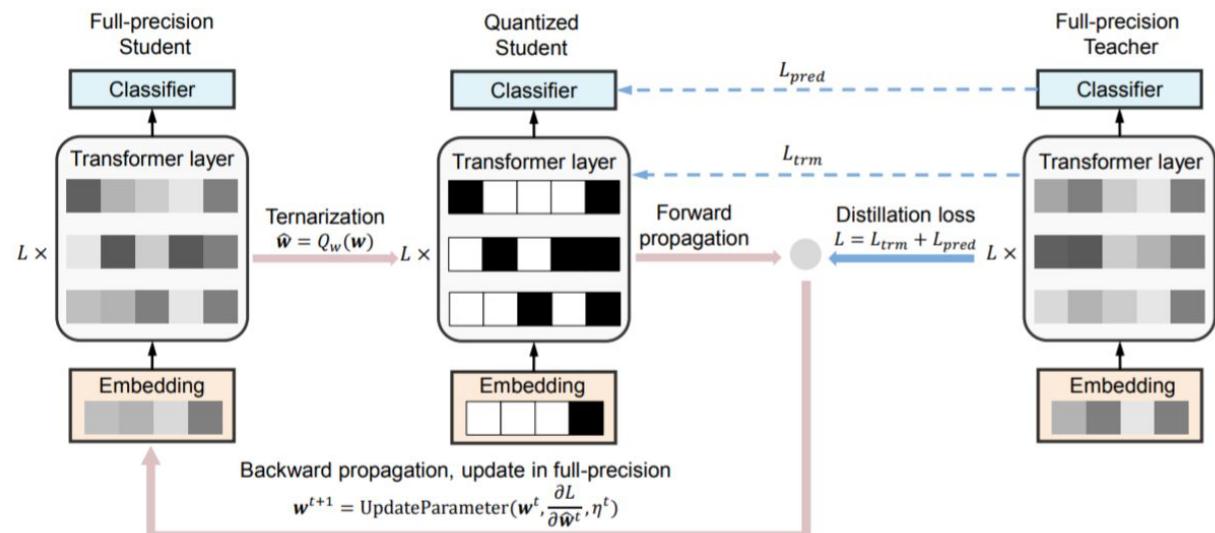
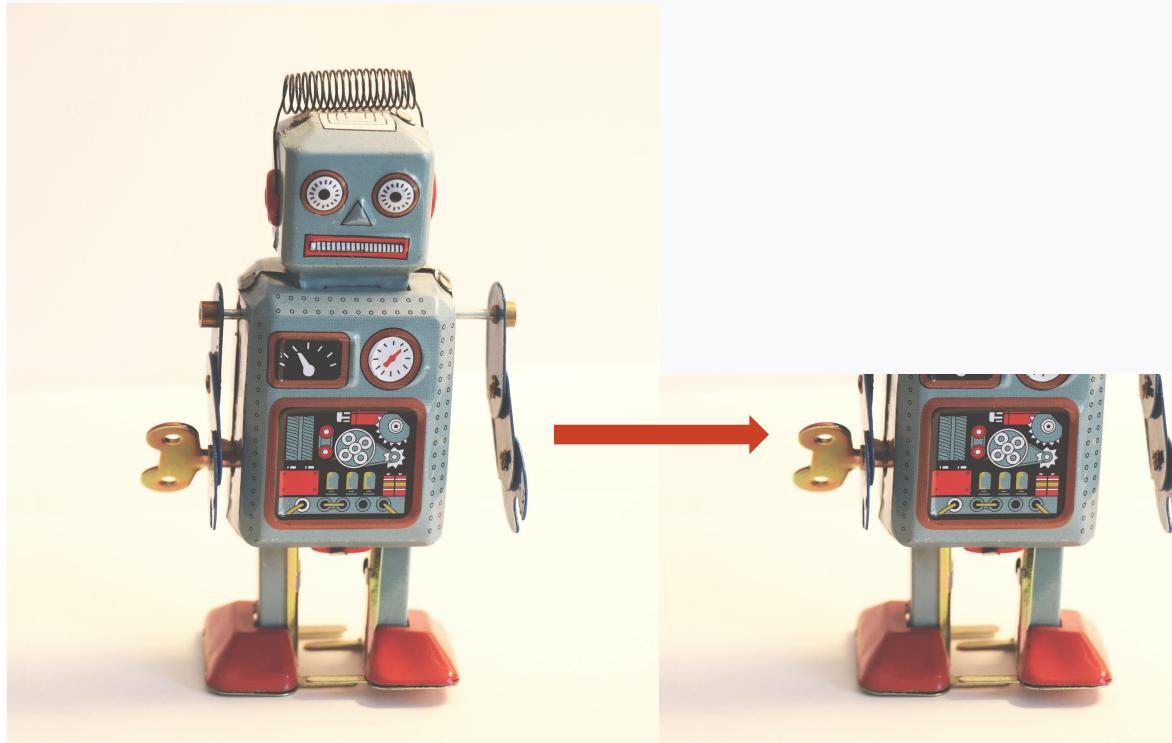


Figure 2: Depiction of the proposed distillation-aware ternarization of BERT model.

Pruning



Source: unsplash.com

Pruning

Definition

Pruning removes “unimportant” weights from a network:

$$a = (W \odot M) x$$

The diagram illustrates the computation of activation a from input x . It shows three components: W (model weight), M (pruning mask), and x (input). Arrows indicate the flow of data: one arrow from W to M , another from M to the activation a , and a third from x to a . The label "activation" is placed below a , and "model weight" is placed below W .

Main Questions [\(Hassibi and Stork\)](#)

- Which weights should be eliminated?
- How should the remaining weights be adjusted?
- How can such network pruning be done in an efficient way?

Pruning

Early Work

LeCun et al., 1990

OBD: [Optimal Brain Damage](#)

Hassibi and Stork, 1993

OBS: [Second order derivatives for network](#)

pruning: [Optimal Brain Surgeon](#)

Pruning based on second-order derivatives

Main idea:

- Start with a “reasonably large” network
- Train it to convergence
- Prune in multiple iterations, based on second-order derivatives:
 - OBD: prune and train
 - OBS: prune and update weights based on second-order statistics

Pruning

Early Work

LeCun et al., 1990

OBD: [Optimal Brain Damage](#)

Hassibi and Stork, 1993

OBS: [Second order derivatives for network](#)

pruning: [Optimal Brain Surgeon](#)

Pruning based on [second-order derivatives](#)

Main idea:

- Start with a “reasonably large” network
- Train it to convergence
- Prune in multiple iterations, based on second-order derivatives:
 - OBD: prune and train
 - OBS: prune and update weights based on second-order statistics

Why do we not train this smaller architecture instead?

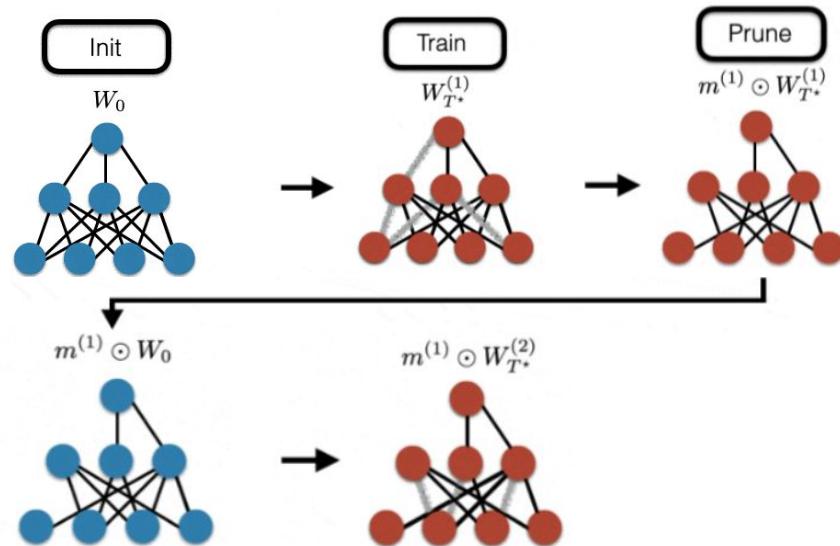
Pruning The LTH

Frankle and Carbin, 2018

The Lottery Ticket Hypothesis: Finding Sparse,
Trainable Neural Networks

The Lottery Ticket Hypothesis. A randomly-initialized, dense neural network contains a subnet-work that is initialized such that—when trained in isolation—it can match the test accuracy of the original network after training for at most the same number of iterations.

Searching for Tickets: **One-Shot** Magnitude Pruning



Frankle & Carbin, 2019
Viz: @RobertTLange

Source: <https://roberttlange.github.io/posts/2020/06/lottery-ticket-hypothesis/>

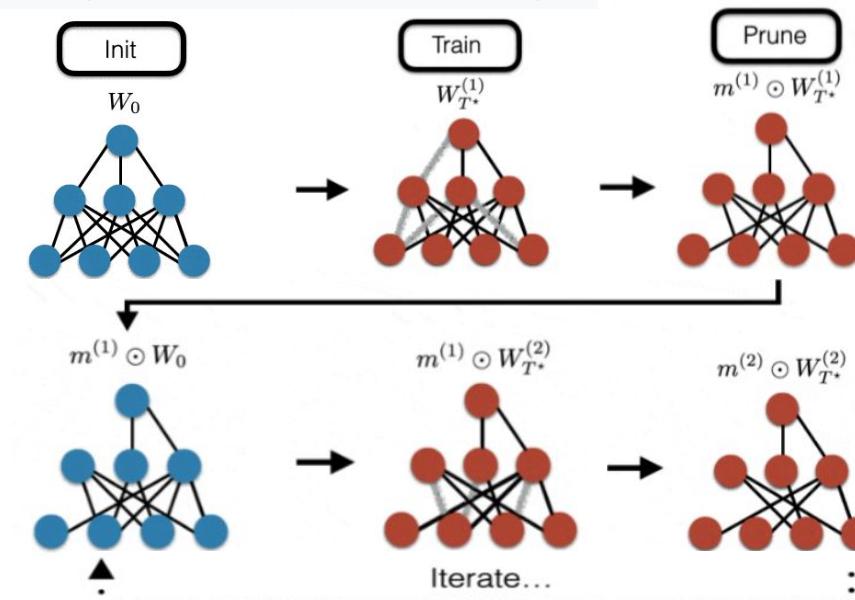
Pruning The LTH

Frankle and Carbin, 2018

The Lottery Ticket Hypothesis: Finding Sparse,
Trainable Neural Networks

The Lottery Ticket Hypothesis. A randomly-initialized, dense neural network contains a subnet-work that is initialized such that—when trained in isolation—it can match the test accuracy of the original network after training for at most the same number of iterations.

Searching for Tickets: **Iterative** Magnitude Pruning



Frankle & Carbin, 2019
Viz: @RobertTLange

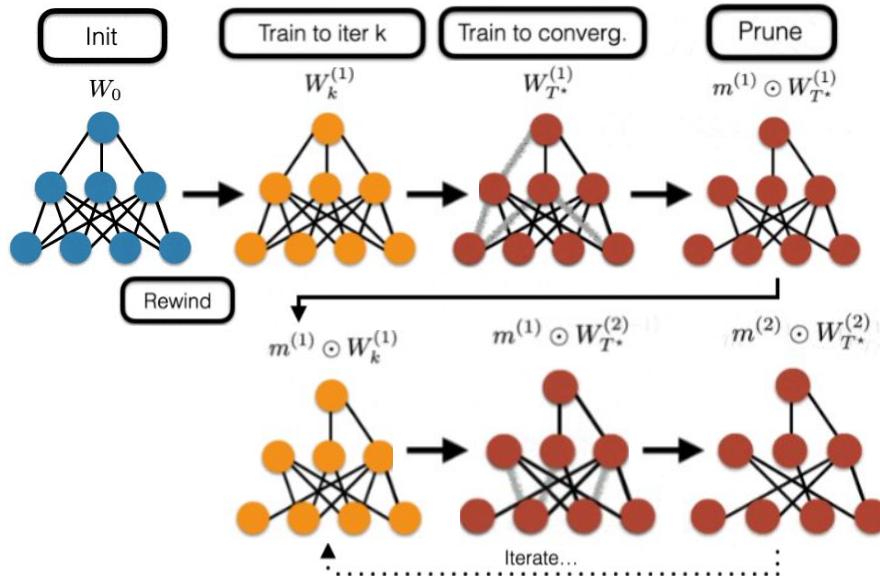
Pruning The LTH

Frankle et al, 2019

Stabilizing the Lottery Ticket Hypothesis

The Lottery Ticket Hypothesis. A randomly-initialized, dense neural network contains a subnet-work that is initialized such that—when trained in isolation—it can match the test accuracy of the original network after training for at most the same number of iterations.

Searching for Tickets: **Iterative** Magnitude Pruning **with Rewinding**



Frankle et al., 2019
Viz: @RobertTLange

Source: <https://roberttlange.github.io/posts/2020/06/lottery-ticket-hypothesis/>

Pruning

The LTH, ctd

Brix et al., 2020

Successfully Applying the Stabilized Lottery
Ticket Hypothesis to the Transformer
Architecture

Sparsity Memory		MP		LT		SLT		CLT		SLT-MP		MP-SLT	
		BLEU	TER	BLEU	TER	BLEU	TER	BLEU	TER	BLEU	TER	BLEU	TER
0%	234 MB	26.8	64.5	26.8	64.5	26.8	64.5	26.8	64.5	26.8	64.5	26.8	64.5
10%	226 MB	26.8	64.5	26.7	64.6	26.8	64.9	26.9	64.7	n/a	n/a	26.8	64.5
20%	206 MB	26.7	64.5	26.2	65.3	26.9	64.6	27.0	64.5	n/a	n/a	26.7	64.5
30%	184 MB	26.4	65.0	26.0	65.3	26.9	64.8	26.9	64.7	n/a	n/a	26.4	65.0
40%	161 MB	26.5	64.8	25.8	65.7	27.1	65.1	26.8	65.0	n/a	n/a	26.5	64.8
50%	137 MB	26.4	65.0	25.4	66.3	26.6	65.2	26.7	65.2	26.4 [†]	64.9 [†]	26.4	65.0
60%	112 MB	25.9	65.5	24.9	66.5	26.4	65.7	26.8	65.0	26.4 [†]	65.1 [†]	25.9	65.5
70%	86 MB	25.7	65.8	24.2	67.6	25.6	66.9	26.2	65.8	26.2 [‡]	65.3 [‡]	25.6	66.0
80%	59 MB	24.8	66.8	23.2	68.4	24.8	67.7	24.1	67.9	25.6 [‡]	65.9 [‡]	24.6	67.2
85%	46 MB	23.9	67.7	22.3	69.8	23.7	68.5	23.7	68.0	24.9 [‡]	66.4 [‡]	23.9	67.9
90%	31 MB	22.9	69.0	20.9	72.0	21.7	71.4	21.6	70.6	23.5 [‡]	68.4 [‡]	22.4	69.8
95%	17 MB	20.2	72.9	18.1	75.4	17.4	77.1	18.2	73.3	20.5 [‡]	72.3 [‡]	18.5	75.5
98%	7 MB	15.8	78.9	13.3	81.2	11.0	86.9	14.6	78.2	16.1 [‡]	79.2 [‡]	13.5	82.6

Table 1: En→De translation: BLEU [%] and TER [%] scores of the final model at different sparsity levels, evaluated on newstest2014. For SLT-MP, models marked with † are trained with SLT pruning, models marked with ‡ are trained with MP. For MP-SLT, the MP model with 60% sparsity was used for SLT pruning. For each sparsity level, the best score is highlighted.

MP = Magnitude Pruning

LT = Lottery Ticket

SLT = Stabilized Lottery Ticket

CLT = Constant Lottery Ticket

Pruning

Sanh et al., 2020

[Movement Pruning: Adaptive Sparsity by Fine-Tuning](#)

Movement Pruning

- **First-order** strategy: “instead of selecting weights that are far from zero, we retain connections that are moving away from zero during the training process”
- The pruning mask \mathbf{M} is learnt together with the model parameters.
 - **hard** version: $\mathbf{M} = \text{Top}_{\mathbf{v}}(\mathbf{S})$, where score \mathbf{S} is learnt and \mathbf{v} is a hyperparameter.
 - **soft** version: $\mathbf{M} = (\mathbf{S} > \tau)$, where score \mathbf{S} is learnt and threshold τ is a hyperparameter.

Pruning & Hardware

Hooker, 2020

The Hardware Lottery

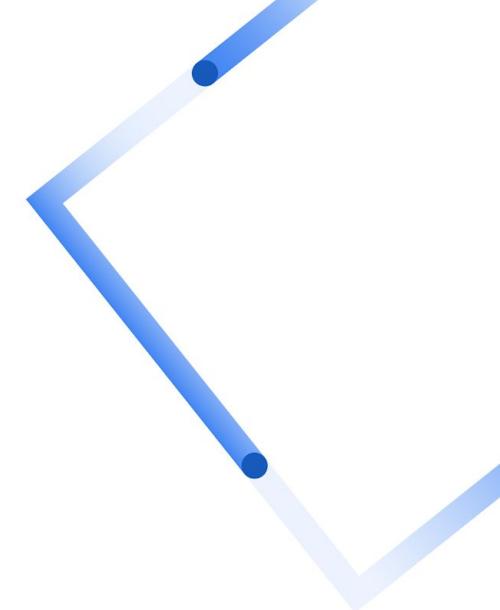
On standard hardware:

	Unstructured Pruning	Structured Pruning
Storage	✓	✓
Inference	✗	✓
Flexibility	✓	✗

- “In many ways, hardware is catching up to the present state of ML research.”
- There is research for specialized software kernels to support unstructured sparsity (see paper for references).

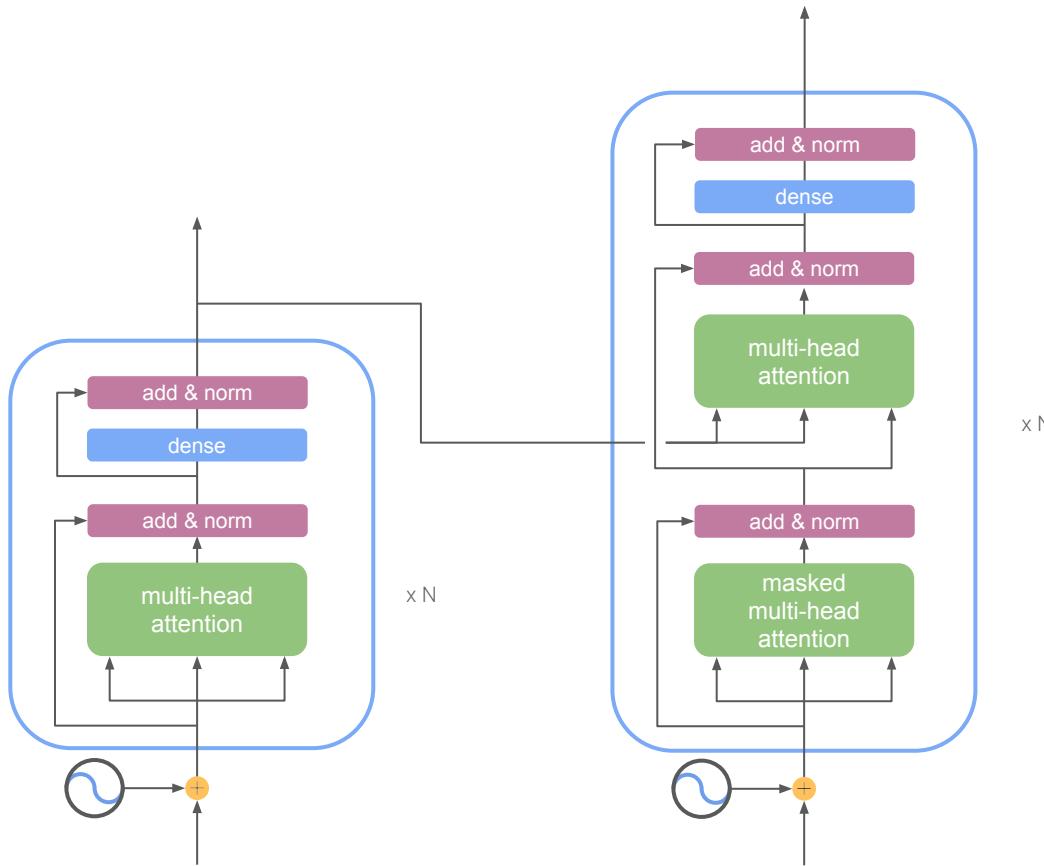
04

Efficient Attention



Recap

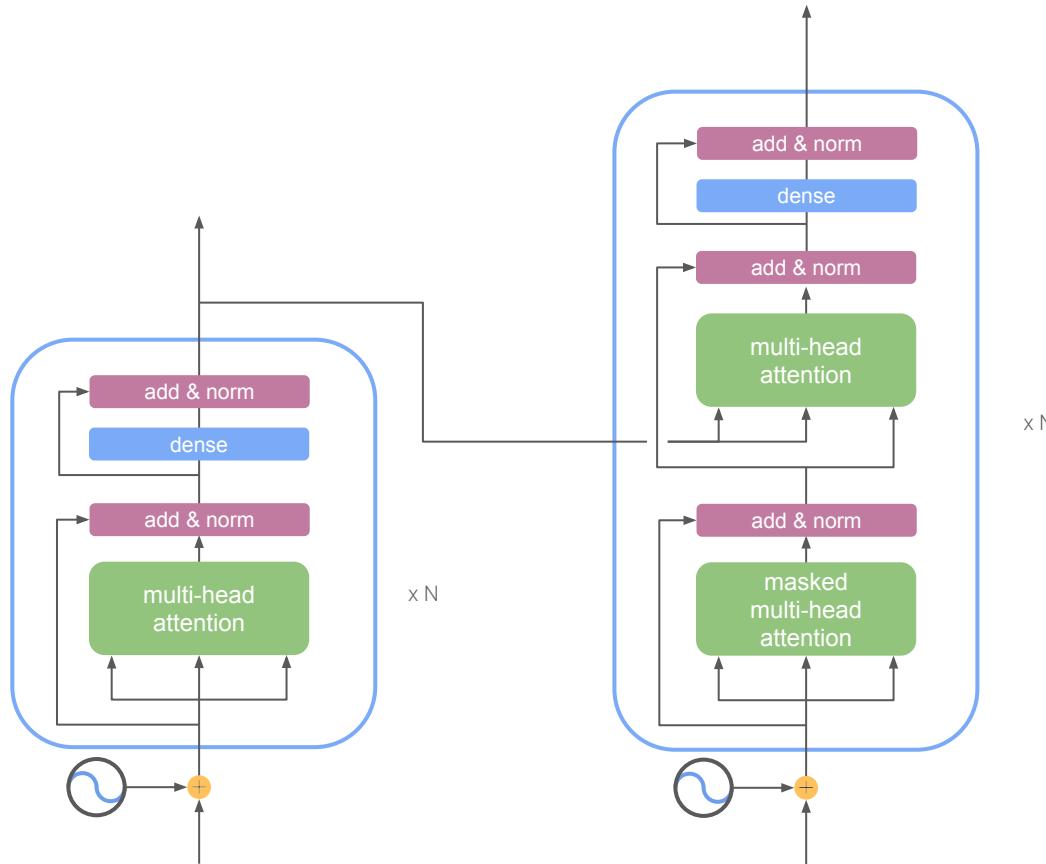
The Transformer architecture



Recap

The Transformer architecture

Quadratic bottleneck in sequence length due to multi-head attention



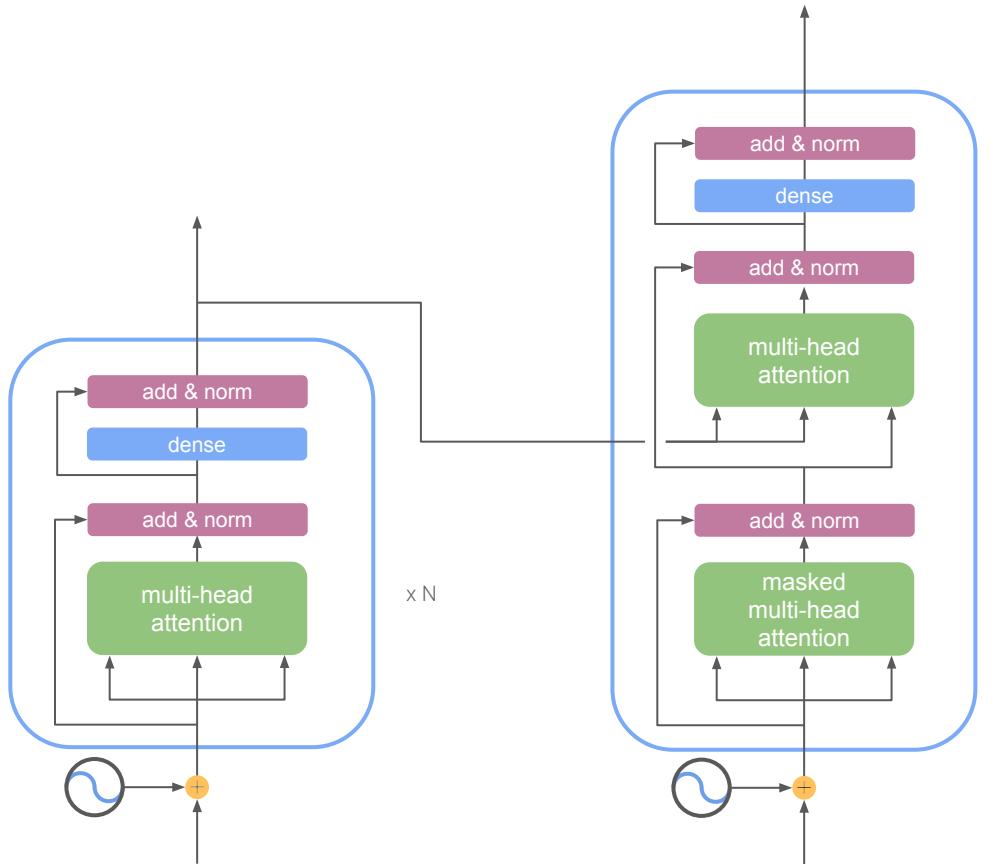
Recap

The Transformer architecture

Quadratic bottleneck in sequence length due to multi-head attention

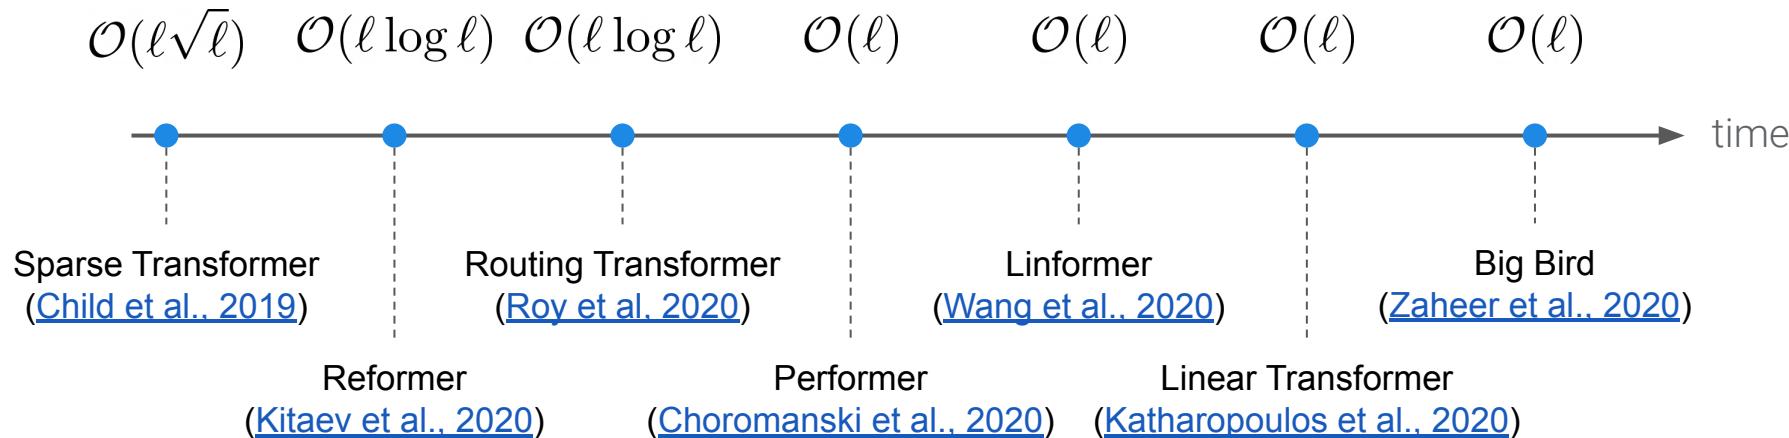
This poses a serious problem when large sequences are required, e.g.:

- Long-range dependencies
- Character-level models
- Speech processing
- High-resolution image processing



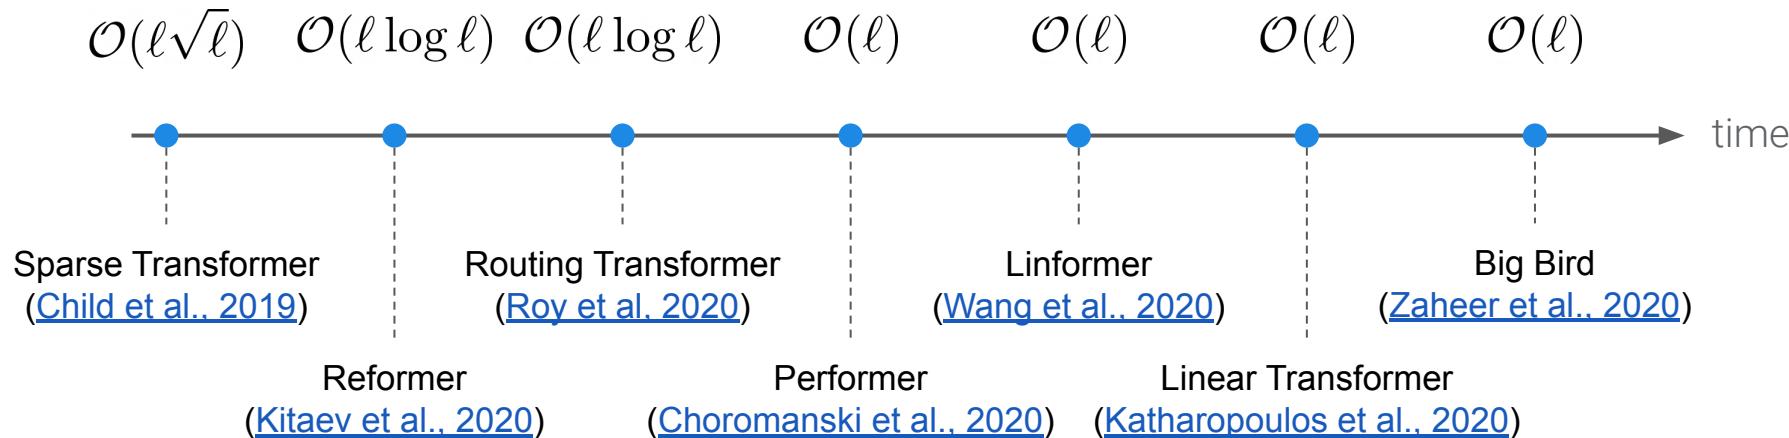
Efficient Attention

In the past months, there has been much progress in making self-attention more efficient



Efficient Attention

In the past months, there has been much progress in making self-attention more efficient

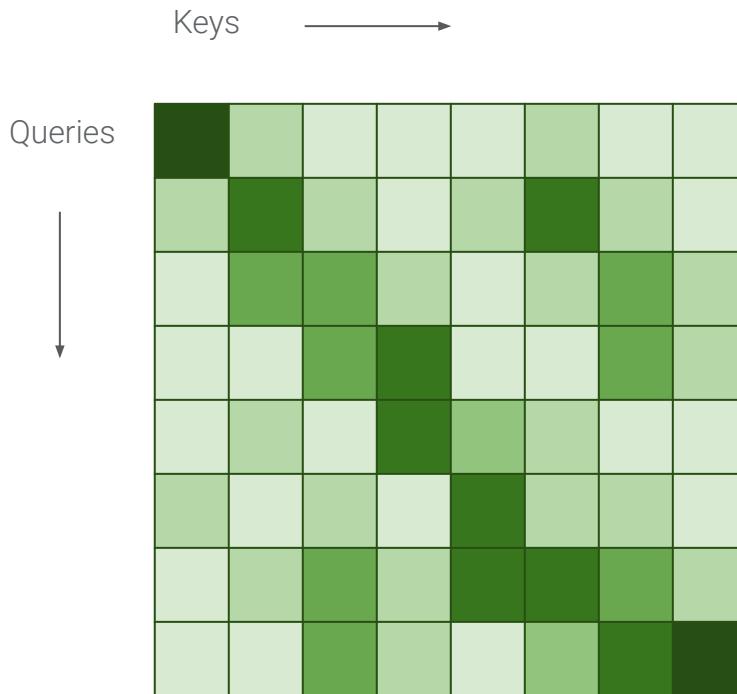


We are going to cover some ideas that make this possible

Beyond a Dense Attention Matrix

Goal:

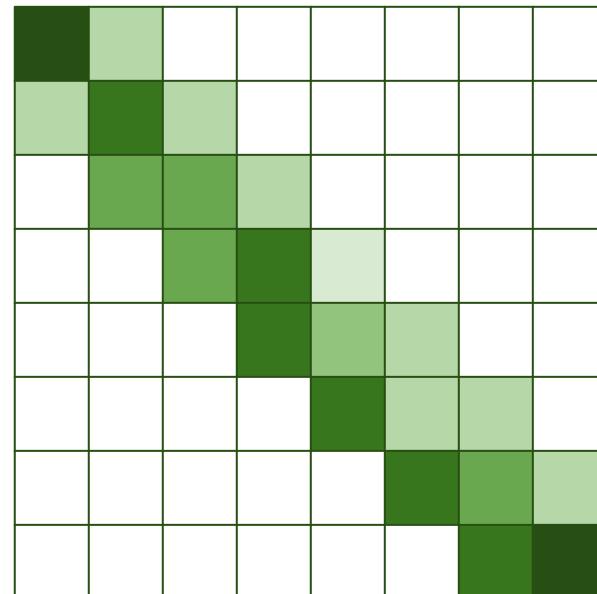
Approximate the computation of attention
via more efficient operations



Efficient Attention

A wide range of recent techniques!

- Data-Independent Patterns
 - Blockwise Transformer ([Qiu et al., 2019](#))
 - Sparse Transformer ([Child et al., 2019](#))
 - Longformer ([Beltagy et al., 2020](#))
 - Big Bird ([Zaheer et al., 2020](#))

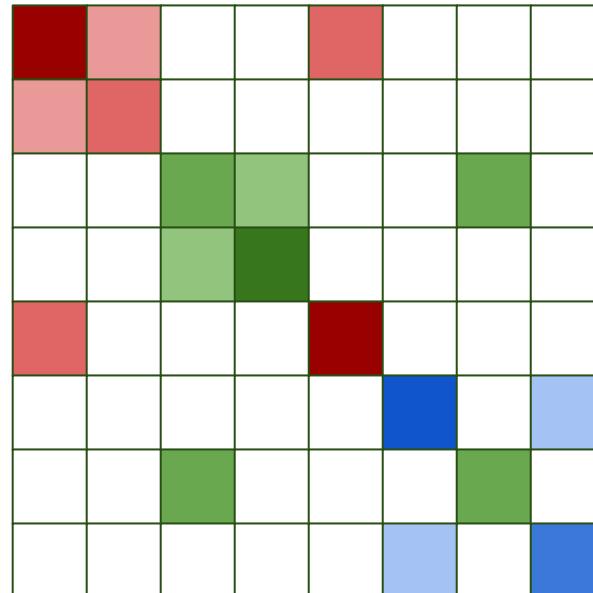


Taxonomy inspired by [Tay et al., 2020](#)

Efficient Attention

A wide range of recent techniques!

- Data-Independent Patterns
- **Data-Dependent Patterns**
 - Linformer ([Wang et al., 2020](#))
 - Reformer ([Kitaev et al., 2020](#))
 - Routing Transformer ([Roy et al., 2020](#))
 - Clustered Attention ([Vyas et al., 2020](#))
 - Sinkhorn Transformer ([Tay et al., 2020](#))

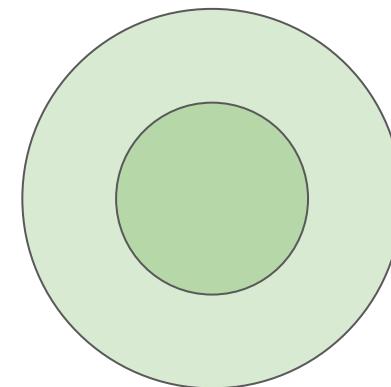


Taxonomy inspired by [Tay et al., 2020](#)

Efficient Attention

A wide range of recent techniques!

- Data-Independent Patterns
- Data-Dependent Patterns
- **Kernels and Alternative Attention Mechanisms**
 - Linear Transformer ([Katharopoulos et al., 2020](#))
 - Random Feature Attention ([Anonymous, 2020](#))
 - Performer ([Choromanski et al., 2020](#))
 - Synthesizer ([Tay et al., 2020](#))

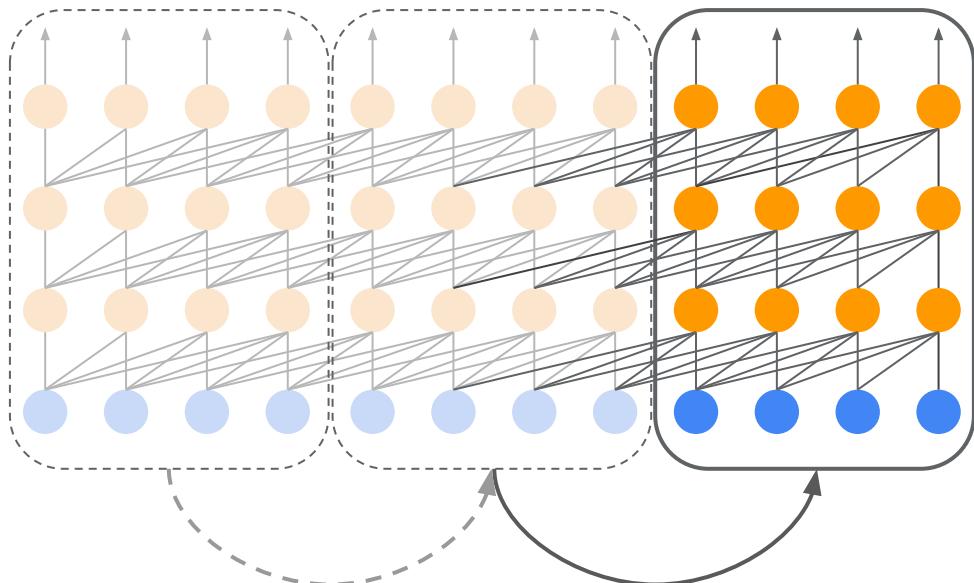


Taxonomy inspired by [Tay et al., 2020](#)

Efficient Attention

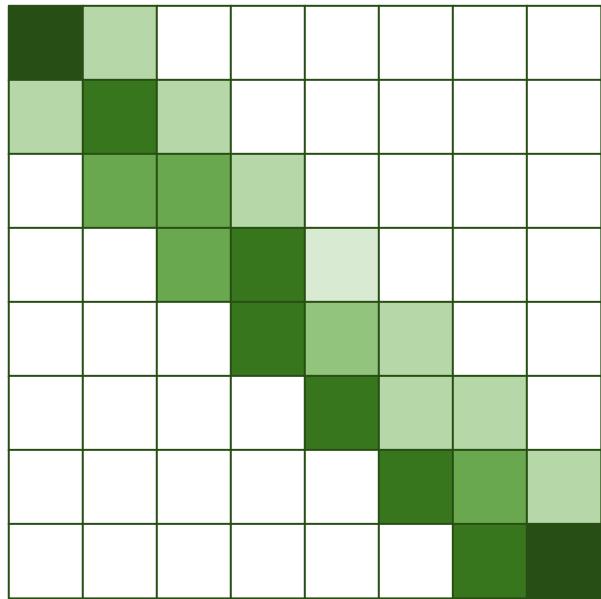
A wide range of recent techniques!

- Data-Independent Patterns
- Data-Dependent Patterns
- Alternative Attention Mechanisms
- **Recurrence**
 - Transformer XL ([Dai et al., 2019](#))
 - Compressive Transformers ([Rae et al., 2019](#))



Taxonomy inspired by [Tay et al., 2020](#)

Data-Independent Patterns



Data-Independent Patterns

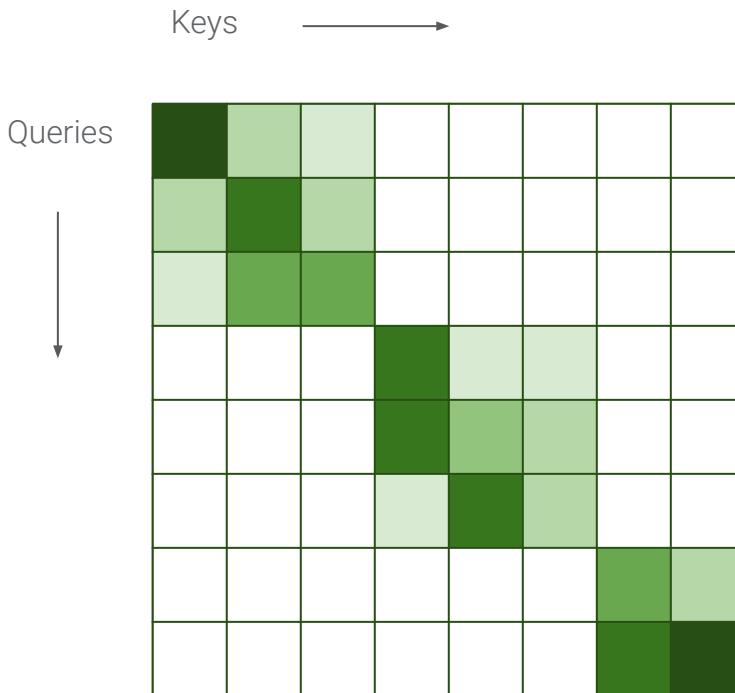
Blockwise Patterns

Divide sequence into local blocks and restrict attention within them

Examples:

Blockwise Transformer ([Qiu et al., 2019](#))

Local Attention ([Parmar et al., 2018](#))



Data-Independent Patterns

Strided Patterns

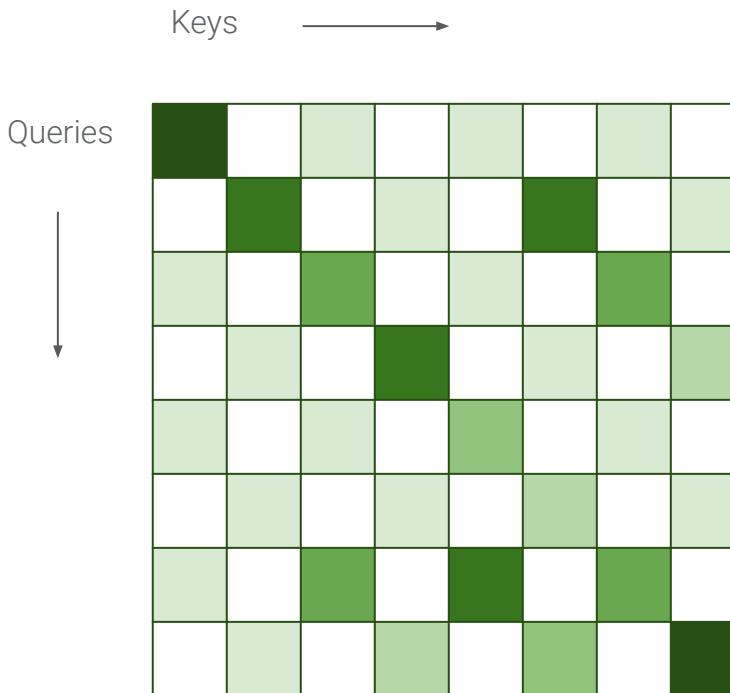
Skip some query/key pairs.

Quadratic in sequence length / stride

Examples:

Sparse Transformer ([Child et al., 2019](#))

Longformer ([Beltagy et al, 2020](#))



Data-Independent Patterns

Diagonal Patterns

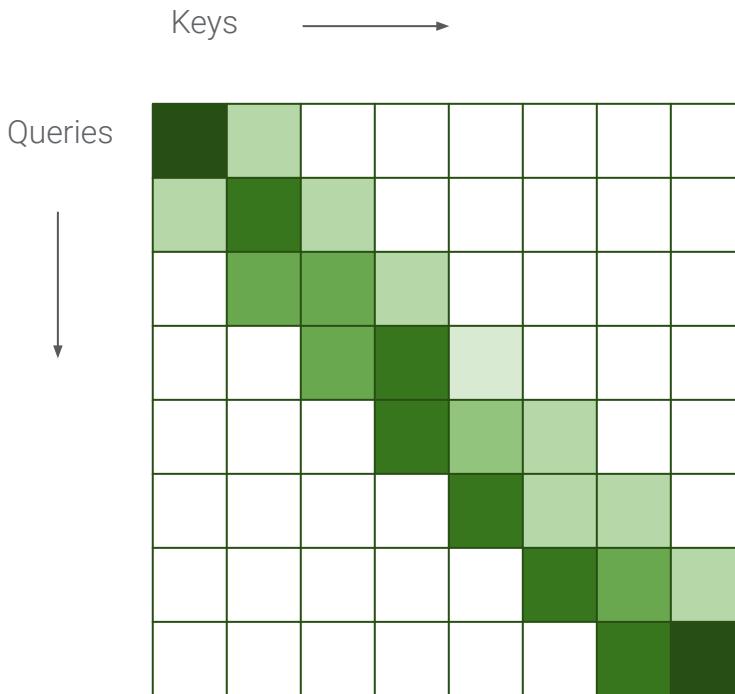
Compute attention over the diagonal.

Linear in sequence length and window size.

Examples:

Longformer ([Beltagy et al, 2020](#))

Big Bird ([Zaheer et al., 2020](#))



Data-Independent Patterns

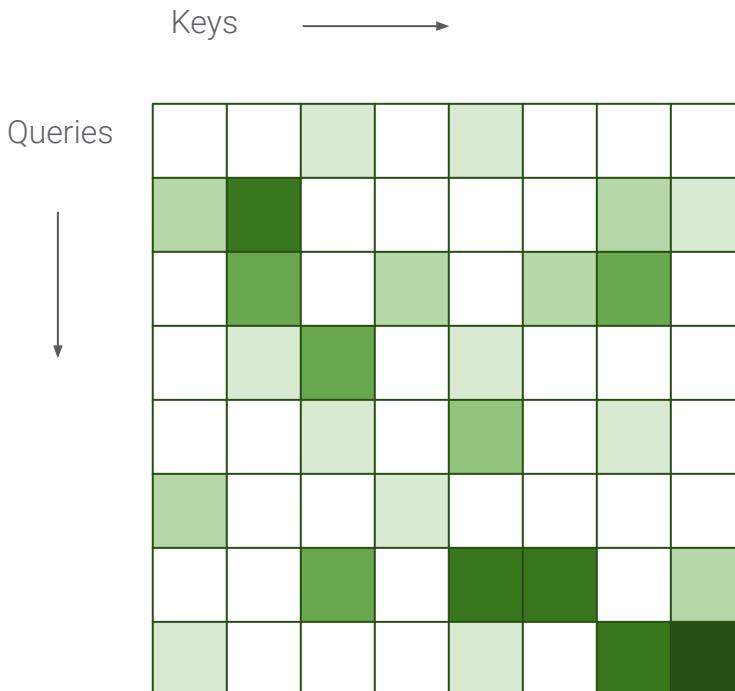
Random Patterns

Compute attention over random query/key pairs.

Linear in number of points.

Examples:

Big Bird ([Zaheer et al., 2020](#))



Data-Independent Patterns

Global Attention

Applied to one or a few special tokens, often prepended to the sequence.

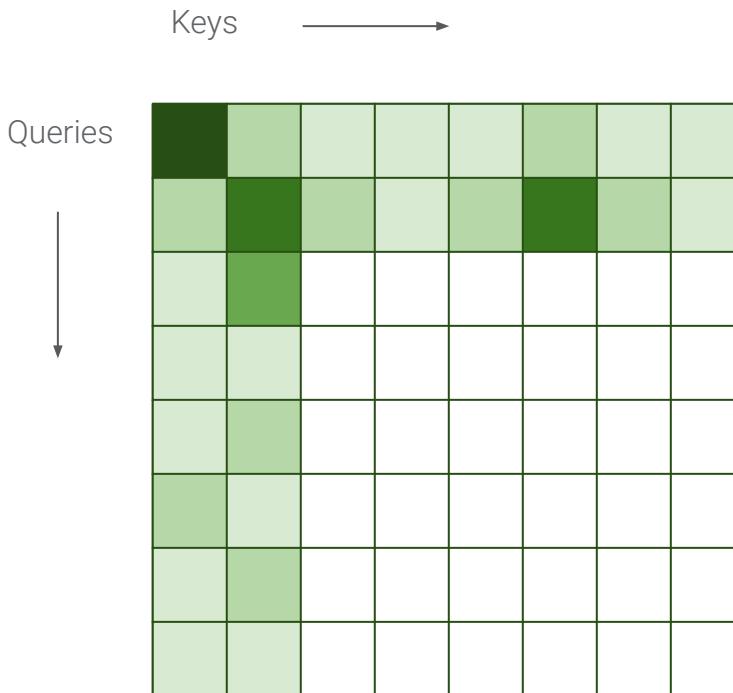
Usually combined with other patterns

Examples:

Big Bird ([Zaheer et al., 2020](#))

Longformer ([Beltagy et al., 2020](#))

ETC ([Ainslie et al., 2020](#))



Data-Independent Patterns

Combination of Patterns

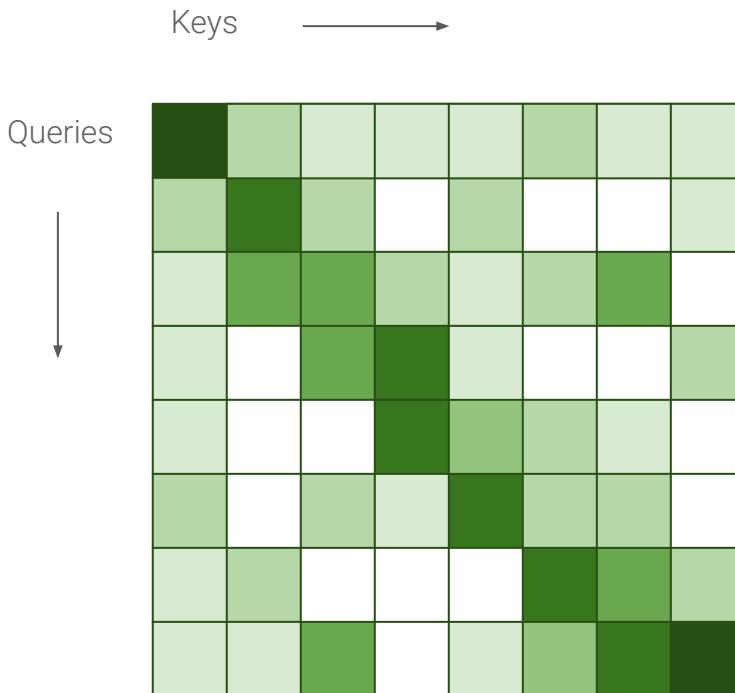
Combine multiple patterns

(e.g. Global + Diagonal + Random)

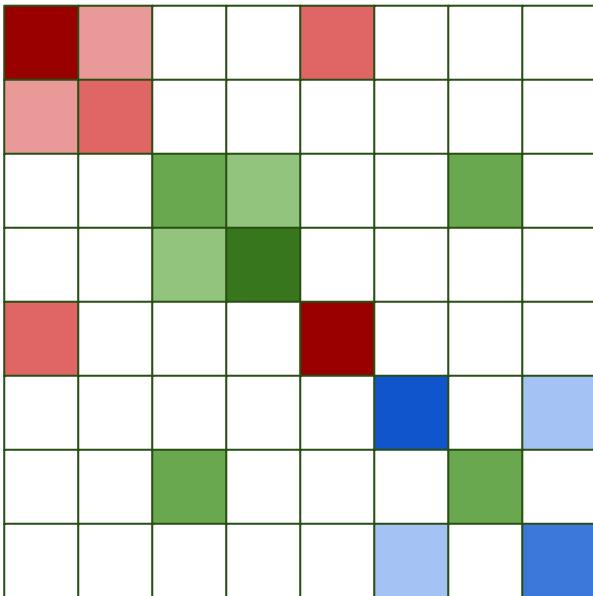
Examples:

Big Bird ([Zaheer et al., 2020](#))

Longformer ([Beltagy et al., 2020](#))



Data-Dependent Patterns



Data-Dependent Patterns

Buckets

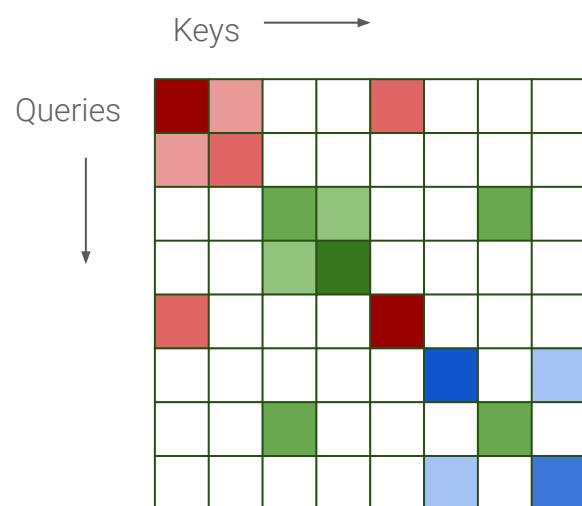
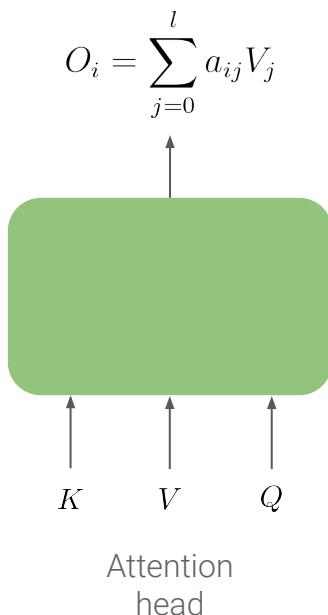
Create buckets/clusters and compute attention within.

Ideally, buckets should contain the highests attention weights in the matrix

Examples:

Reformer ([Kitaev et al., 2020](#))

Routing Transformer ([Roy et al., 2020](#))



Data-Dependent Patterns

Buckets: Hashing

Locality-Sensitive Hashing (LSH)

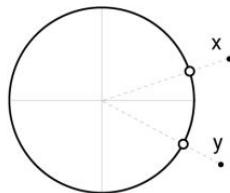
Key idea: take a random projection matrix R , compute hash for a vector x through:

$$h(x) = \arg \max([xR; -xR])$$

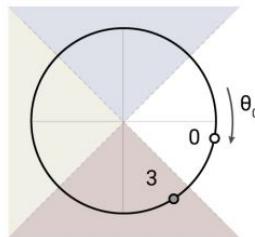
Examples:

Reformer ([Kitaev et al., 2020](#))

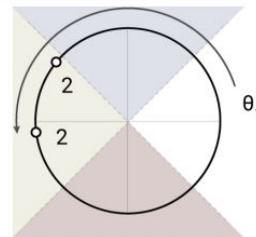
Sphere Projected Points



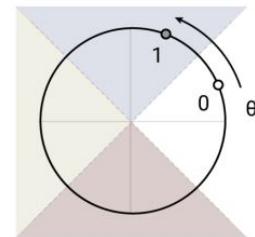
Random Rotation 0



Random Rotation 1

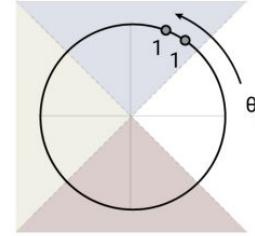
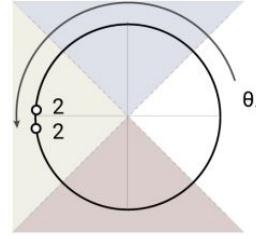
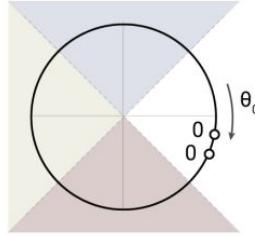
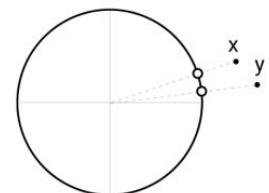


Random Rotation 2



x: 0 2 1

y: 3 2 0



x: 0 2 1

y: 0 2 1

Data-Dependent Patterns

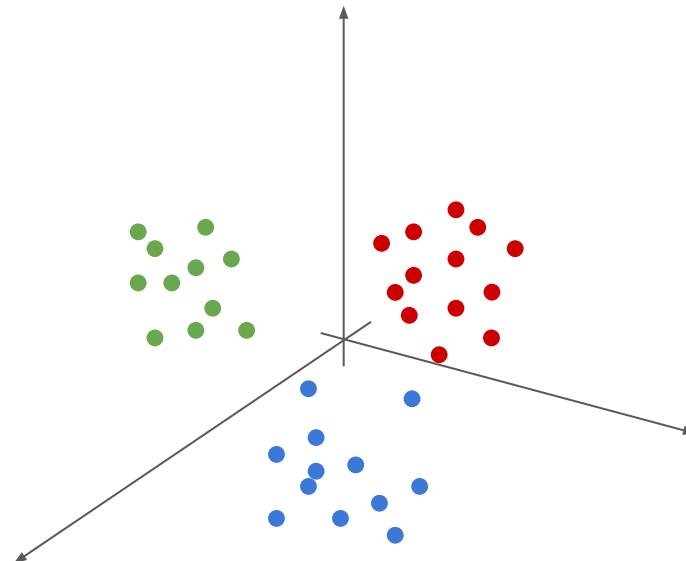
Buckets: Clustering

E.g. online k-means

Examples:

Routing Transformer ([Roy et al., 2020](#))

Clustered Attention ([Vyas et al., 2020](#))



Data-Dependent Patterns

Sorting and blocking

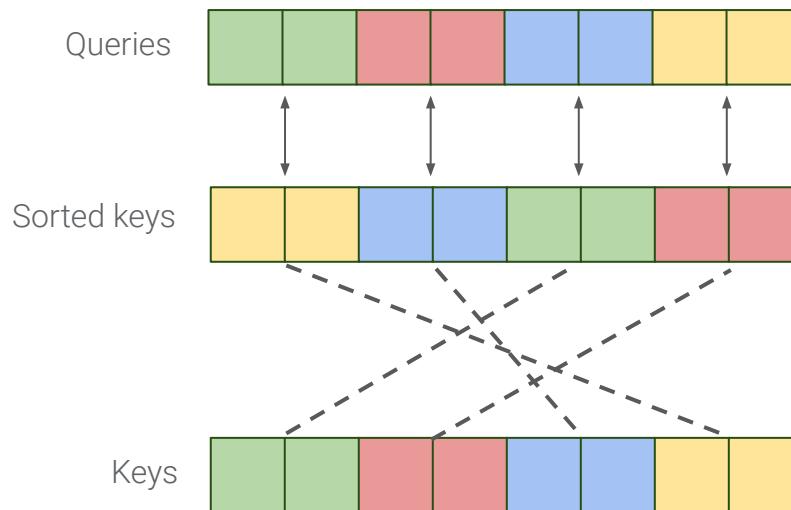
E.g. Sparse Sinkhorn Attention

Key ideas:

- A differentiable sorting network that learns to rearrange blocked inputs, using the Sinkhorn balancing mechanism to create a permutation matrix
- Attention is computed only on local neighborhoods (before and after sorting)

Examples:

Sinkhorn Transformer ([Tay et al., 2020](#))



Data-Dependent Patterns

Compression

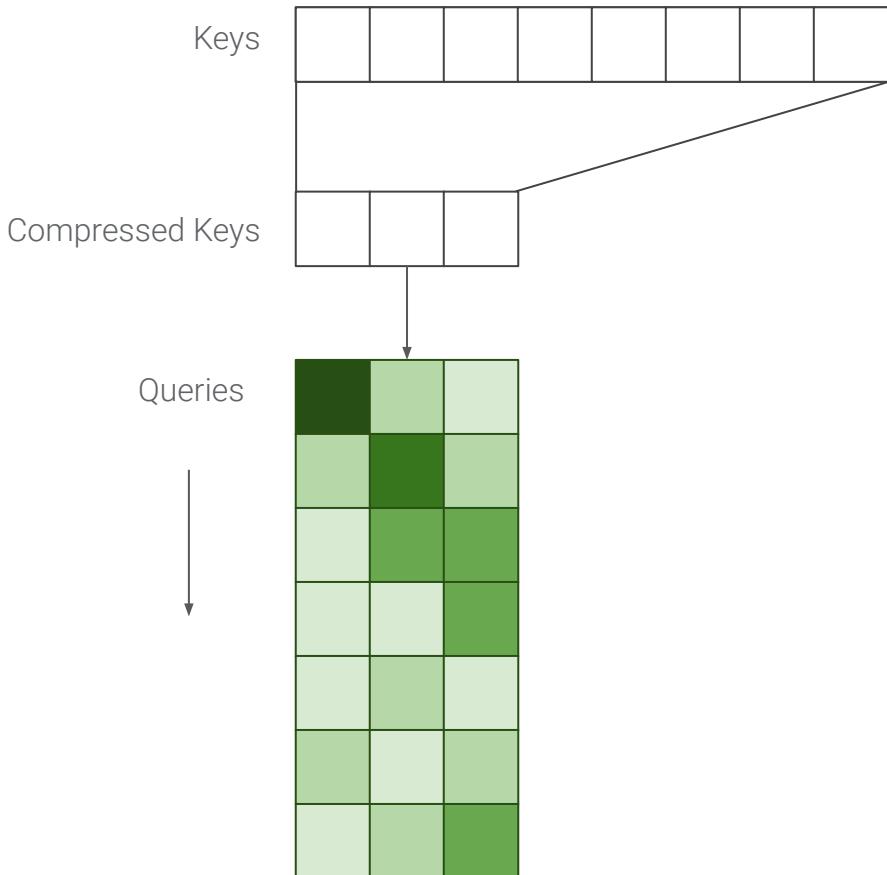
E.g. pooling, strided convolution, low-rank projections with learnable weights

Examples:

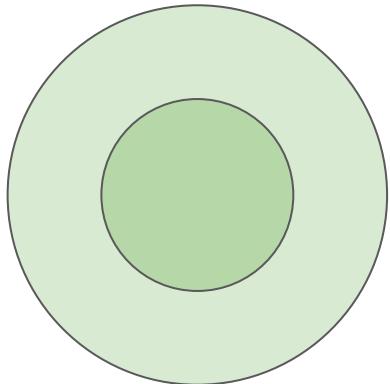
Compressed Attention ([Liu et al., 2018](#))

Linformer ([Wang et al., 2020](#))

Synthesizers ([Tay et al., 2020](#))



Kernels and Alternative Attention Mechanisms



Kernels and Alternative Attention Mechanisms

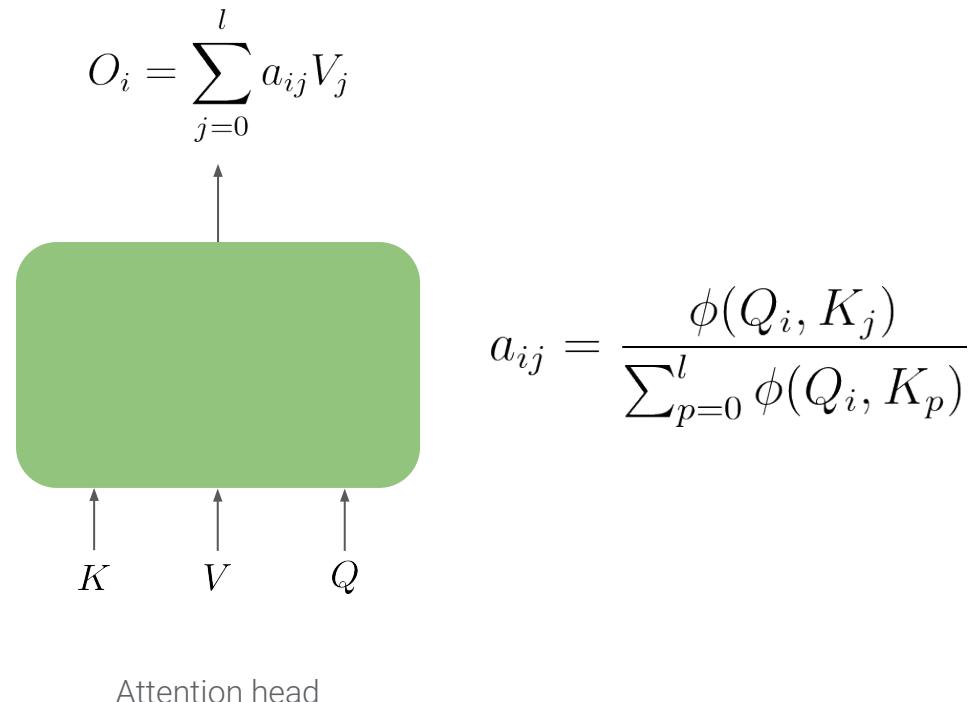
Kernels

Recap: attention in its general form uses a similarity function

$$\phi(Q_i, K_j)$$

Standard transformers use dot product attention:

$$\phi(Q_i, K_j) = \exp\left(\frac{Q_i K_j^\top}{\sqrt{d}}\right)$$



Kernels and Alternative Attention Mechanisms

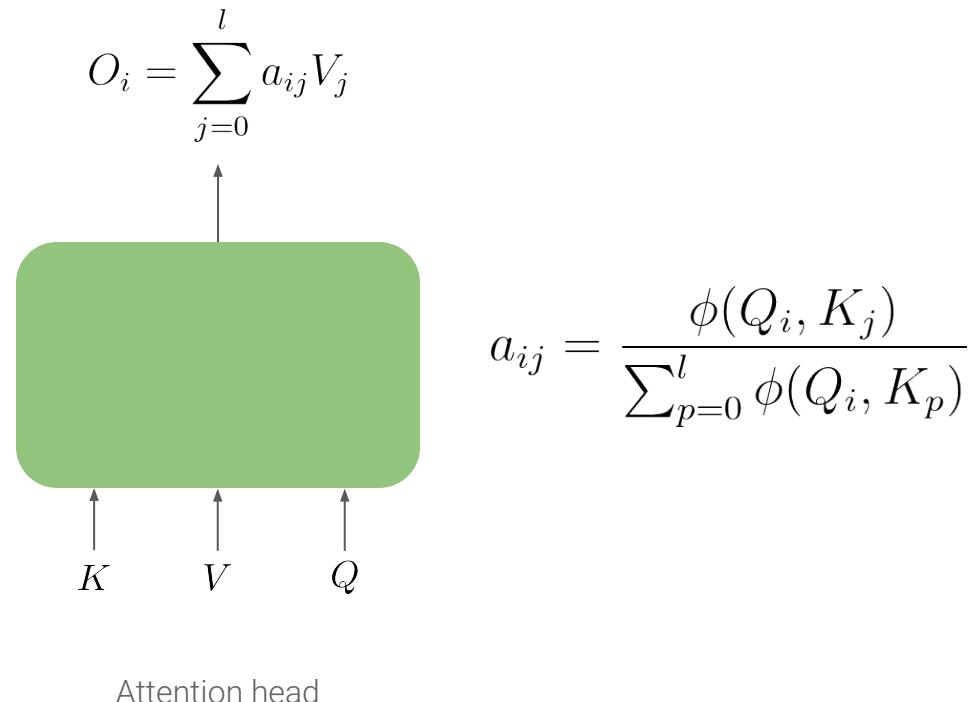
Kernels

Recap: attention in its general form uses a similarity function

$$\phi(Q_i, K_j)$$

However, we can simplify things with a decomposable **kernel**:

$$\phi(Q_i, K_j) = \phi(Q_i)^\top \phi(K_j)$$



Kernels and Alternative Attention Mechanisms

Kernels

Recap: attention in its general form uses a similarity function

$$\phi(Q_i, K_j)$$

However, we can simplify things with a decomposable **kernel**:

$$O_i = \sum_{j=0}^l a_{ij} V_j \quad a_{ij} = \frac{\phi(Q_i, K_j)}{\sum_{p=0}^l \phi(Q_i, K_p)}$$

$$\phi(Q_i, K_j) = \phi(Q_i)^\top \phi(K_j)$$

Kernels and Alternative Attention Mechanisms

Kernels

Recap: attention in its general form uses a similarity function

$$\phi(Q_i, K_j)$$

However, we can simplify things with a decomposable **kernel**:

$$O_i = \sum_{j=0}^l a_{ij} V_j \quad a_{ij} = \frac{\phi(Q_i, K_j)}{\sum_{p=0}^l \phi(Q_i, K_p)}$$

$$O_i = \frac{\sum_{j=0}^l \phi(Q_i)^\top \phi(K_j) V_j}{\sum_{j=0}^l \phi(Q_i)^\top \phi(K_j)}$$

$$\phi(Q_i, K_j) = \phi(Q_i)^\top \phi(K_j)$$

$$O_i = \frac{\phi(Q_i)^\top \sum_{j=0}^l \phi(K_j) V_j}{\phi(Q_i)^\top \sum_{j=0}^l \phi(K_j)}$$

Kernels and Alternative Attention Mechanisms

Kernels

Recap: attention in its general form uses a similarity function

$$\phi(Q_i, K_j)$$

However, we can simplify things with a decomposable **kernel**:

$$O_i = \sum_{j=0}^l a_{ij} V_j \quad a_{ij} = \frac{\phi(Q_i, K_j)}{\sum_{p=0}^l \phi(Q_i, K_p)}$$

$$O_i = \frac{\sum_{j=0}^l \phi(Q_i)^\top \phi(K_j) V_j}{\sum_{j=0}^l \phi(Q_i)^\top \phi(K_j)}$$

$$\phi(Q_i, K_j) = \phi(Q_i)^\top \phi(K_j)$$

$$O_i = \frac{\phi(Q_i)^\top \boxed{\sum_{j=0}^l \phi(K_j) V_j}}{\phi(Q_i)^\top \boxed{\sum_{j=0}^l \phi(K_j)}}$$

Independent of query!

Kernels and Alternative Attention Mechanisms

Kernels

$$O_i = \frac{\phi(Q_i)^\top \sum_{j=0}^l \phi(K_j) V_j}{\phi(Q_i)^\top \sum_{j=0}^l \phi(K_j)}$$



Independent of query!

This allows us to compute attention in **linear** time!

In [Katharopoulos et al., 2020](#):

$$\phi(x) = \text{elu}(x) + 1 = \max(\alpha(e^x - 1), 0) + 1$$

Kernels and Alternative Attention Mechanisms

Kernels

Random Feature Attention ([Anonymous, 2020](#))

$$O_i = \frac{\phi(Q_i)^\top \sum_{j=0}^l \phi(K_j) V_j}{\phi(Q_i)^\top \sum_{j=0}^l \phi(K_j)}$$


Independent of query!

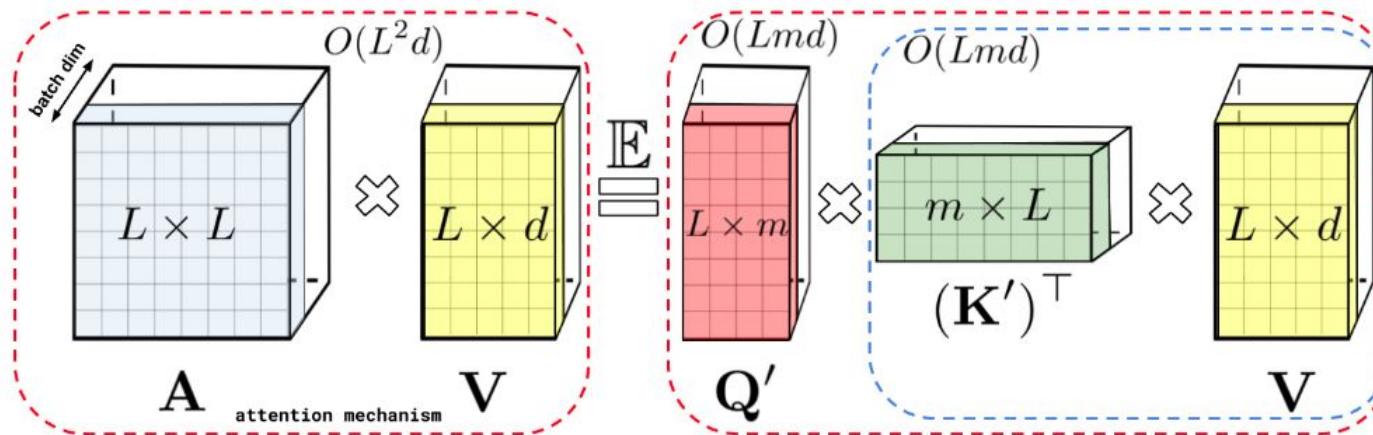
Random features can be used to generate an unbiased estimation of the standard softmax function

Kernels and Alternative Attention Mechanisms

Performer: Generalized Attention and FAVOR ([Choromanski et al., 2020](#))

Rethink attention as $a_{ij} = g(Q_i^\top)K(Q_i^\top, K_j^\top)h(K_j^\top)$, parametrized by a kernel K and functions g and h

This work presents an unbiased, low-variance approximation of attention via random feature map decompositions, with linear time and space complexity.



Kernels and Alternative Attention Mechanisms

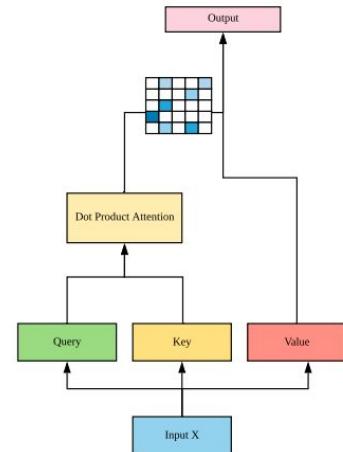
Synthesizers ([Tay et al., 2020](#))

Are token-to-token interactions really necessary?

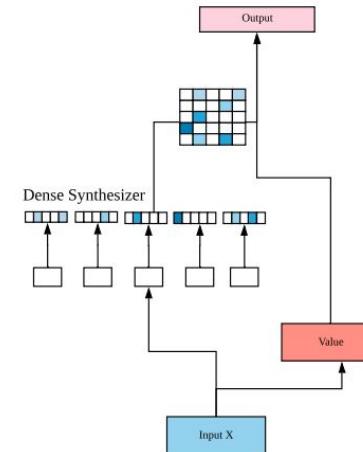
Random attention matrices are surprisingly competitively!

Low-rank alternatives can be used

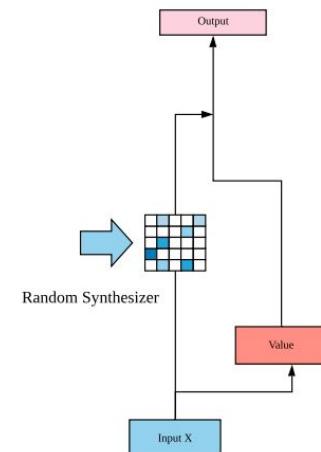
(a) Transformer



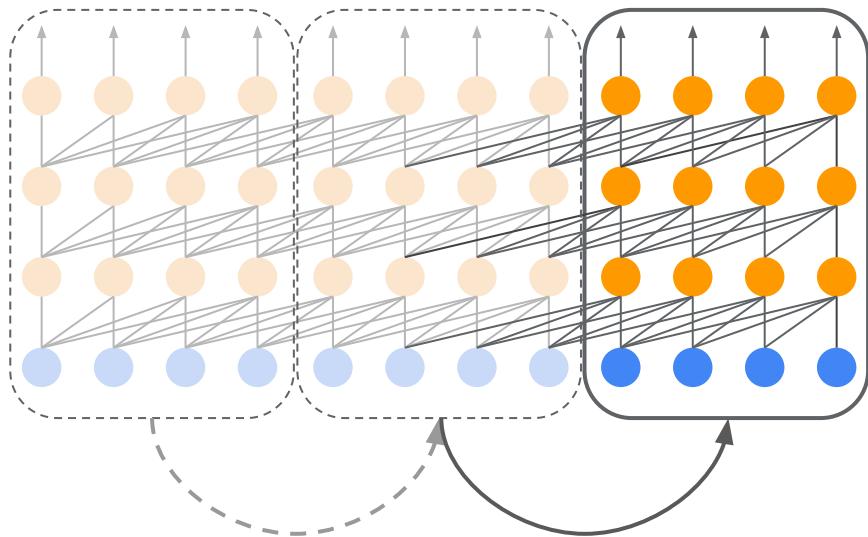
(b) Synthesizer (Dense)



(c) Synthesizer (Random)



Recurrence

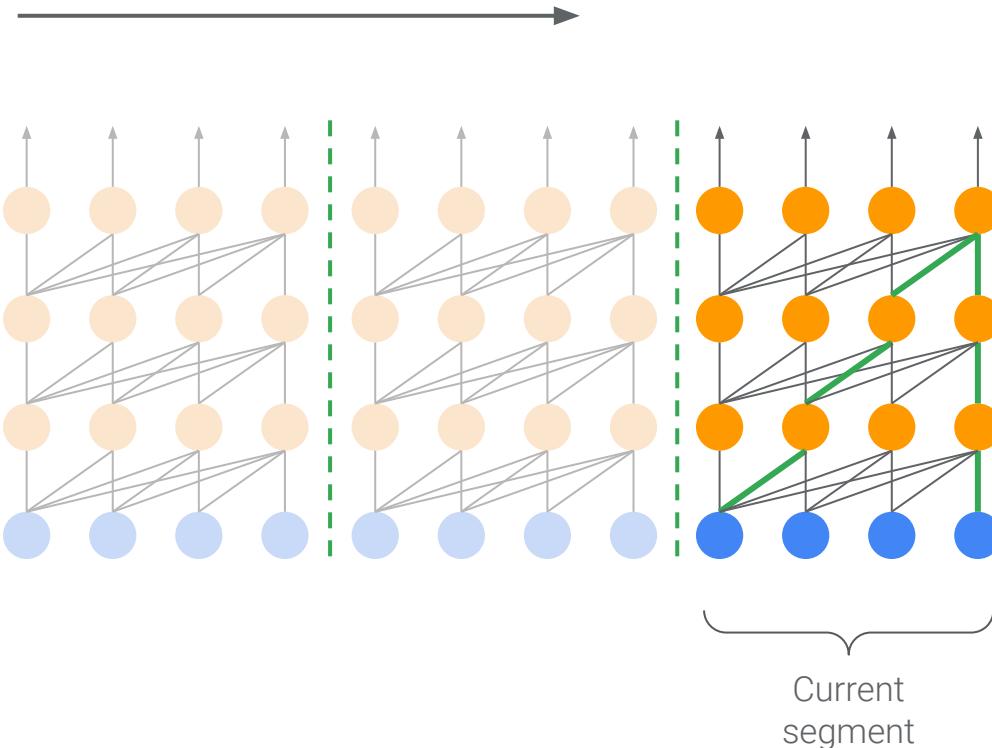


Recurrence

Transformer-XL ([Dai et al., 2019](#))

How can models process long sequences under limited hardware constraints?

A naive approach is to split the sequence into multiple smaller ones and process them separately



Current
segment

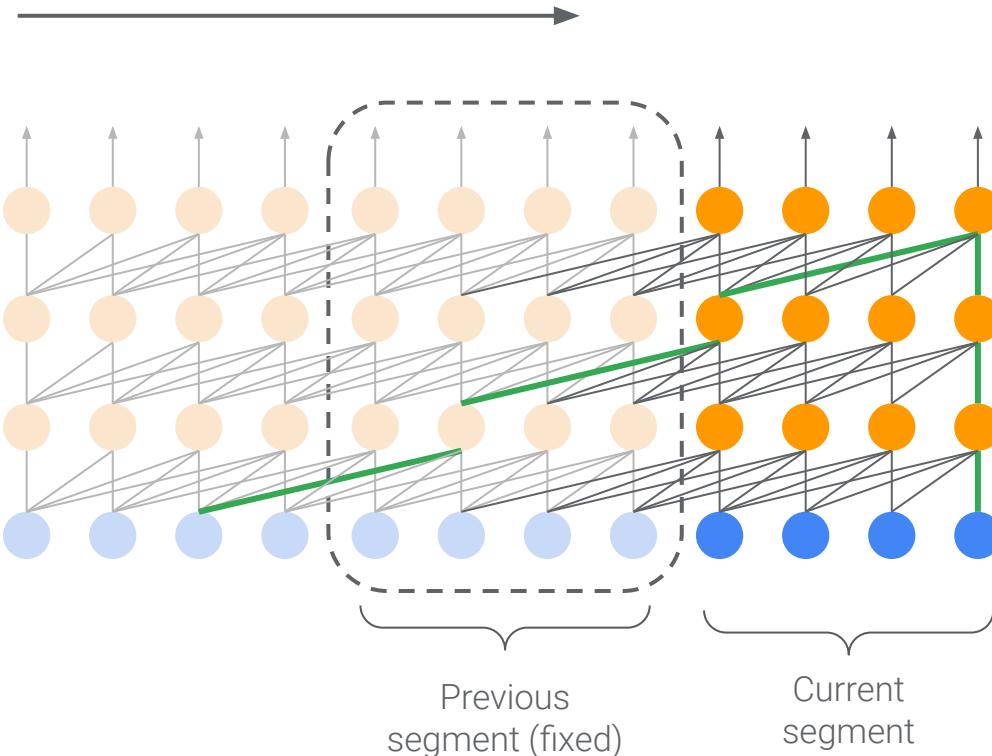
Recurrence

Transformer-XL ([Dai et al., 2019](#))

A better way is to add a segment-level recurrence mechanism

Representations from the previous segment are cached and re-used (no gradients flowing at training)

This increases **receptive field** proportionally to the depth of the transformer



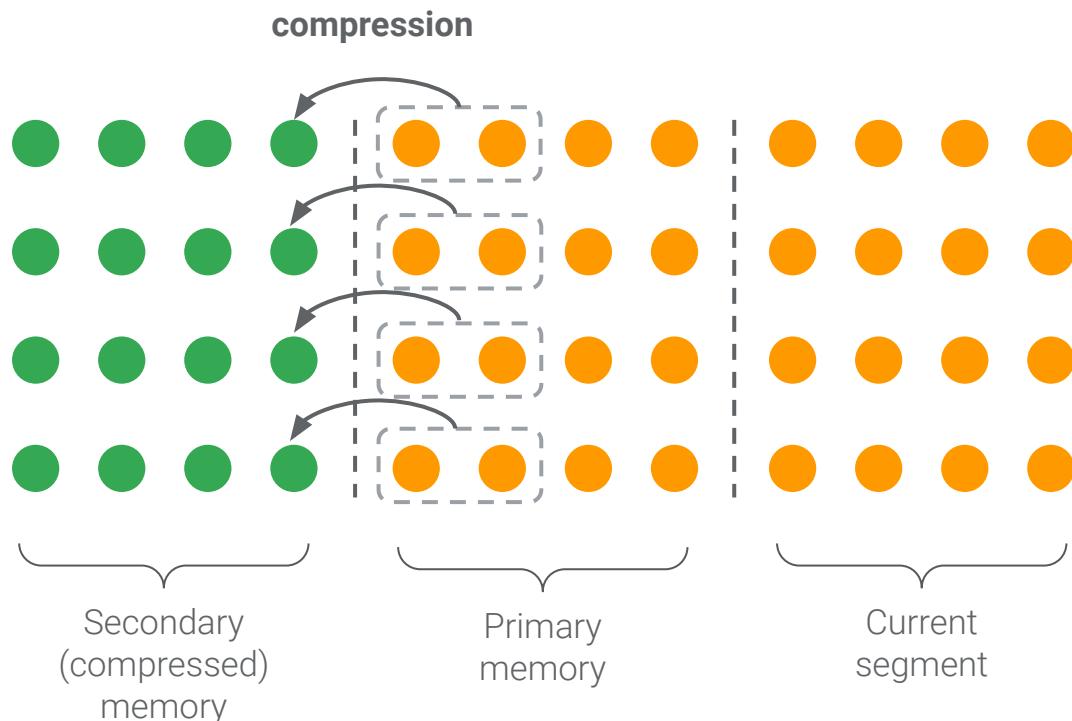
Recurrence

Compressive Transformers

(Rae et al., 2019)

Dual memory system:

- Primary mem. contains activations from previous segment
- Secondary mem. is compresses activations from all previous segments



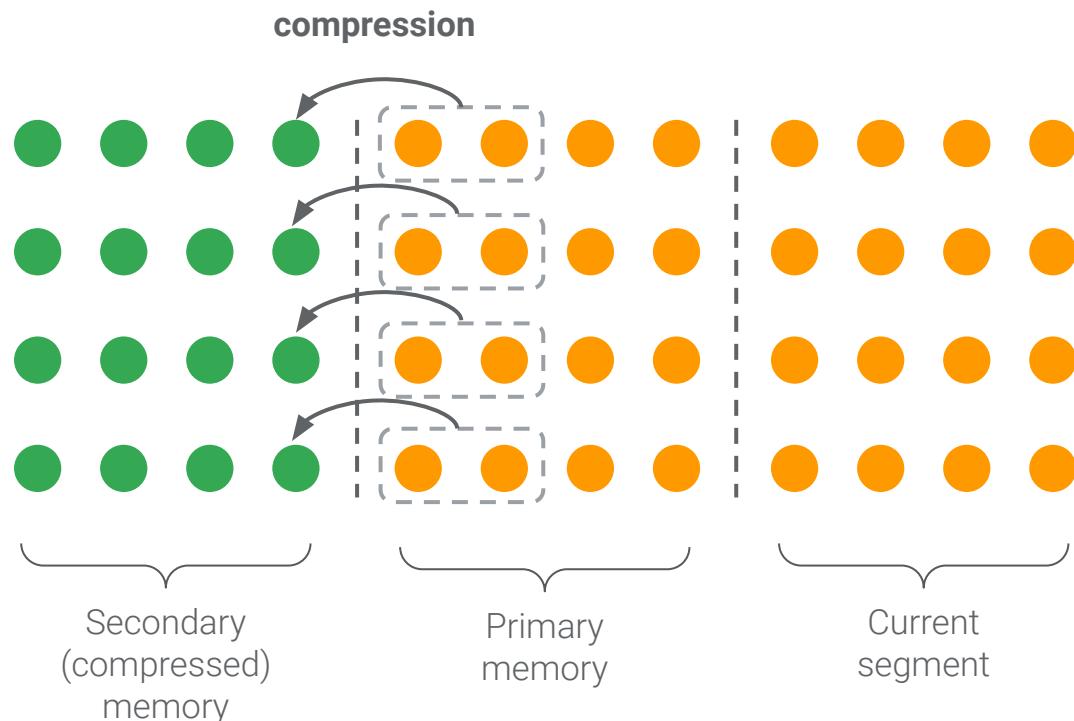
Recurrence

Compressive Transformers

(Rae et al., 2019)

When a new segment comes:

- Primary memory is updated with activations from previous segment
- Secondary memory is updated with the activations from the primary memory, where a **compression function** is applied (e.g. pooling, convolutions, most used)



Overview



Benchmarking

How do these models compare in practice?

The [Long-Range Arena](#): a benchmark for
efficient transformers

Benchmarking

How do these models compare in practice?

The [Long-Range Arena](#): a benchmark for efficient transformers

Longer sequences: 1K-16K

5 tasks:

- List operations (e.g. max, min, median)
- Byte-level text classification
- Byte-level document retrieval
- Image classification
- Long-range spatial dependency

List operations example:

INPUT: [MAX 4 3 [MIN 2 3] 1 0 [MEDIAN 1 5 8 9, 2]] **OUTPUT:** 5

Benchmarking

How do these models compare in practice?

The [Long-Range Arena](#): a benchmark for efficient transformers

Longer sequences: 1K-16K

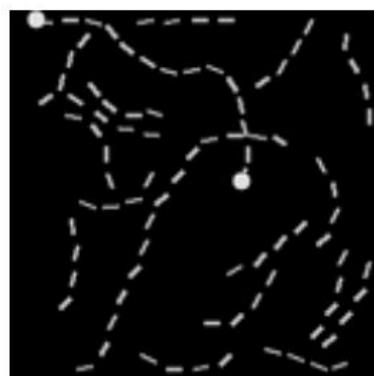
5 tasks:

- List operations (e.g. max, min, median)
- Byte-level text classification
- Byte-level document retrieval
- Image classification
- Long-range spatial dependency

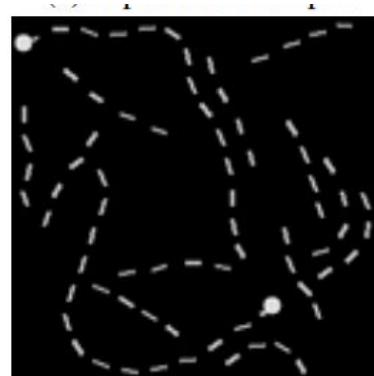
List operations example:

INPUT: [MAX 4 3 [MIN 2 3] 1 0 [MEDIAN 1 5 8 9, 2]] **OUTPUT:** 5

Long-range spatial dependency example:



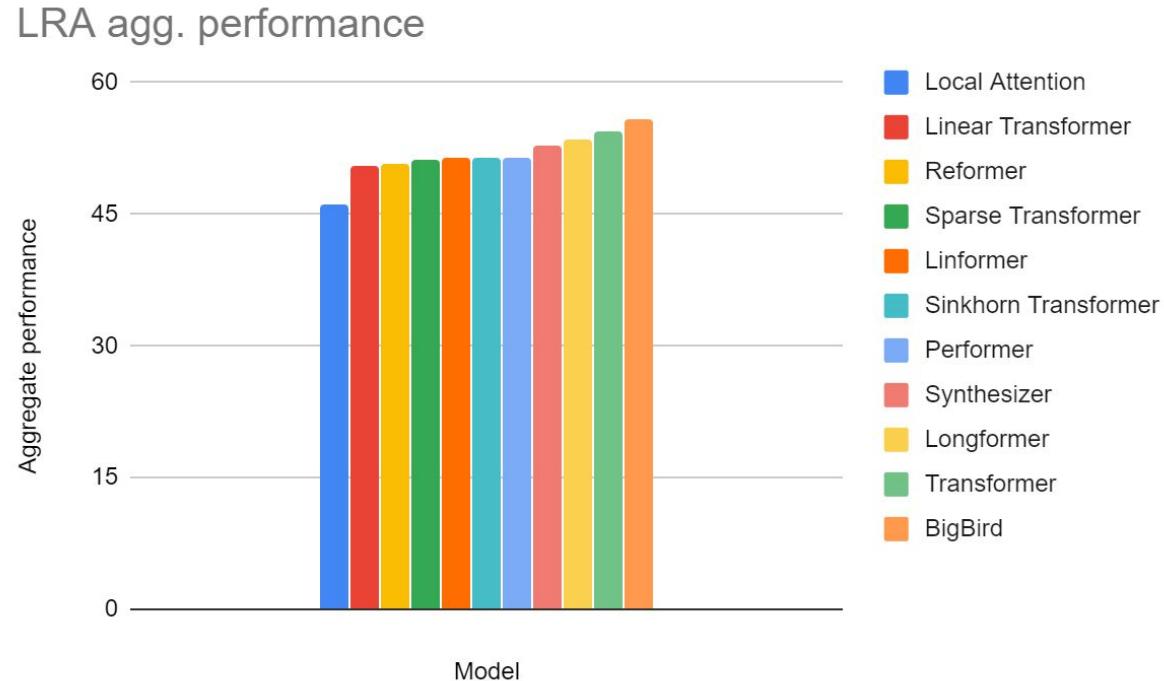
(positive)



(negative)

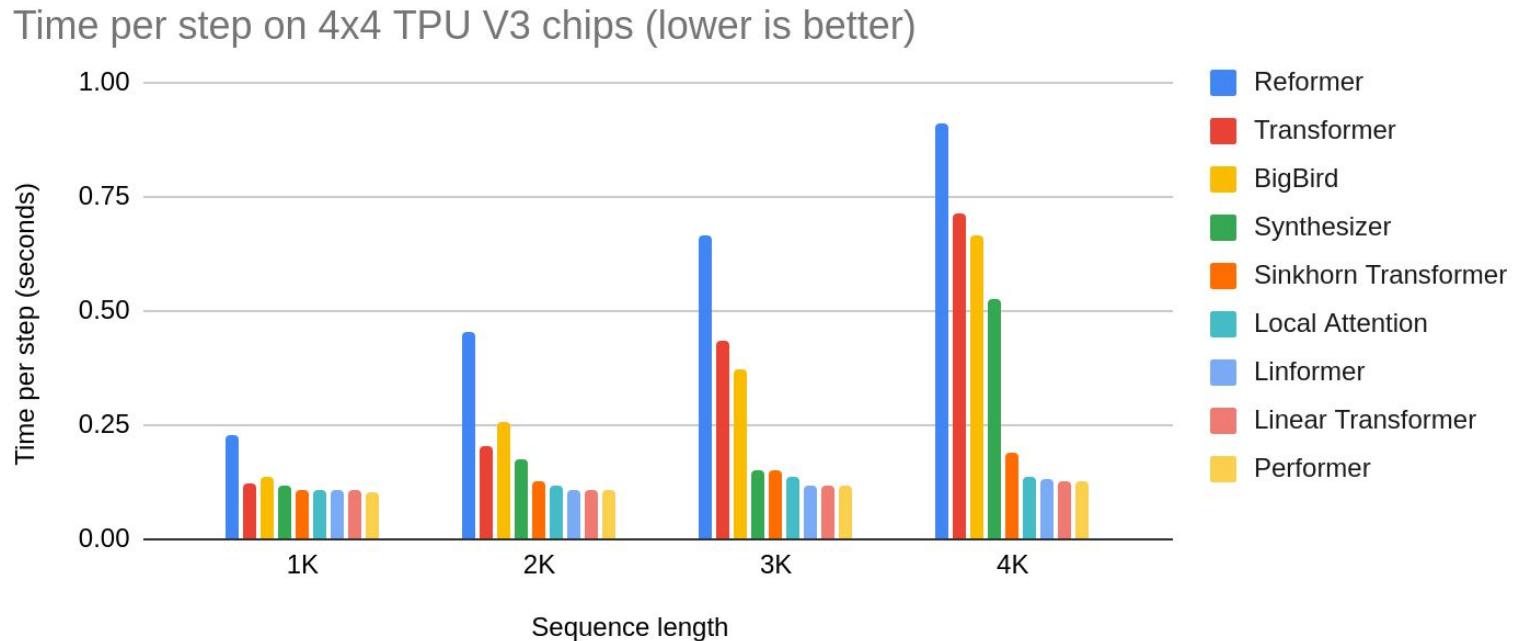
Benchmarking

The [Long-Range Arena](#): a benchmark for efficient transformers



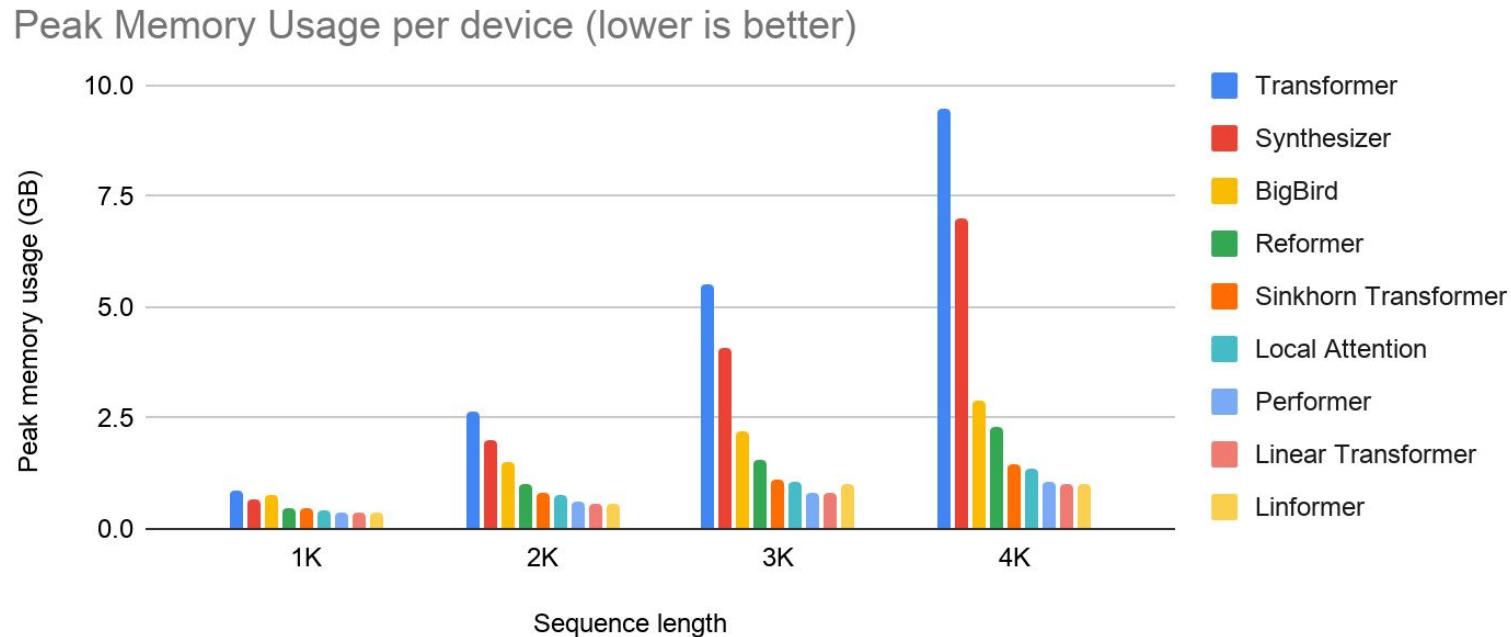
Benchmarking

The [Long-Range Arena](#): a benchmark for efficient transformers



Benchmarking

The [Long-Range Arena](#): a benchmark for efficient transformers

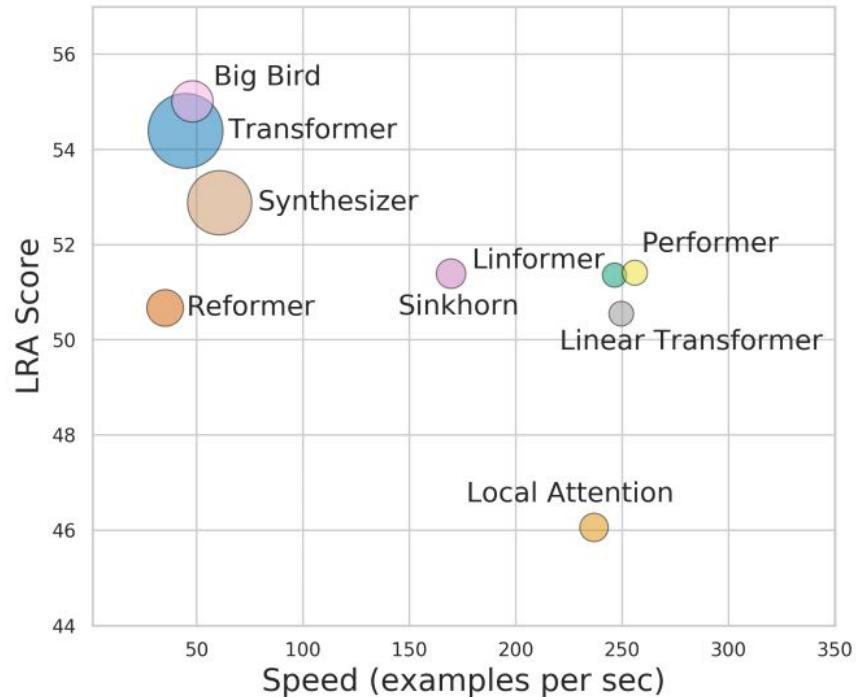


Benchmarking

The [Long-Range Arena](#): a benchmark for efficient transformers

Putting it all together (size of circles corresponds to memory footprint)

Note: these results might be sensitive to implementation details, hardware and hyper-parameters.



Key Takeaways

There has been a surge in ideas for improving the efficiency of attention and transformers, especially for improving their capacity to handle long sequences.

Key Takeaways

There has been a surge in ideas for improving the efficiency of attention and transformers, especially for improving their capacity to handle long sequences.

There has been good progress in recent months: we are now able to compute attention in linear time with respect to sequence length, leading to large speed improvements without much performance drops for large sequences.

Key Takeaways

There has been a surge in ideas for improving the efficiency of attention and transformers, especially for improving their capacity to handle long sequences.

There has been good progress in recent months: we are now able to compute attention in linear time with respect to sequence length, leading to large speed improvements without much performance drops for large sequences.

Future improvements in hardware, e.g. on the efficiency of sparse computations, may make these ideas even more appealing in the long run ([Hooker, 2020](#))

Key Takeaways

There has been a surge in ideas for improving the efficiency of attention and transformers, especially for improving their capacity to handle long sequences.

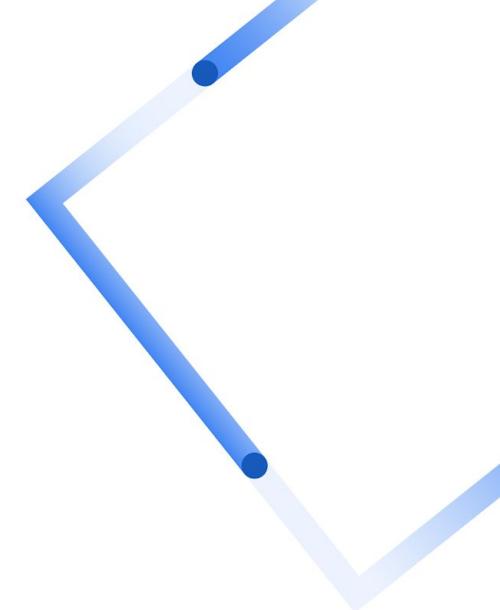
There has been good progress in recent months: we are now able to compute attention in linear time with respect to sequence length, leading to large speed improvements without much performance drops for large sequences.

Future improvements in hardware, e.g. on the efficiency of sparse computations, may make these ideas even more appealing in the long run ([Hooker, 2020](#))

The ideas presented in this section are often orthogonal to each other and to other efforts presented in this tutorial, and can be combined for more efficient models.

05

Case Studies



Efficient Language models



Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks



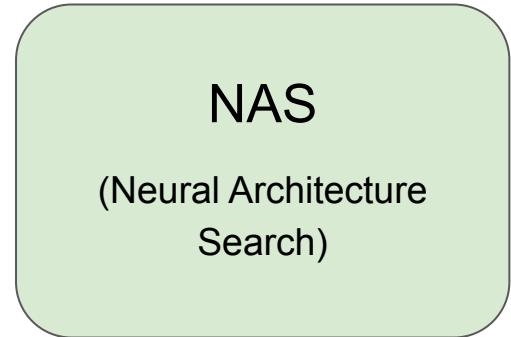
Source: unsplash.com

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer

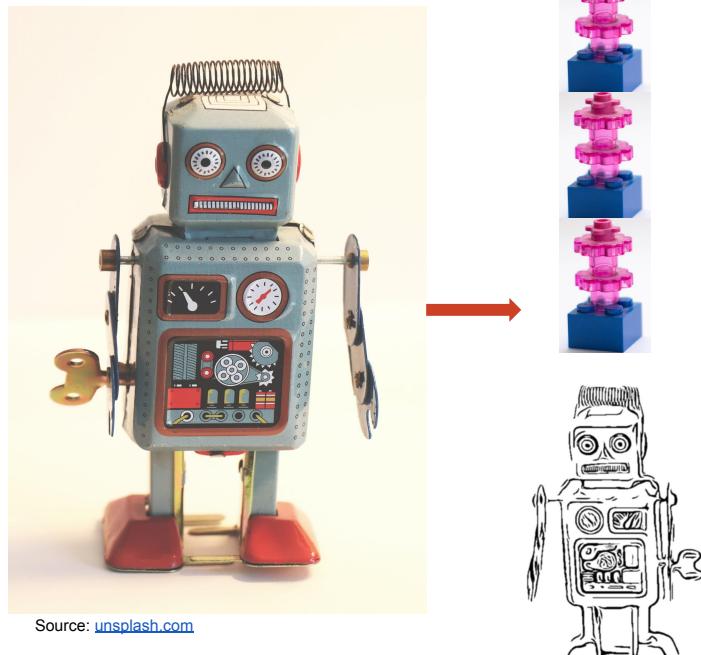


Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pORNN



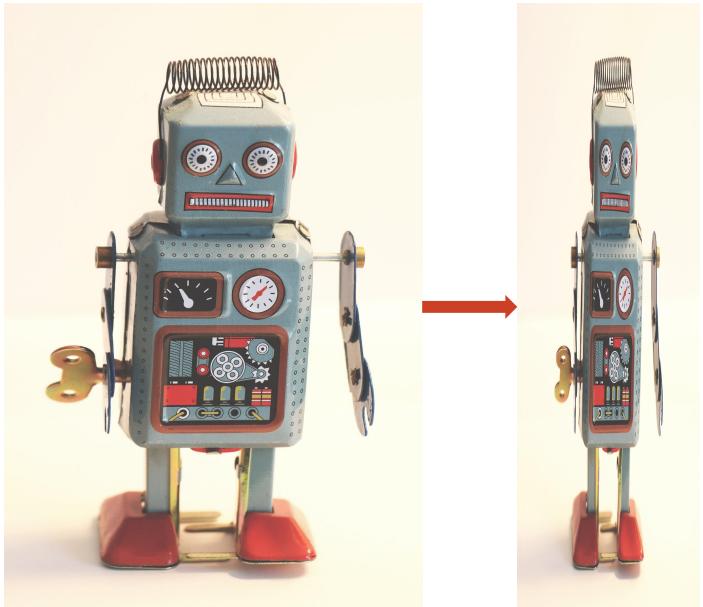
Source: unsplash.com

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pQRNN
- MobileBERT



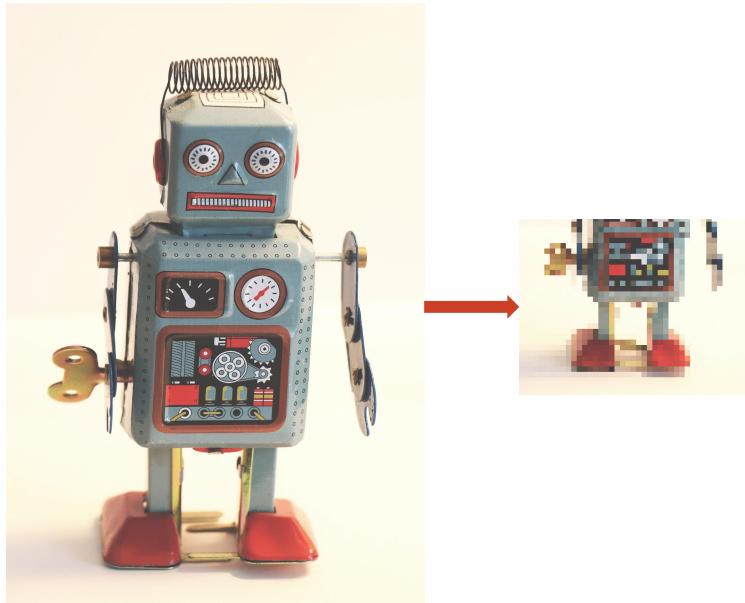
Source: unsplash.com

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pQRNN
- MobileBERT
- Lite Transformer with Long-Short Range Attention



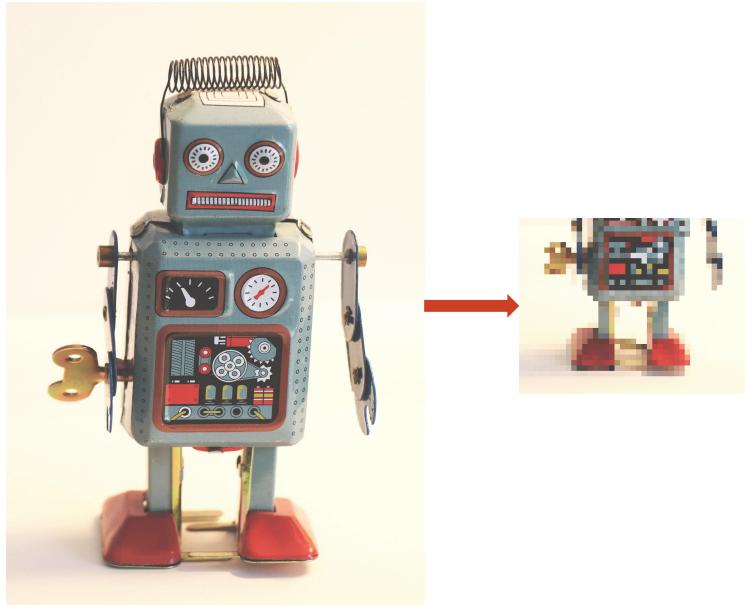
Source: unsplash.com

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pQRNN
- MobileBERT
- Lite Transformer with Long-Short Range Attention
- MicroNet for Efficient Language Modeling



Source: unsplash.com

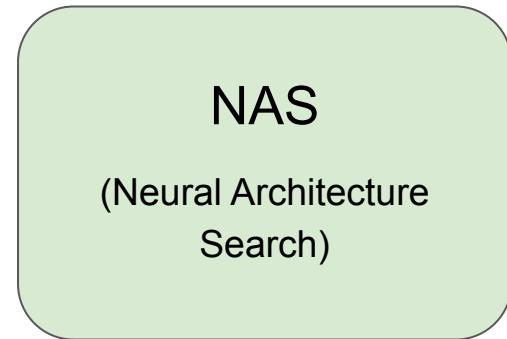
Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention

3) Case studies

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pQRNN
- MobileBERT
- Lite Transformer with Long-Short Range Attention
- MicroNet for Efficient Language Modeling
- Hardware-Aware Transformers



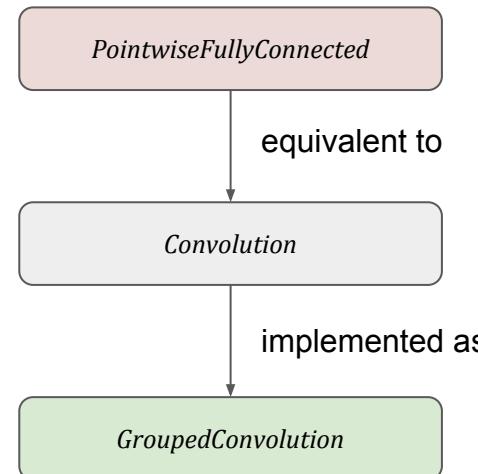
Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pQRNN
- MobileBERT
- Lite Transformer with Long-Short Range Attention
- MicroNet for Efficient Language Modeling
- Hardware-Aware Transformers
- SqueezeBERT

more expensive operation



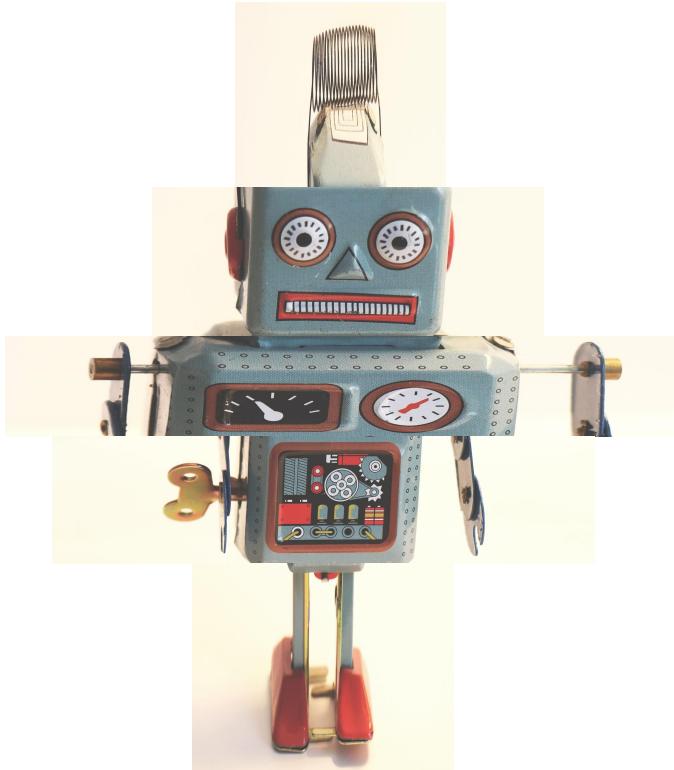
more efficient operation

Towards more efficient NLP

- 1) Core techniques
- 2) Efficient attention
- 3) Case studies**

a) Efficient Language Models

- Natural Language Processing with Small Feed-Forward Networks
- The Evolved Transformer
- PRADO + pQRNN
- MobileBERT
- Lite Transformer with Long-Short Range Attention
- MicroNet for Efficient Language Modeling
- Hardware-Aware Transformers
- SqueezeBERT
- DeLight: Very Deep and Light-weight Transformer



Natural Language Processing with Small Feed-Forward Networks

Botha et al., 2017

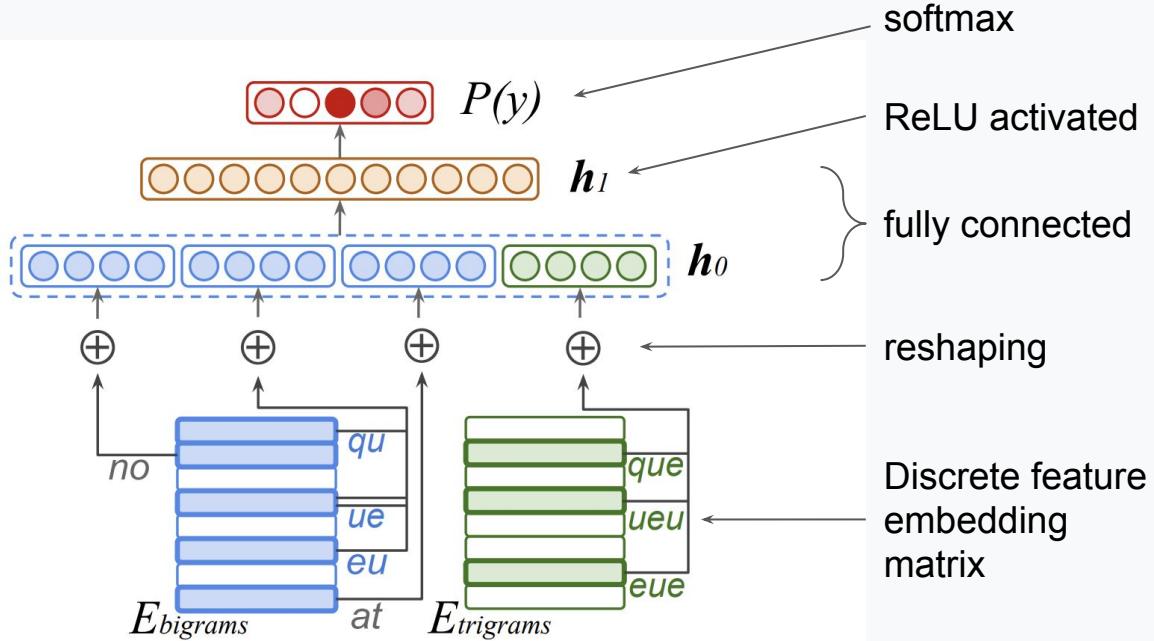
arxiv.org/abs/1708.00214

- Useful accuracies on a variety of tasks
 - Great runtime and memory value in resource constrained environments
-
- Features defined over character n-grams, embeddings learned from scratch
 - Random feature mixing hashing for small feature vocabularies
 - Quantization for embedding weight compression

Natural Language Processing with Small Feed-Forward Networks

Botha et al., 2017

arxiv.org/abs/1708.00214



There was no que at the ...

Figure 1: An example network structure for a model using bigrams of the previous, current and next word, and trigrams of the current word. Does not illustrate hashing.

Natural Language Processing with Small Feed-Forward Networks

Botha et al., 2017

arxiv.org/abs/1708.00214

Example result:

POS Tagging, compared to BTS ([Gillick et al., 2016](#))

- +0.3% accuracy (**95.4%**, near state-of-the-art)
- 6x fewer parameters
- 36x fewer FLOPs

The Evolved Transformer

- Consistent improvement over Transformer on well established WMT and LM1B.
- NAS to search Transformer alternatives
- Large search space from feed-forward sequence models
- Evolutionary architecture search

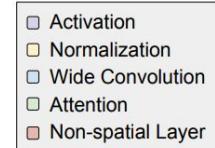
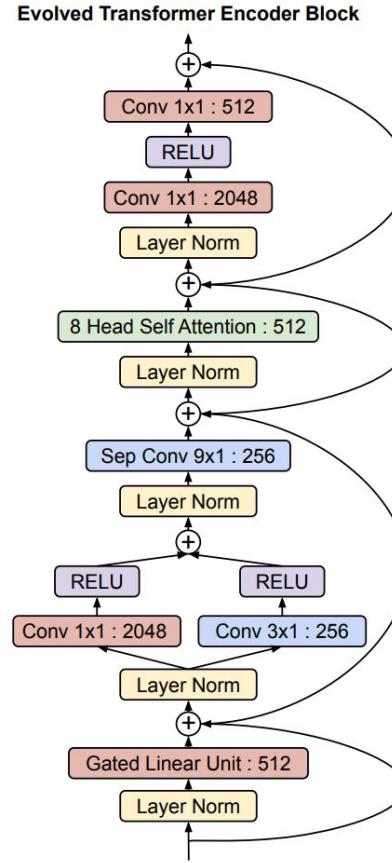
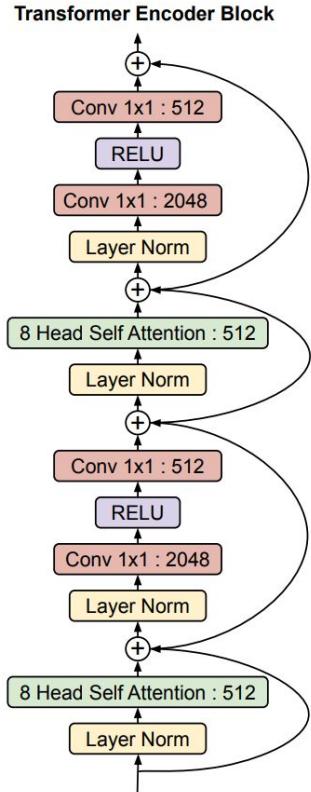
So et al., 2019

arxiv.org/abs/1901.11117

The Evolved Transformer

So et al., 2019

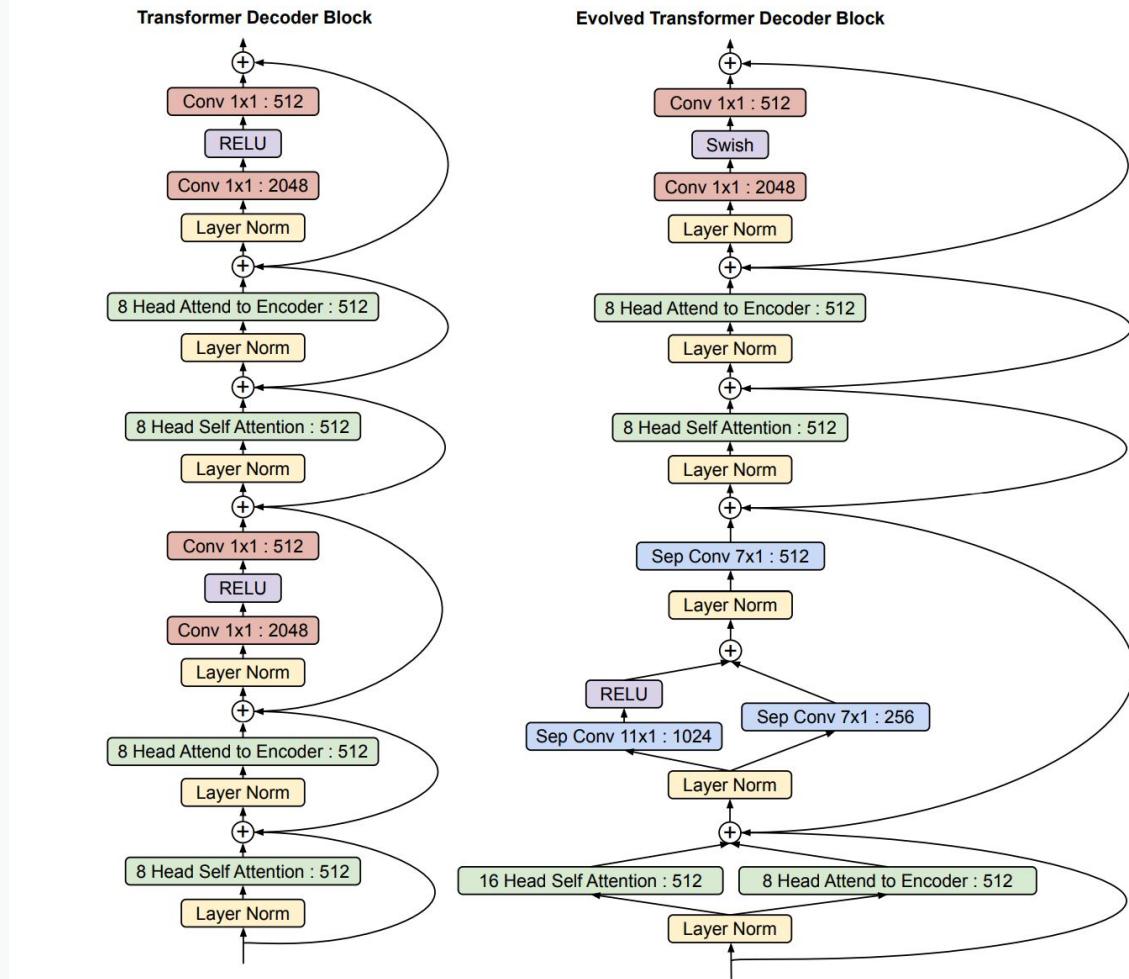
arxiv.org/abs/1901.11117



The Evolved Transformer

So et al., 2019

arxiv.org/abs/1901.11117

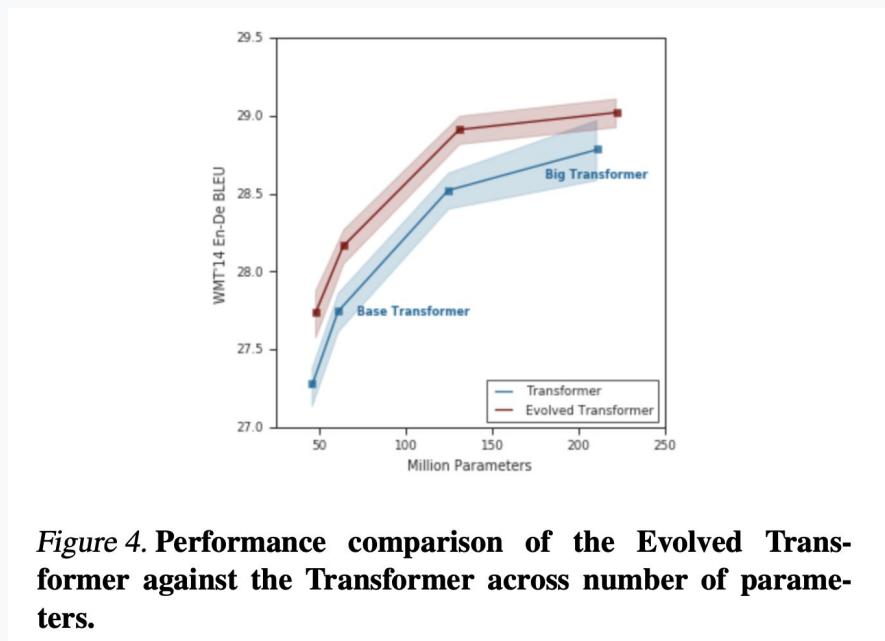


The Evolved Transformer

So et al., 2019

arxiv.org/abs/1901.11117

Same quality as original “big” Transformer with 37.6% fewer parameters and outperforms Transformer by 0.7 BLEU at a mobile-friendly model size of ~7M params.



PRADO + pQRNN

PRADO: Projection Attention Networks for Document Classification On-Device

- Combines trainable projections with attention and convolutions
- With only 200 Kilobytes in size, outperformed prior CNN and LSTM models and achieved near state of the art performance on multiple long document classification tasks.

Kaliamoorthi et al., 2019-2020

www.aclweb.org/anthology/D19-1506/

<https://ai.googleblog.com/2020/09/advancing-nlp-with-efficient-projection.html>

PRADO + pQRNN

Kaliamoorthi et al., 2019-2020

www.aclweb.org/anthology/D19-1506/

<https://ai.googleblog.com/2020/09/advancing-nlp-with-efficient-projection.html>

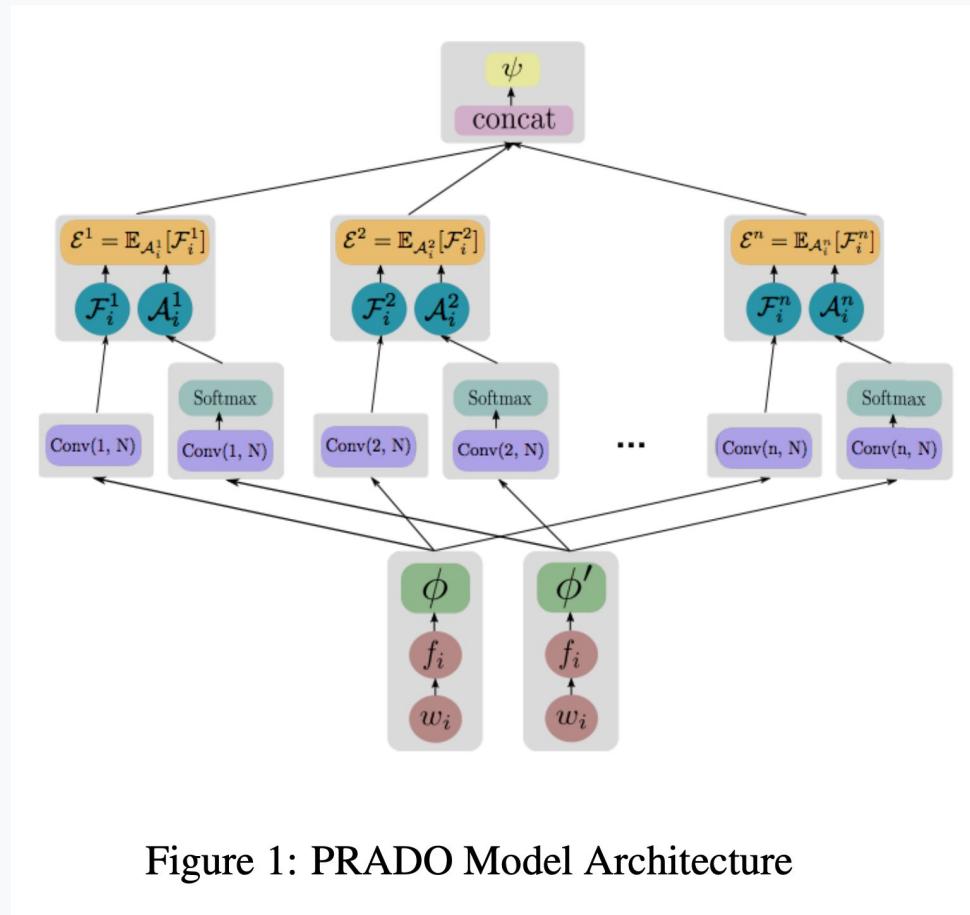
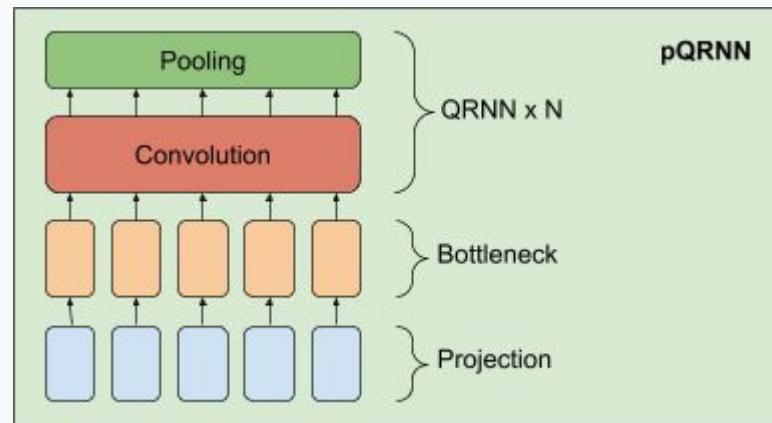


Figure 1: PRADO Model Architecture

PRADO + pQRNN

pQRNN

- A projection layer with a quasi-RNN encoder
- Same projection layer used in PRADO



Kaliamoorthi et al., 2019-2020

www.aclweb.org/anthology/D19-1506/

<https://ai.googleblog.com/2020/09/advancing-nlp-with-efficient-projection.html>

- pQRNN is also quantized

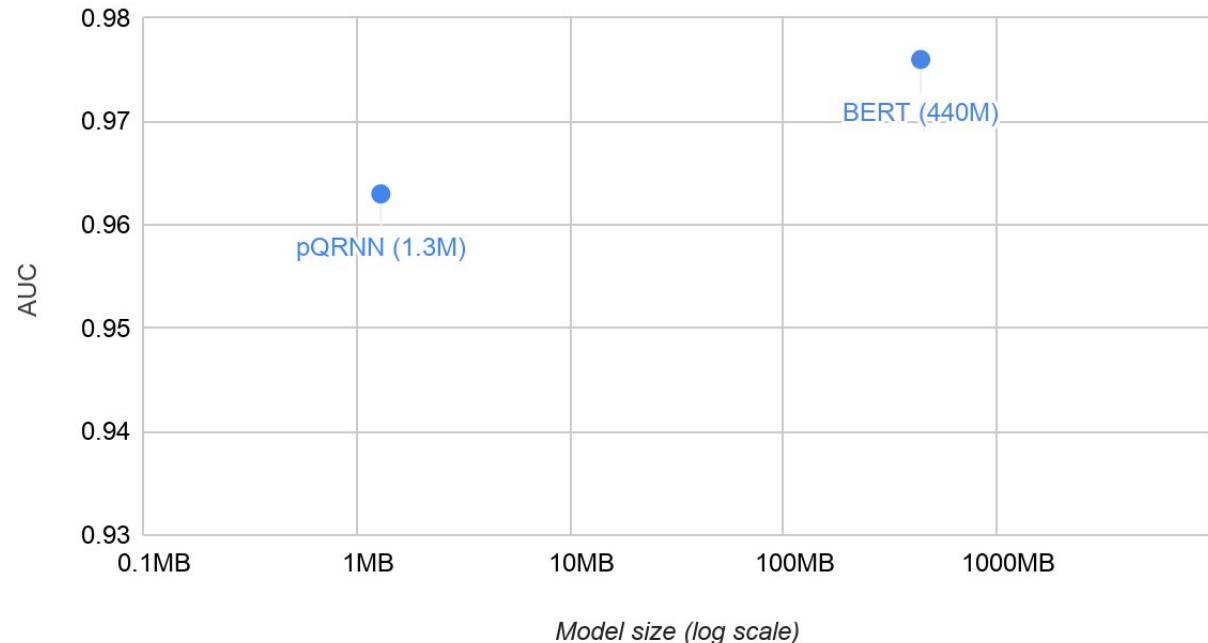
PRADO + pQRNN

Kaliamoorthi et al., 2019-2020

www.aclweb.org/anthology/D19-1506/

<https://ai.googleblog.com/2020/09/advancing-nlp-with-efficient-projection.html>

Model Comparison



MobileBERT

- Designed for running on mobile phones with acceptable latency
- Inverted-Bottleneck BERT_{LARGE} teacher
- Distilled into a compact MobileBERT student
- As deep as BERT_{LARGE}, but narrower
- Task-agnostic compression (task specific fine-tuning performed directly on the compact model)

Sun et al., 2020

arxiv.org/abs/2004.02984

MobileBERT

Sun et al., 2020

arxiv.org/abs/2004.02984

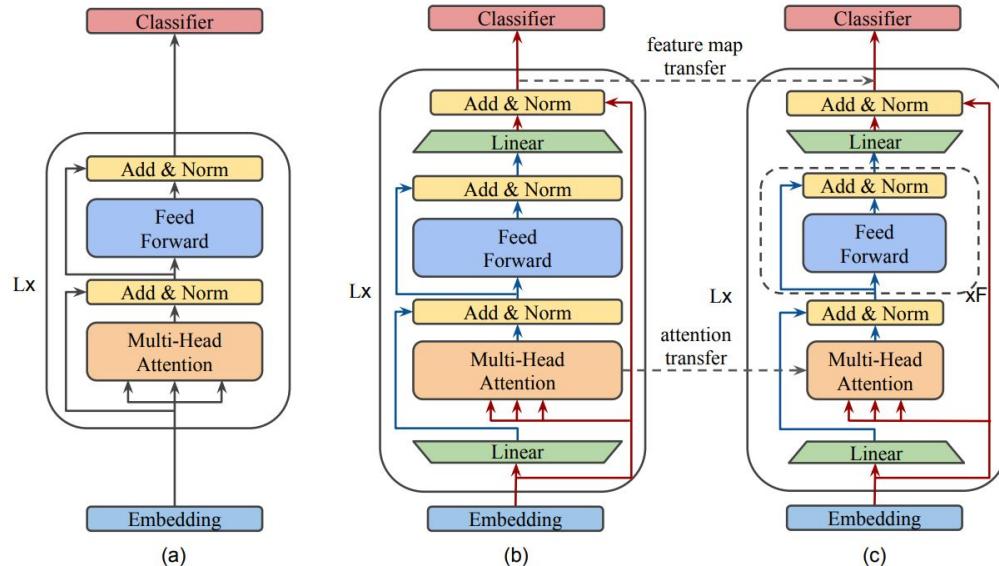


Figure 1: Illustration of three models: (a) BERT; (b) Inverted-Bottleneck BERT (IB-BERT); and (c) MobileBERT. In (b) and (c), **red lines** denote **inter-block flows** while **blue lines** **intra-block flows**. MobileBERT is trained by layer-to-layer imitating IB-BERT.

MobileBERT

Sun et al., 2020

arxiv.org/abs/2004.02984

		BERT _{LARGE}	MobileBERT	MobileBERT _{TINY}
embedding	$h_{\text{embedding}}$	1024		128
	no-op		3-convolution	
	h_{inter}	1024		512
body	Linear	h_{input} h_{output}	$\left[\begin{array}{c} 1024 \\ 16 \\ 1024 \\ 1024 \\ 4096 \\ 1024 \end{array} \right] \times 24$ $\left[\begin{array}{c} 512 \\ 128 \\ 512 \\ 4 \\ 128 \\ 128 \end{array} \right] \times 24$ $\left[\begin{array}{c} 512 \\ 128 \\ 128 \\ 4 \\ 128 \\ 128 \end{array} \right] \times 24$ $\left[\begin{array}{c} 512 \\ 128 \\ 128 \\ 512 \\ 128 \\ 128 \end{array} \right] \times 2$	$\left[\begin{array}{c} 512 \\ 128 \\ 512 \\ 4 \\ 128 \\ 128 \end{array} \right] \times 24$ $\left[\begin{array}{c} 128 \\ 512 \\ 128 \\ 128 \\ 128 \\ 512 \end{array} \right] \times 4$ $\left[\begin{array}{c} 128 \\ 512 \\ 128 \\ 128 \\ 128 \\ 512 \end{array} \right] \times 2$
	MHA	h_{input} #Head h_{output}		
	FFN	h_{input} h_{FFN} h_{output}		
	Linear	h_{input} h_{output}		
#Params		334M	25.3M	15.1M

MobileBERT

- 4.3x smaller, 5.5x faster than BERT_{BASE}
- 77.7 GLUE score ~ BERT_{BASE}
- 90.0/79.2 SQuAD v1.1/v2.0 F1 ~ BERT_{BASE}
- 62 ms latency on a Pixel 4 phone

	#Params	#FLOPS	Latency	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m/mm	QNLI	RTE	GLUE
				8.5k	67k	3.7k	5.7k	364k	393k	108k	2.5k	
BERT _{BASE}	109M	22.5B	342 ms	52.1	93.5	88.9	85.8	71.2	84.6/83.4	90.5	66.4	78.3
MobileBERT _{TINY}	15.1M	3.1B	40 ms	46.7	91.7	87.9	80.1	68.9	81.5/81.6	89.5	65.1	75.8
MobileBERT	25.3M	5.7B	62 ms	50.5	92.8	88.8	84.4	70.2	83.3/82.6	90.6	66.2	77.7
MobileBERT w/o OPT	25.3M	5.7B	192 ms	51.1	92.6	88.8	84.8	70.5	84.3/83.4	91.6	70.4	78.5

Sun et al., 2020

arxiv.org/abs/2004.02984

Lite Transformer with Long-Short Range Attention

Wu et al., 2020

arxiv.org/abs/2004.11886

- Long-Short Range Attention (LSRA)
 - Local context modeling by convolution
 - Long distance modeling by attention
- 2.5× reduced computation vs Transformer base
- 18.2× smaller with pruning and quantization
- 0.5 higher BLUE compared to Evolved Transformer, without the 250 GPU-year NAS cost.

MicroNet for Efficient Language Modeling

Yan et al., 2020

arxiv.org/abs/2005.07877

NeurIPS 2019 MicroNet Challenge¹

Language modeling track: train efficient word-level language models on the Wikitext-103 Dataset² (word-level perplexity < 35)

Score = Normalized Parameter Storage +
Normalized Math Operations

(Normalized by LSTM Rae et al, 2018 arxiv.org/abs/1803.10049)

¹Gale et al 2019, micronet-challenge.github.io

²Merity et al 2016, arxiv.org/abs/1609.07843

MicroNet for Efficient Language Modeling

Yan et al., 2020

arxiv.org/abs/2005.07877

- Core Language Model
 - Transformer-XL
 - Short Context Group Joint Optimization
 - Adaptive Embedding and Softmax
 - Hebbian Updates
- Compression Techniques
 - Knowledge Distillation
 - Pruning
 - Quantization

MicroNet for Efficient Language Modeling

- 90-fold reduction in parameter size and a 36-fold reduction in math operations compared to the MicroNet baseline

Yan et al., 2020

arxiv.org/abs/2005.07877

Hardware-Aware Transformers

- Neural Architecture Search
 - Train a *SuperTransformer* to cover a large space
 - Evolutionary search with hardware latency constraint to find a specialized *SubTransformer*
- Speed up and smaller size over baseline Transformer, and low search cost

Wang et al., 2020

arxiv.org/abs/2005.14187

Hardware-Aware Transformers

Wang et al., 2020

arxiv.org/abs/2005.14187

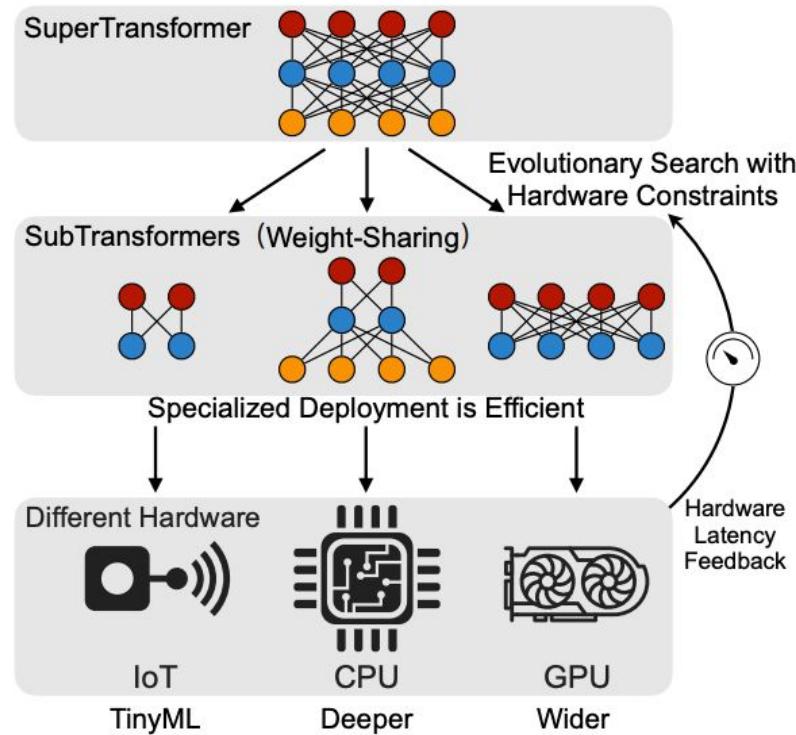


Figure 1: Framework for searching Hardware-Aware Transformers. We first train a SuperTransformer that contains numerous sub-networks, then conduct an evolutionary search with hardware latency feedback to find one **specialized** SubTransformer for each hardware.

Hardware-Aware Transformers

SubTransformer search

- Evolutionary search
- Find a satisfactory SubTransformer given a latency requirement
- Latency predictor trained for offline latency estimation (fast and accurate)

Wang et al., 2020

arxiv.org/abs/2005.14187

Hardware-Aware Transformers

WMT'14 results on Raspberry Pi-4:

- 3× speedup, 3.7× smaller size over baseline Transformer
- 2.7× speedup, 3.6× smaller size over Evolved Transformer with 12,041× less search cost

Wang et al., 2020

arxiv.org/abs/2005.14187

SqueezeBERT

- Replace several operations in self-attention layers with grouped convolutions
- Much faster inference on mobile devices

Iandola et al., 2020

arxiv.org/abs/2006.11316

SqueezeBERT

- Previous takeaways from CV into NLP (already adopted in MobileBERT)
 - Bottleneck layers
 - High-information flow residual connections
- New contributions from CV incorporated into SqueezeBERT's self-attention
 - Convolutions
 - Grouped convolutions

Landola et al., 2020

arxiv.org/abs/2006.11316

SqueezeBERT

- Results

- 4.3x faster than BERT-base (while MobileBERT is reported as 3.0x faster than BERT-base) on a Pixel 3 phone.
- GLUE score 76.9 (vs 79.0 for BERT-base)

Iandola et al., 2020

arxiv.org/abs/2006.11316

DeLighT: Very Deep and Light-weight Transformer

- More efficient parameter allocation within and across Transformer blocks
- Similar performance with substantially fewer parameters compared to baseline transformers.

Mehta et al., 2020

arxiv.org/abs/2008.00623

DeLighT: Very Deep and Light-weight Transformer

Mehta et al., 2020
arxiv.org/abs/2008.00623

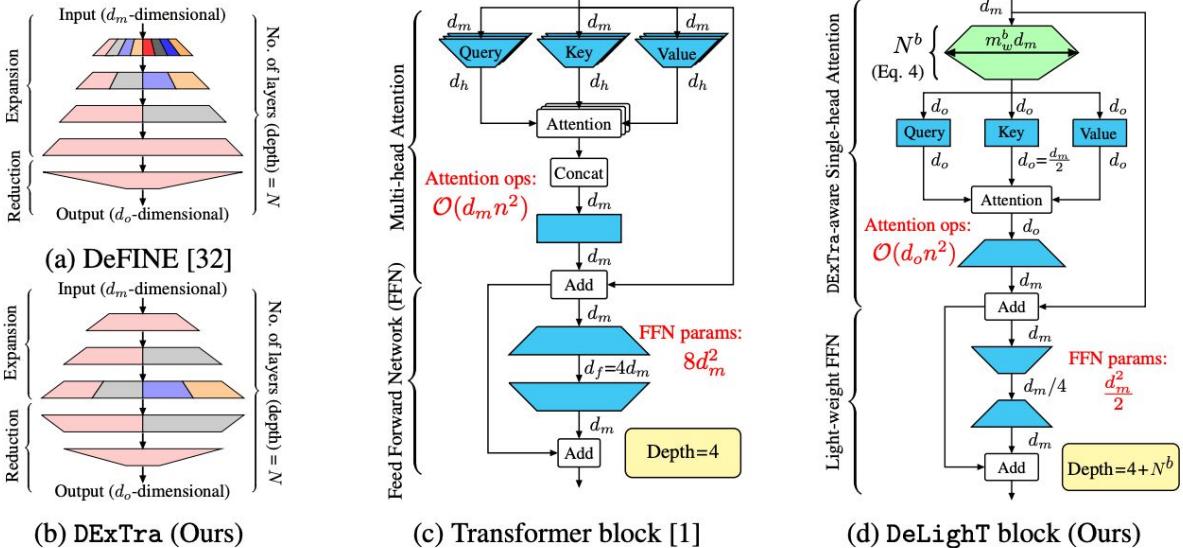
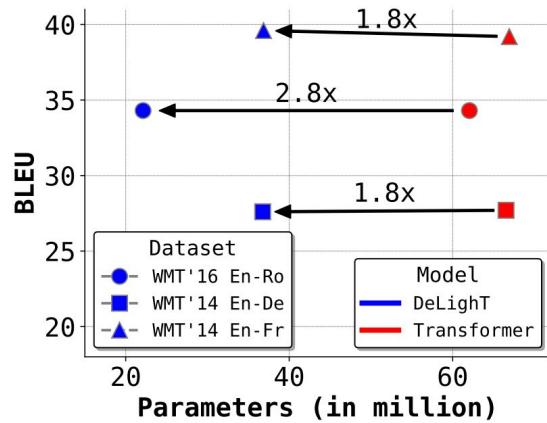


Figure 1: (a, b) compares the DeFINE unit with DExTra. Compared to the DeFINE unit, DExTra uses group linear transformations with more groups to learn wider representations with fewer parameters. Different colors are used to show groups in group linear transformations. For simplicity, we have not shown feature shuffling in (b). (c, d) Block-wise comparison between the standard transformer block and the DeLighT block. With DExTra, the number of operations in computing attention are reduced by half while the number of parameters (and operations) in the FFN are reduced by $16\times$. Layers with learnable parameters (Linear and DExTra) are shown in color. The shape of linear layers indicate their operation (expansion, reduction, etc.).

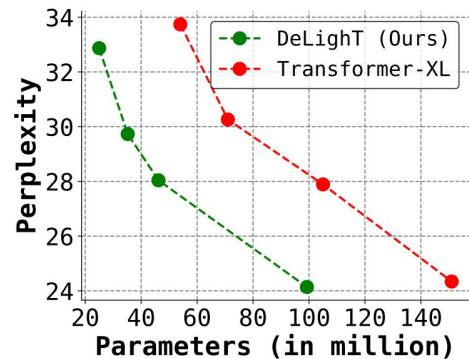
DeLighT: Very Deep and Light-weight Transformer

Mehta et al., 2020

arxiv.org/abs/2008.00623



(b) DeLighT vs. Transformer-XL



Retrieval



Towards more efficient NLP

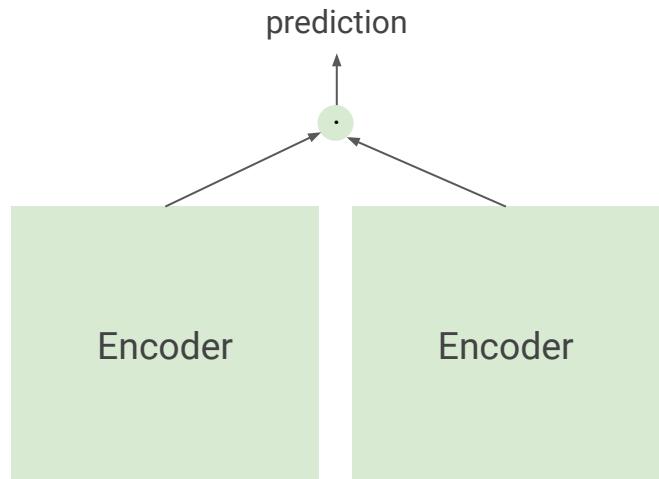
- 1) Core techniques
- 2) Efficient attention

3) Case studies

- a) Efficient Language Models

b) Retrieval

- Sentence Embeddings using Siamese BERT-Networks



Towards more efficient NLP

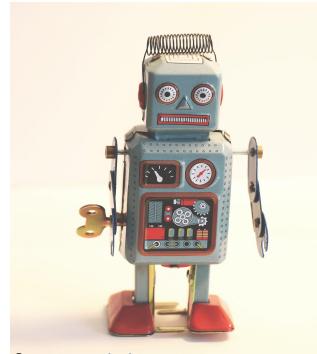
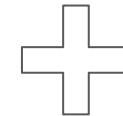
- 1) Core techniques
- 2) Efficient attention

3) Case studies

- a) Efficient Language Models

b) Retrieval

- Sentence Embeddings using Siamese BERT-Networks
- Generalization through Memorization: Nearest Neighbor Language Models



Source: unsplash.com

Towards more efficient NLP

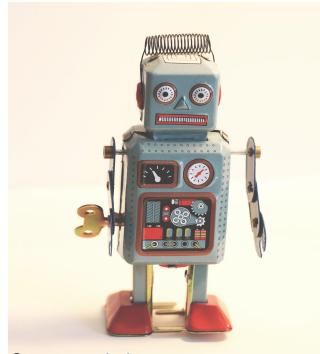
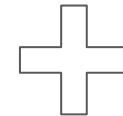
- 1) Core techniques
- 2) Efficient attention

3) Case studies

- a) Efficient Language Models

b) Retrieval

- Sentence Embeddings using Siamese BERT-Networks
- Generalization through Memorization: Nearest Neighbor Language Models
- REALM: Retrieval-Augmented Language Model Pre-Training



Source: unsplash.com

Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks

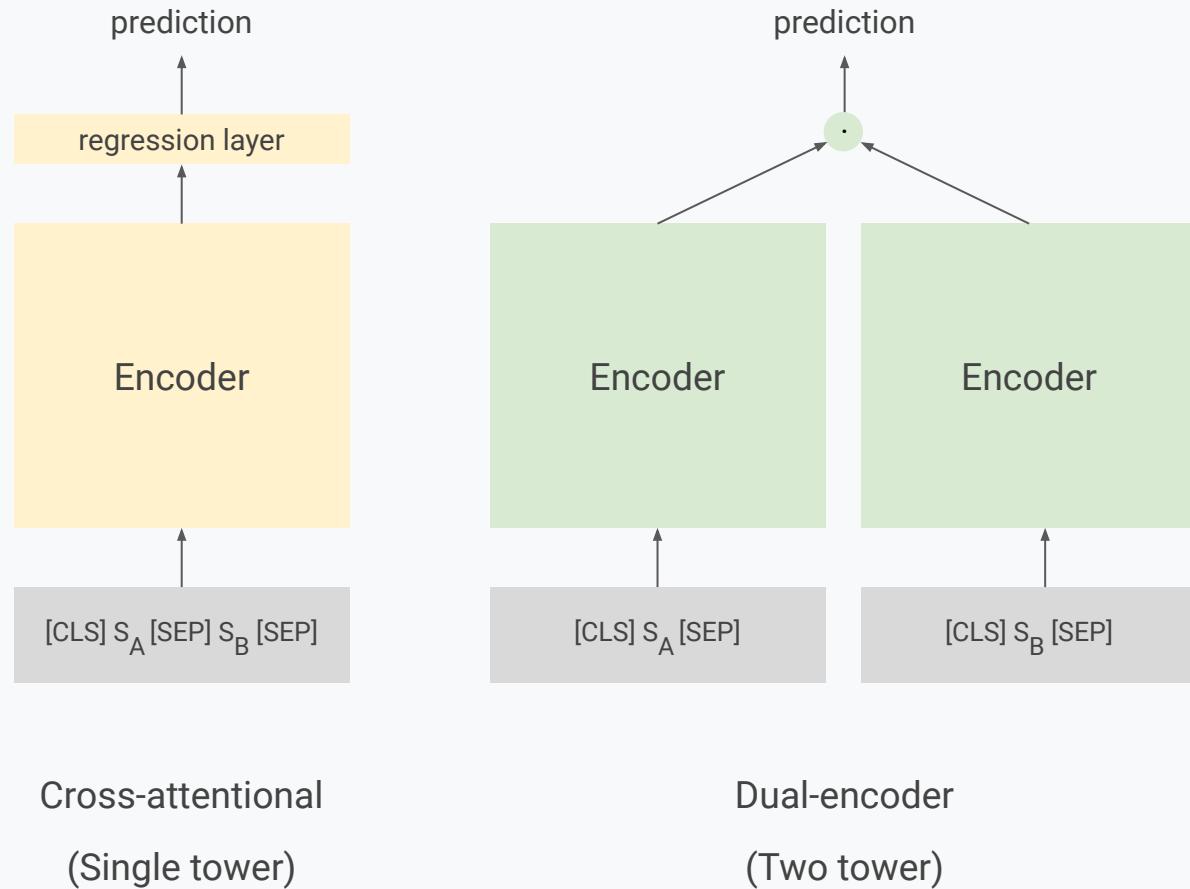
Reimers et al., 2019

arxiv.org/abs/1908.10084

- Cross-attention, single tower models such as BERT have set state-of-the-art results on sentence-pair tasks such as STS.
- For sentence-retrieval tasks, cross-attention model requires expensive re-encoding the entire retrieval corpus.
- Sentence-BERT modifies the pretrained encoder to perform a single inference per input sentence, followed by cheap pairwise comparisons e.g. cosine similarity.

Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks

Reimers et al., 2019
arxiv.org/abs/1908.10084



Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks

Reimers et al., 2019
arxiv.org/abs/1908.10084

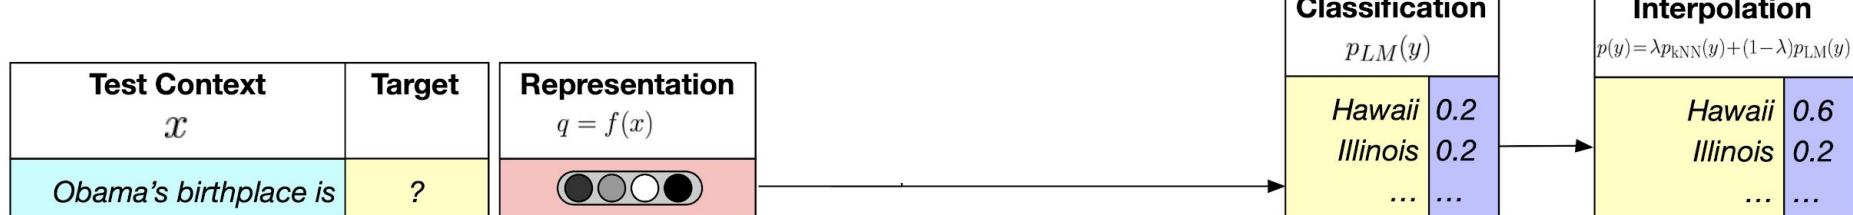
- Finding the most similar sentence in a collection of 10,000 sentences on a V100 GPU
 - BERT (cross-attention): 65 hours
 - SBERT (dual encoder): 5 seconds
- Can also be combined with Maximum Inner Product Search tools for sublinear scaling
 - <https://github.com/google-research/google-research/tree/master/scann>
 - <https://github.com/facebookresearch/faiss>
 - <https://github.com/spotify/annoy>

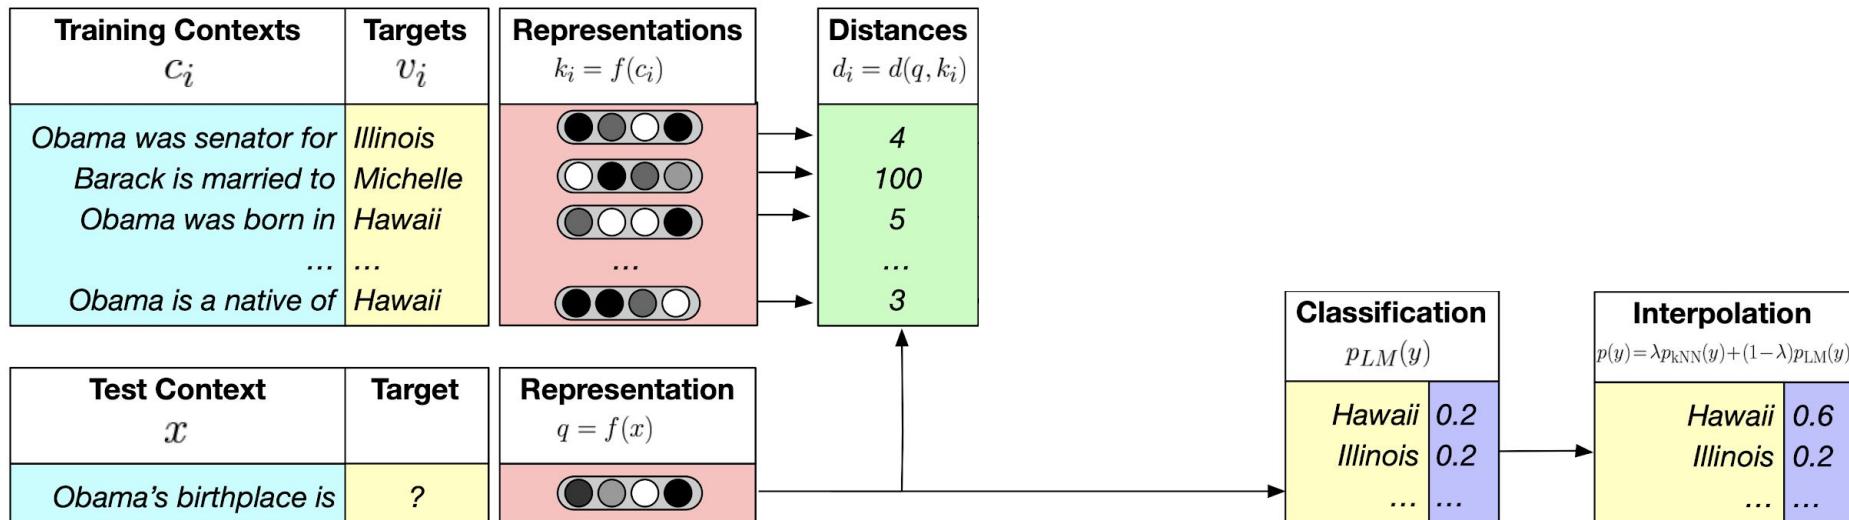
Generalization through Memorization: Nearest Neighbor Language Models

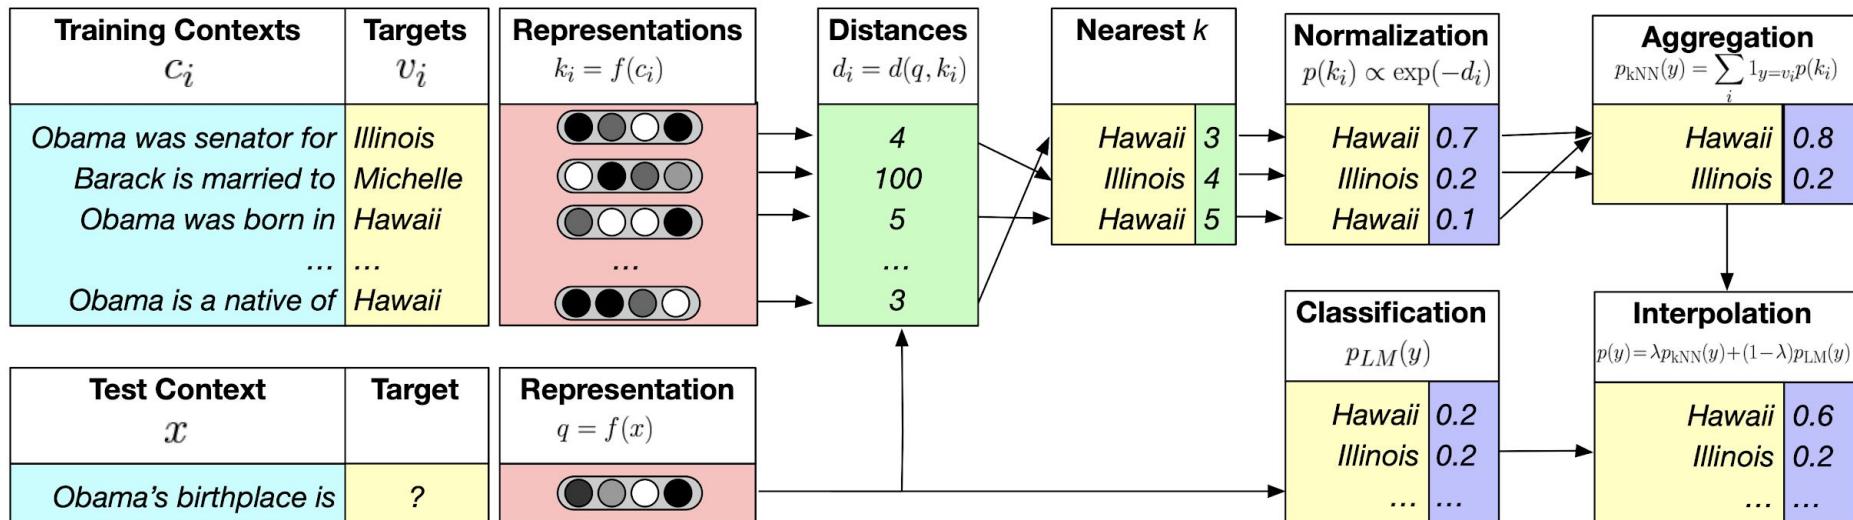
Khandelwal et al., 2019

arxiv.org/abs/1911.00172

- Introduces kNN-LMs, which extends a pre-trained neural language model (LM) by linearly interpolating it with a k-nearest neighbors (kNN) model.
- Allows for efficiently scaling up to larger training sets and for effective domain adaptation





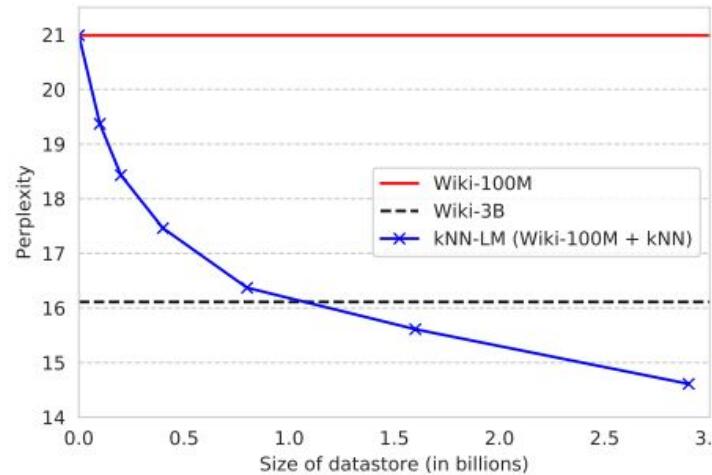


Generalization through Memorization: Nearest Neighbor Language Models

Khandelwal et al., 2019
arxiv.org/abs/1911.00172

Training Data	Datastore	Perplexity (\downarrow)	
		Dev	Test
WIKI-3B	-	16.11	15.17
WIKI-100M	-	20.99	19.59
WIKI-100M	WIKI-3B	14.61	13.73

Table 3: Experimental results on WIKI-3B. The model trained on 100M tokens is augmented with a datastore that contains about 3B training examples, outperforming the vanilla LM trained on the entire WIKI-3B training set.



(a) Effect of datastore size on perplexities.

REALM: Retrieval-Augmented Language Model Pre-Training

- Language model pre-training can capture world knowledge by storing it implicitly in the network parameters, but storage space is limited by the network size (prompting for ever-larger networks).
- REALM introduces a latent knowledge retriever to augment the language model, and shows for the first time how to pretrain it in an unsupervised manner.

Guu et al., 2020

arxiv.org/abs/2002.08909

REALM: Retrieval-Augmented Language Model Pre-Training

Guu et al., 2020

arxiv.org/abs/2002.08909

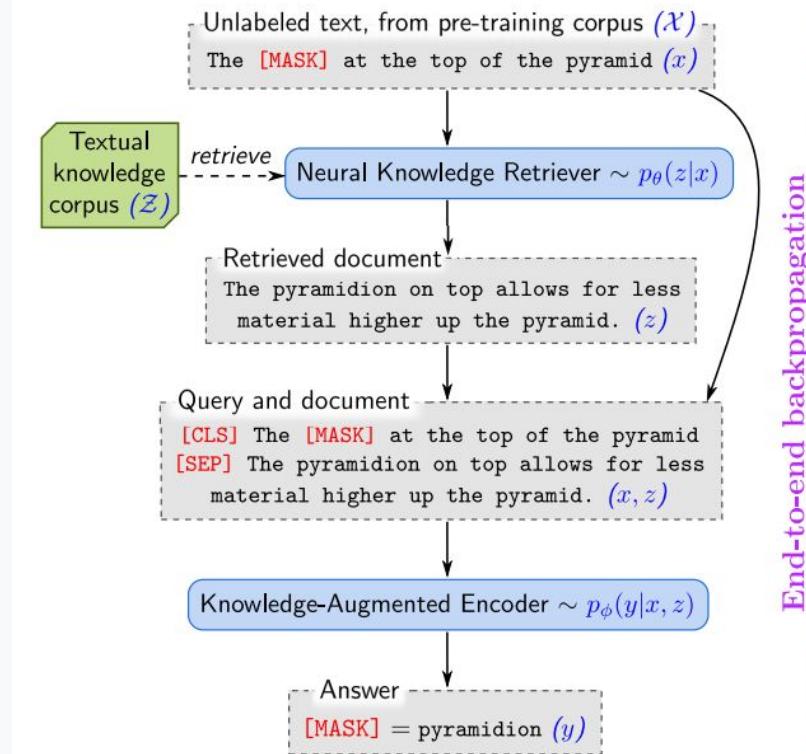


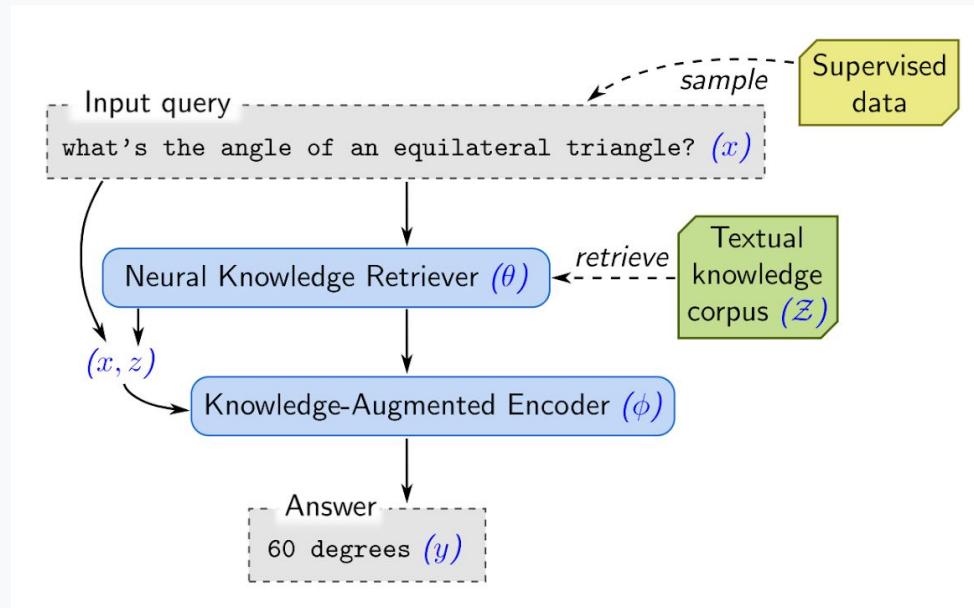
Figure 1. REALM augments language model pre-training with a **neural knowledge retriever** that retrieves knowledge from a **textual knowledge corpus**, \mathcal{Z} (e.g., all of Wikipedia). Signal from the language modeling objective backpropagates all the way through the retriever, which must consider millions of documents in \mathcal{Z} —a significant computational challenge that we address.

REALM: Retrieval-Augmented Language Model Pre-Training

Guu et al., 2020

arxiv.org/abs/2002.08909

- Fine-tuning for open-domain question answering



REALM: Retrieval-Augmented Language Model Pre-Training

Guu et al., 2020

arxiv.org/abs/2002.08909

- State-of-the-art Open-QA, with a relatively small model size (e.g. REALM outperforms T5-11b while being 30 times smaller)

Table 1. Test results on Open-QA benchmarks. The number of train/test examples are shown in parentheses below each benchmark. Predictions are evaluated with exact match against any reference answer. Sparse retrieval denotes methods that use sparse features such as TF-IDF and BM25. Our model, REALM, outperforms all existing systems.

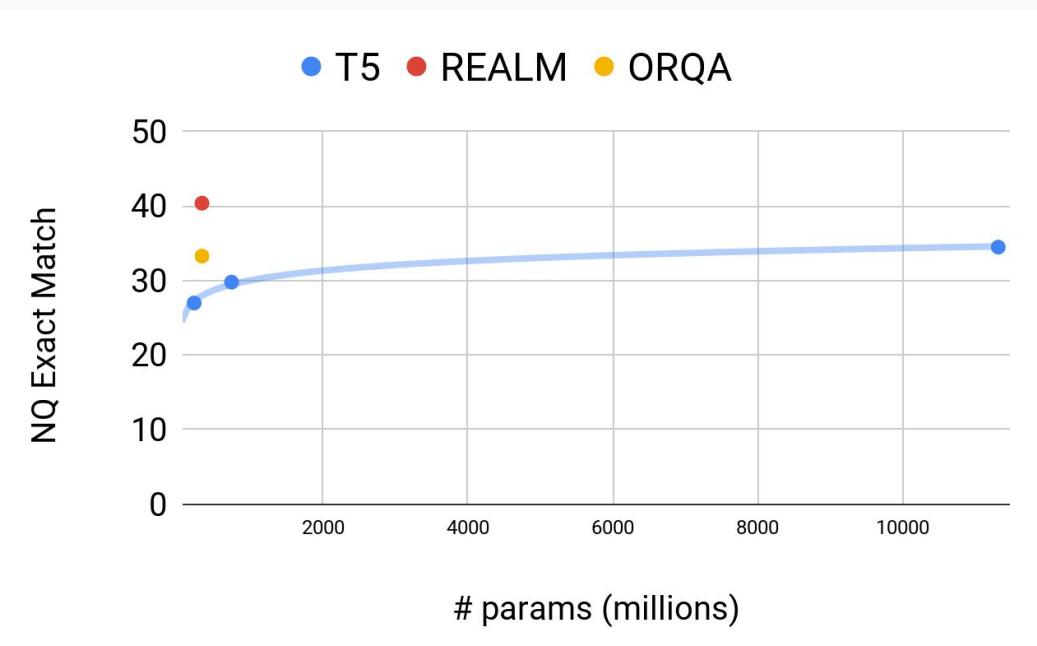
Name	Architectures	Pre-training	NQ (79k/4k)	WQ (3k/2k)	CT (1k /1k)	# params
BERT-Baseline (Lee et al., 2019)	Sparse Retr.+Transformer	BERT	26.5	17.7	21.3	110m
T5 (base) (Roberts et al., 2020)	Transformer Seq2Seq	T5 (Multitask)	27.0	29.1	-	223m
T5 (large) (Roberts et al., 2020)	Transformer Seq2Seq	T5 (Multitask)	29.8	32.2	-	738m
T5 (11b) (Roberts et al., 2020)	Transformer Seq2Seq	T5 (Multitask)	34.5	37.4	-	11318m
DrQA (Chen et al., 2017)	Sparse Retr.+DocReader	N/A	-	20.7	25.7	34m
HardEM (Min et al., 2019a)	Sparse Retr.+Transformer	BERT	28.1	-	-	110m
GraphRetriever (Min et al., 2019b)	GraphRetriever+Transformer	BERT	31.8	31.6	-	110m
PathRetriever (Asai et al., 2019)	PathRetriever+Transformer	MLM	32.6	-	-	110m
ORQA (Lee et al., 2019)	Dense Retr.+Transformer	ICT+BERT	33.3	36.4	30.1	330m
Ours (\mathcal{X} = Wikipedia, \mathcal{Z} = Wikipedia)	Dense Retr.+Transformer	REALM	39.2	40.2	46.8	330m
Ours (\mathcal{X} = CC-News, \mathcal{Z} = Wikipedia)	Dense Retr.+Transformer	REALM	40.4	40.7	42.9	330m

REALM: Retrieval-Augmented Language Model Pre-Training

Guu et al., 2020

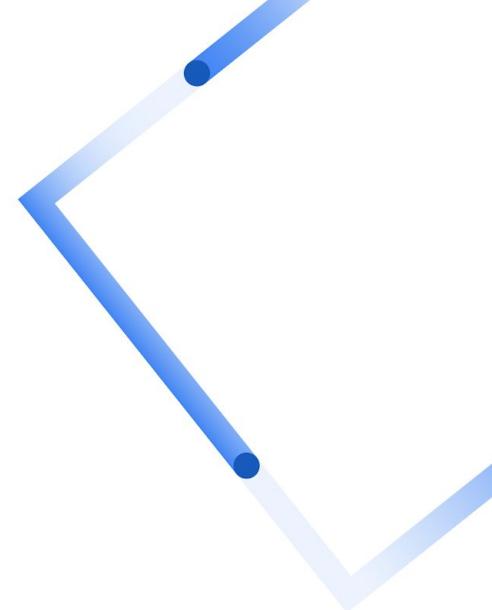
arxiv.org/abs/2002.08909

- State-of-the-art Open-QA, with a relatively small model size (e.g. REALM outperforms T5-11b while being 30 times smaller)



06

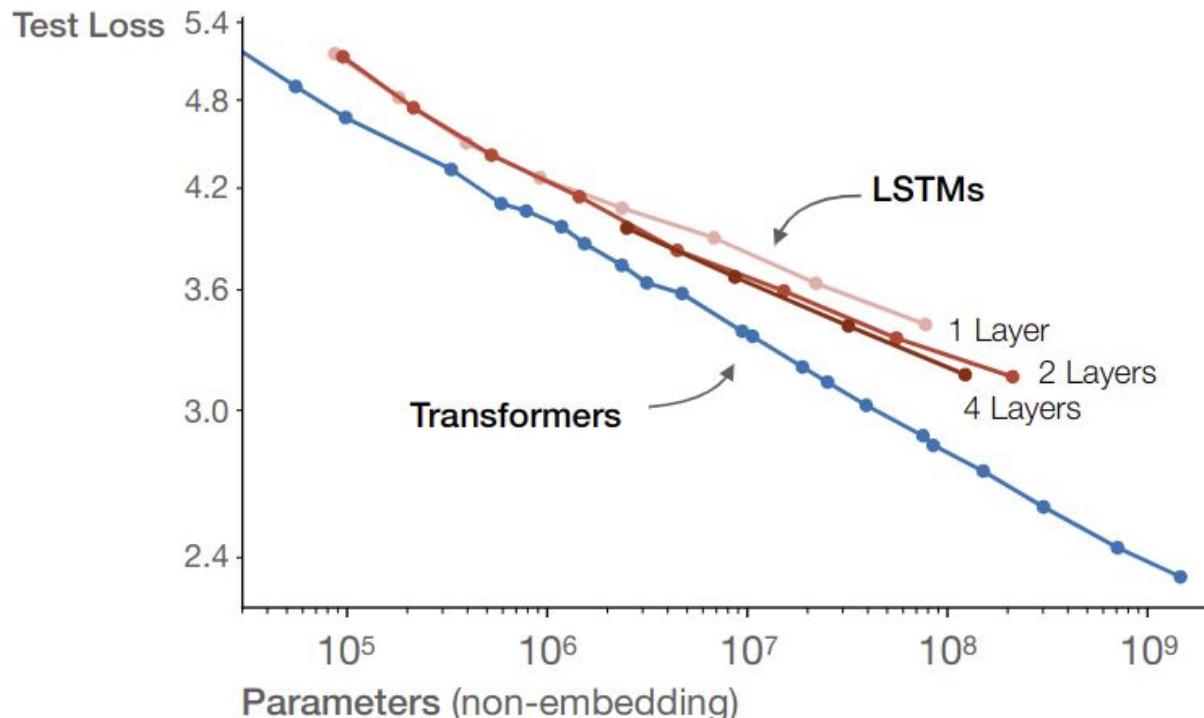
Scaling in Practice



Why Do We Need Scale?

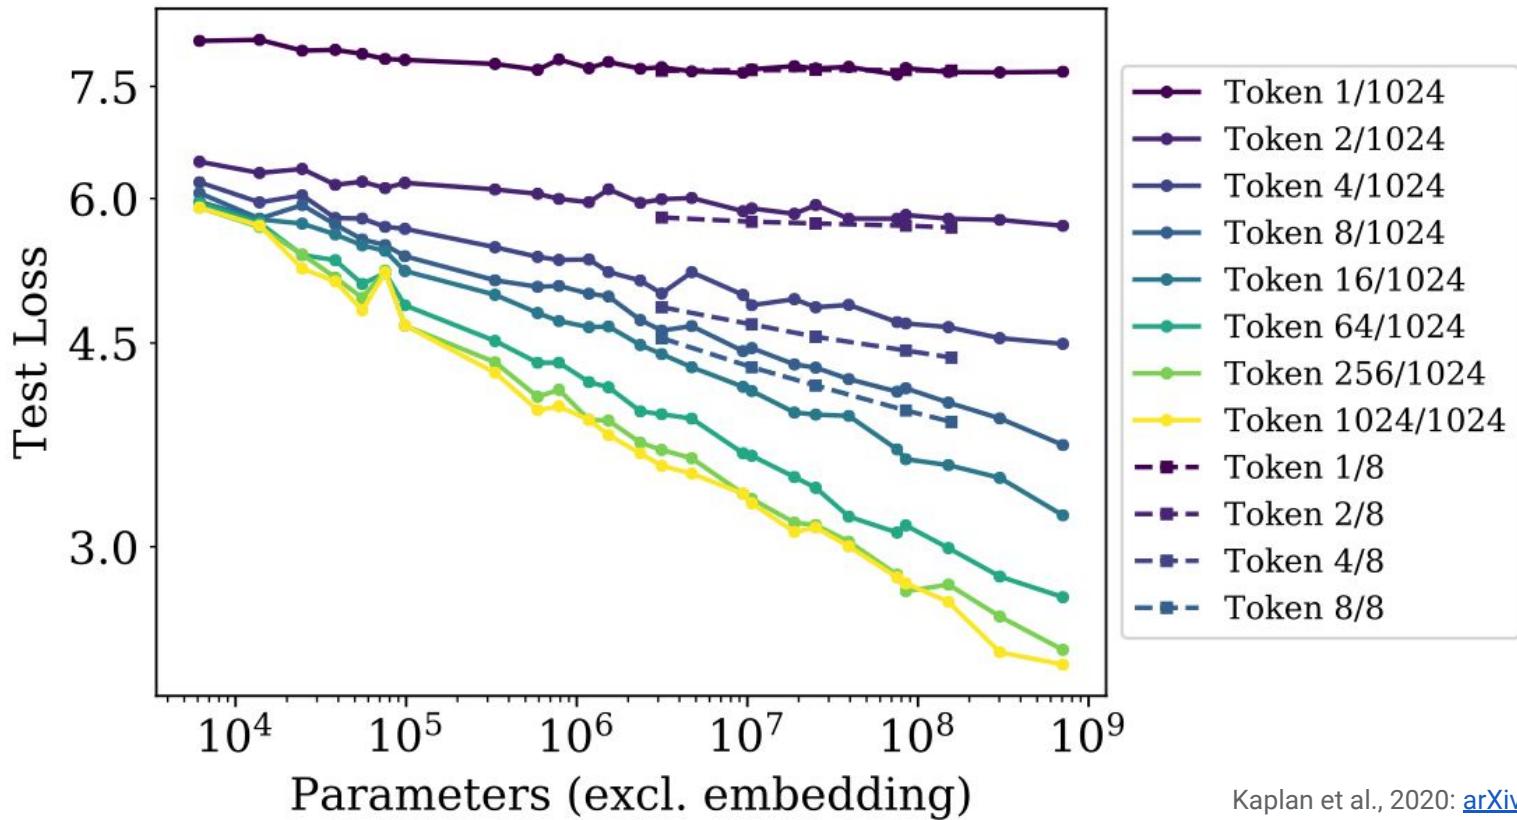
Scale More Important Than Architecture

**Transformers asymptotically outperform LSTMs
due to improved use of long contexts**



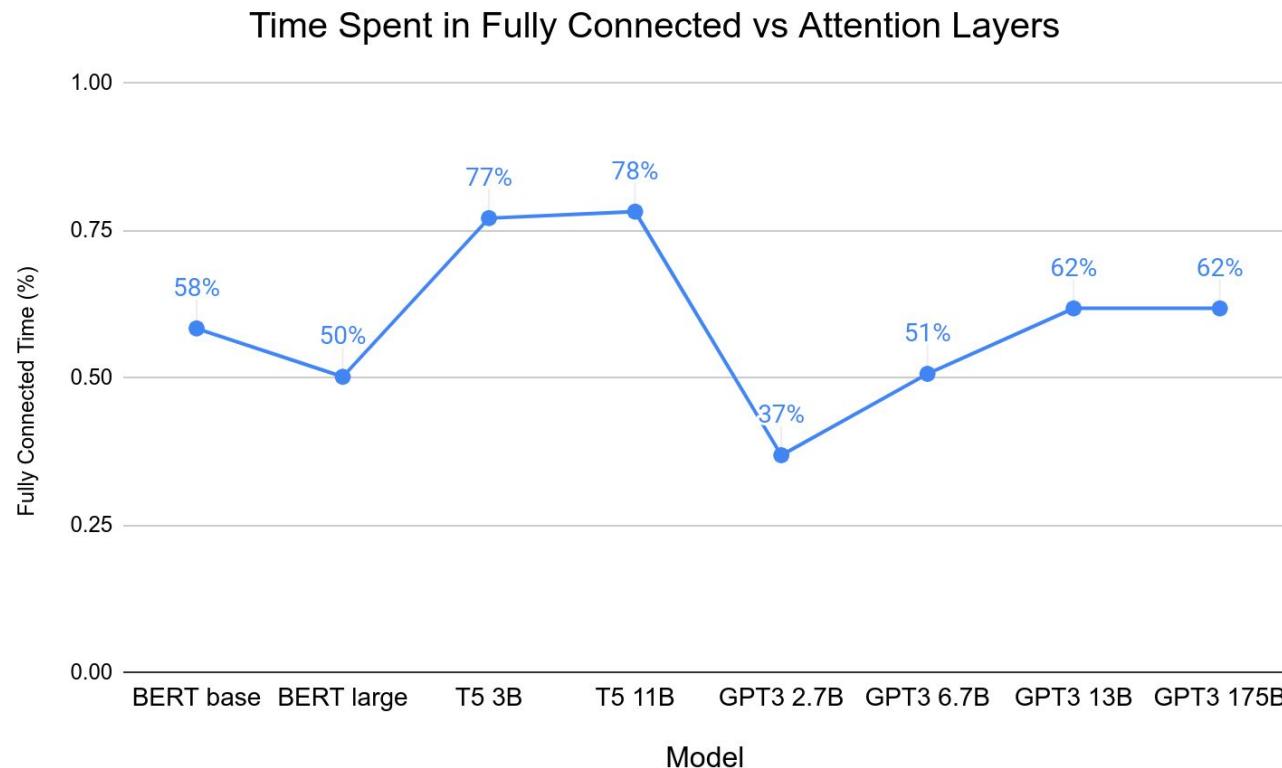
Kaplan et al., 2020: [arXiv](#)

Attention Size vs Model Size vs Test Loss



Kaplan et al., 2020: [arXiv](#)

Attention vs Fully Connected Time for Various Transformers



Conclusions From Measuring Scaling

- Performance increases further and further the more parameters a model has
- Attention is very important for efficiency: Transformers scale better than LSTMs
- Attention has diminishing returns (on general “internet data”)
- Size of data and model are more important than architecture

Practical Considerations

Experimental vs Theoretical Perspective

Theoretical:

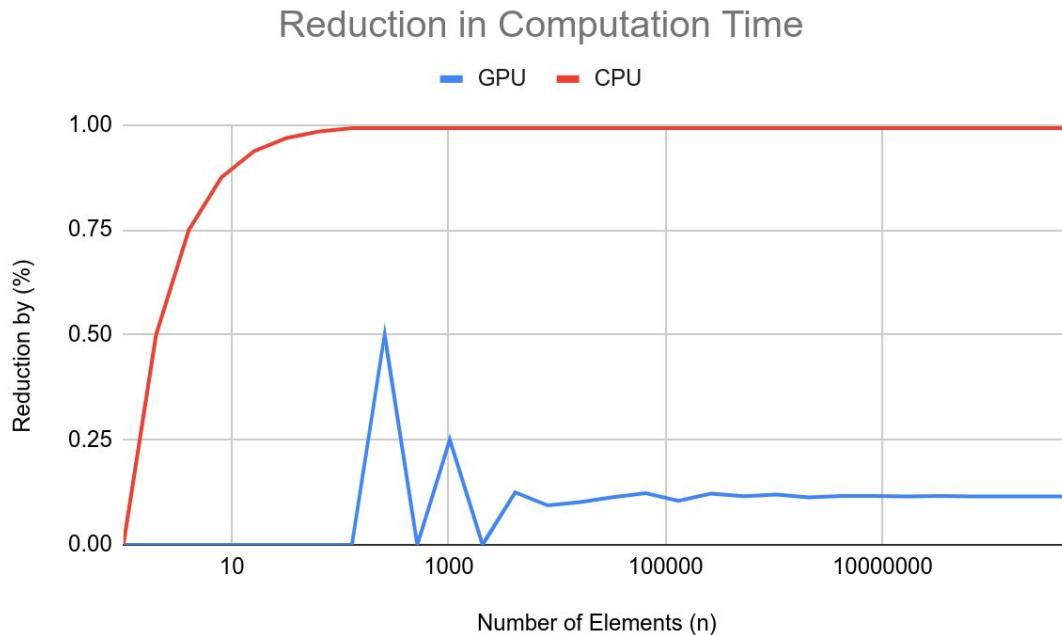
- FLOPS/Operation Complexity/Memory: $O(n)$ better than $O(n^2)$; 100 FLOPS better than 1000
- (Possibly) Analysis of occupancy, memory access patterns for certain hardware

Experimental:

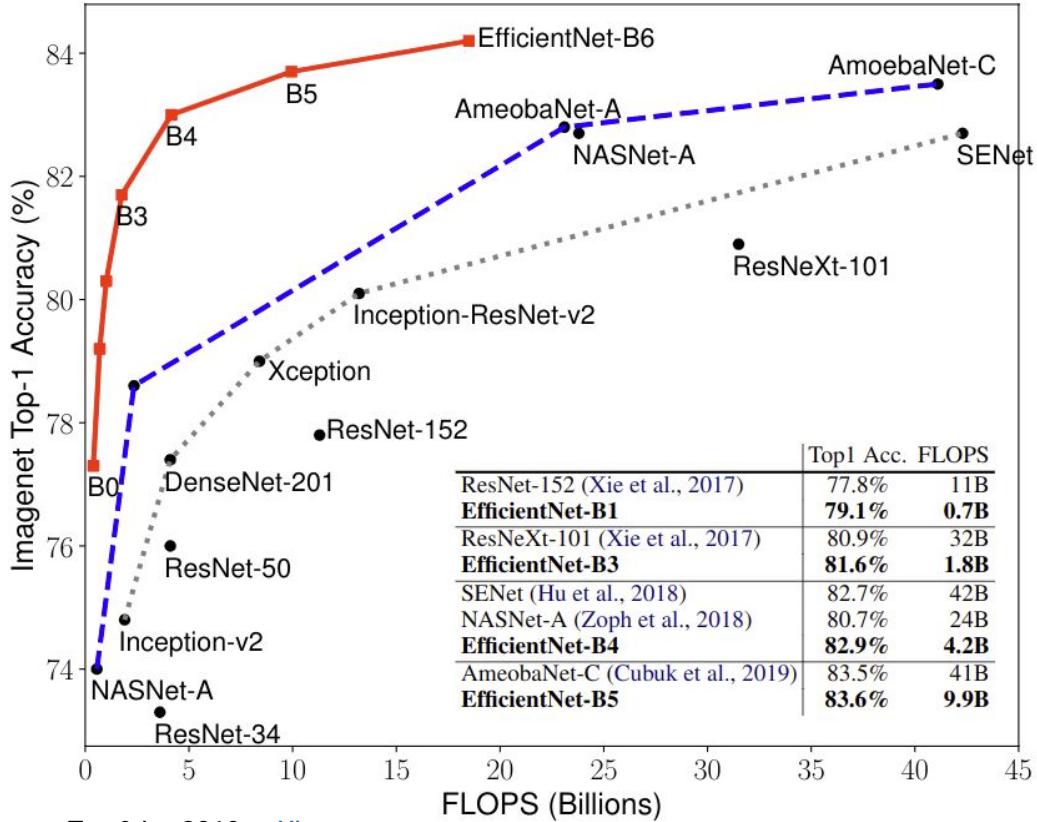
- Three criteria:
 - a. Does it fit into my GPU/TPU/Accelerator?
 - b. Is it faster than other methods?
 - c. Can most people use it (+62% of PhD students)?
- Device oriented walltime/memory: CPU for inference, GPU/TPU for training

Theory vs Practice

- Algorithm: (1) Divide matrix B into chunks of 128; (2) take the maximum element, set others to zero; (3) perform matrix multiply $A \times B = C$ and skip all zero elements



Theory vs Practice



Jeremy Howard
@jeremyphoward

Has anyone seen EfficientNet train faster than resnet50 to the same accuracy? I know the paper says it can be faster for inference, and it has less params, but experiments I've seen for training so far show it train much slower.
e.g. RN50 is 67% for this:



6:26 AM · Jun 5, 2019 · Twitter Web App

53 Retweets 230 Likes

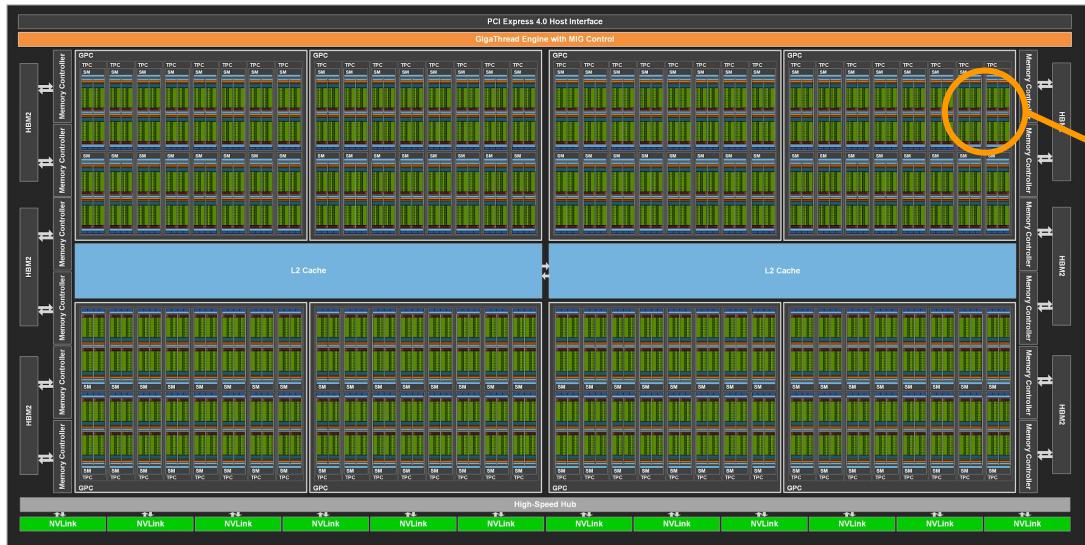


Ross Wightman @wightman · Jun 5, 2019

Replying to @jeremyphoward
All of the models in this family (MNASNet, FBNet, MobileNet-v3, EfficientNet) are pretty challenging and slow to train to spec'd accuracy. Epoch wise, finding the right hyper-params/techniques, and GPU memory use... tricks that work with ResNet, etc. also don't work as well here

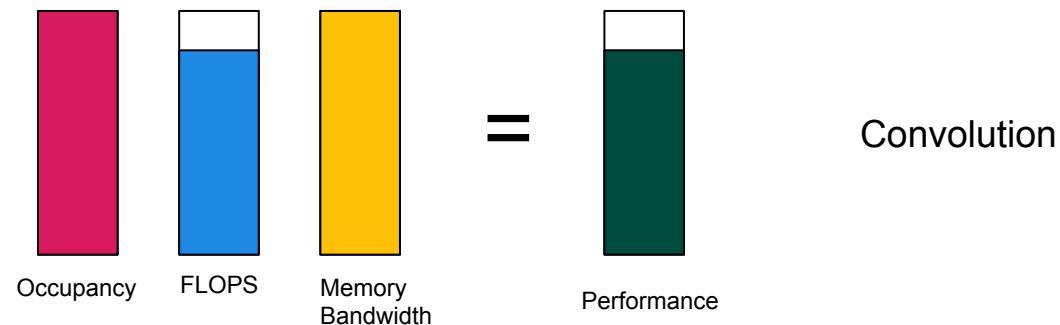
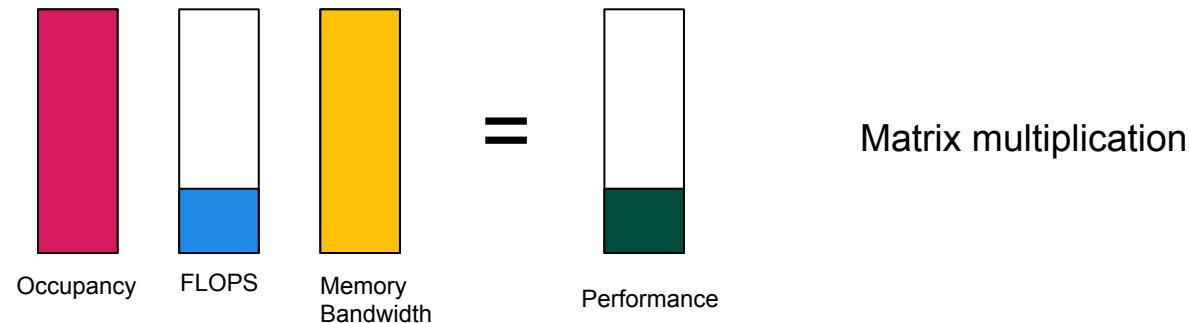


GPU Architecture

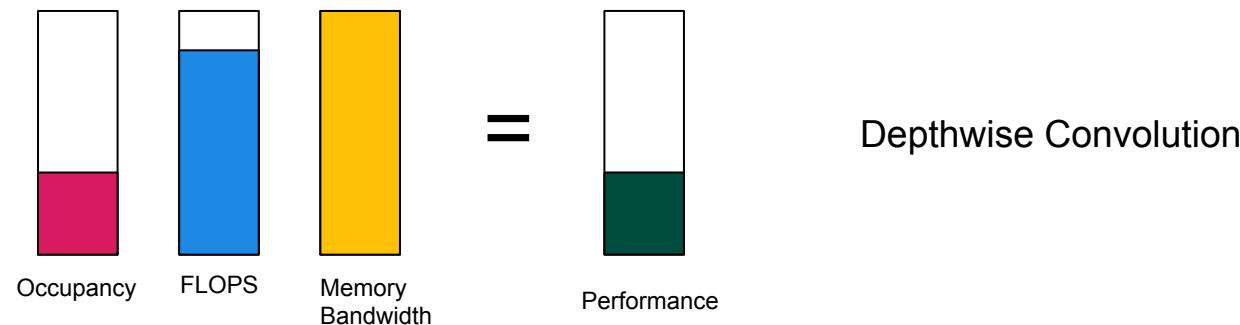
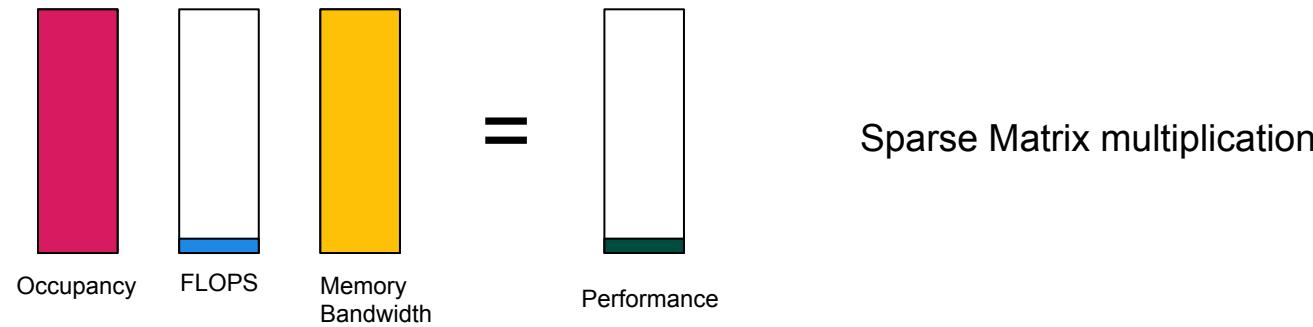


Ampere Architecture ([NVIDIA](#))

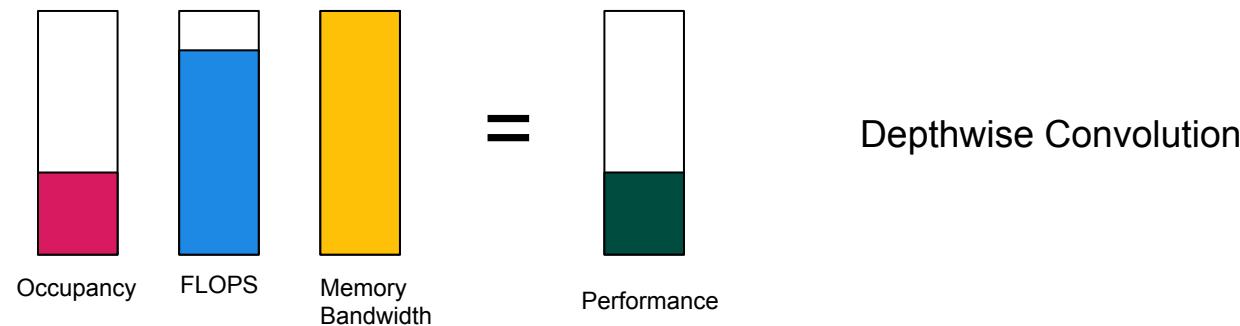
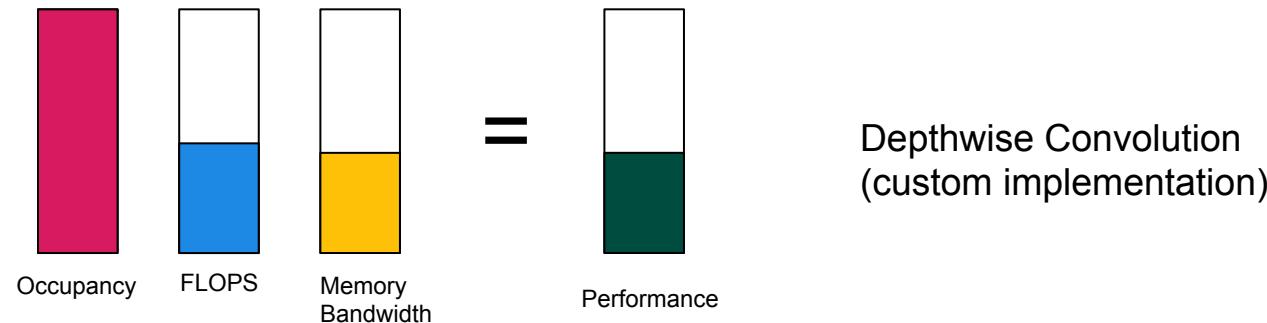
Occupancy vs Memory Bandwidth vs FLOPS



Occupancy vs Memory Bandwidth vs FLOPS

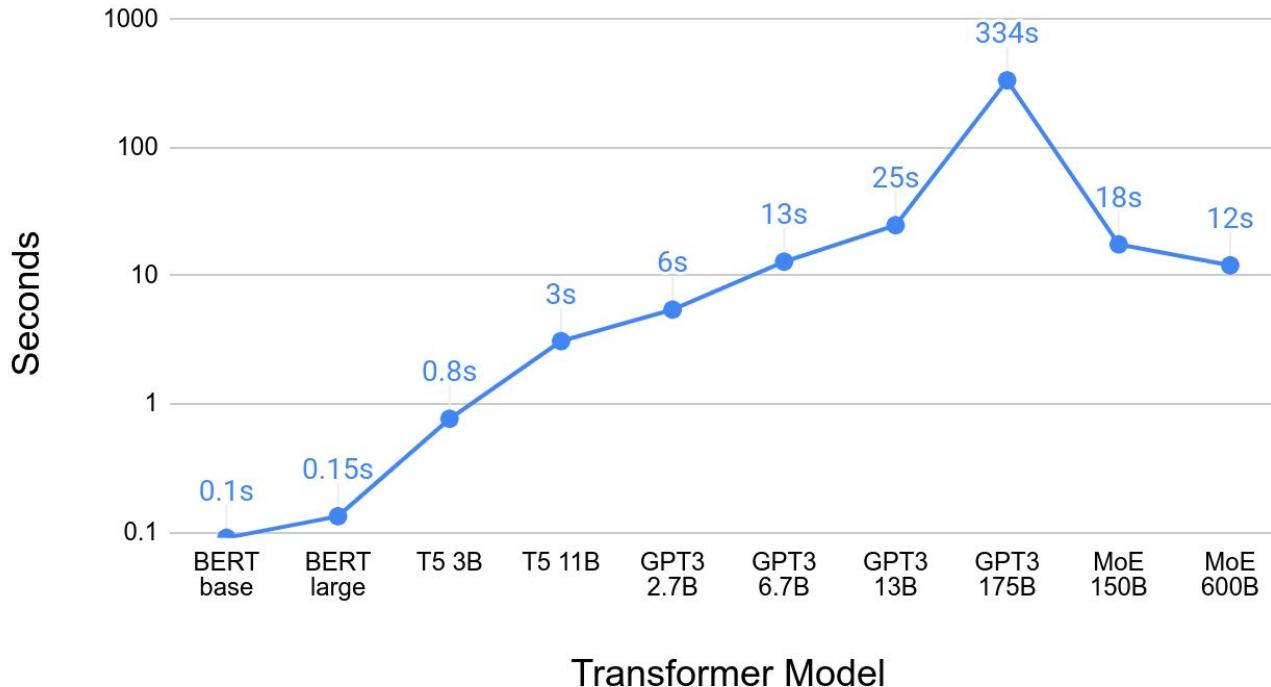


Occupancy vs Memory Bandwidth vs FLOPS



BERT Large vs BERT Base

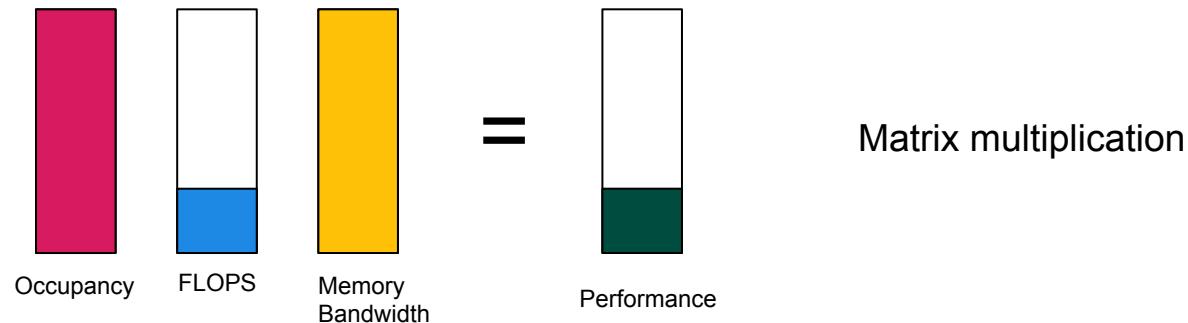
Mini-batch Time in Seconds for Transformer Models (Training)



BERT Base is 3.1x smaller than BERT Large but only trains 1.5x faster.

BERT Base is too small to saturate modern GPUs.

Better Performance at Lower Occupancy



Conclusion

- Occupancy, and FLOPS/memory bandwidth utilization are important for runtime performance
- Understanding of hardware needed for performance analysis
- Even with deep understanding of hardware, it is difficult to analyze performance theoretically
- Runtime performance of different algorithms can often only be understood if they are run on the actual device
- Conclusion: To estimate deep neural network runtime performance, it is best to run the network and measure its performance directly.

Memory Optimizations

Resources: Academia vs Industry



Tim Dettmers
@Tim_Dettmers

I want to get an accurate picture of GPU resources that PhD students have access to. PhD students, please respond and share with other students.

"What is the largest GPU system that you have access to?"

Please pick option +16 GPU only if your cluster has a +50 Gb/s interconnect

1-2 GPU desktop

43.1%

3-4 GPU desktop

18.8%

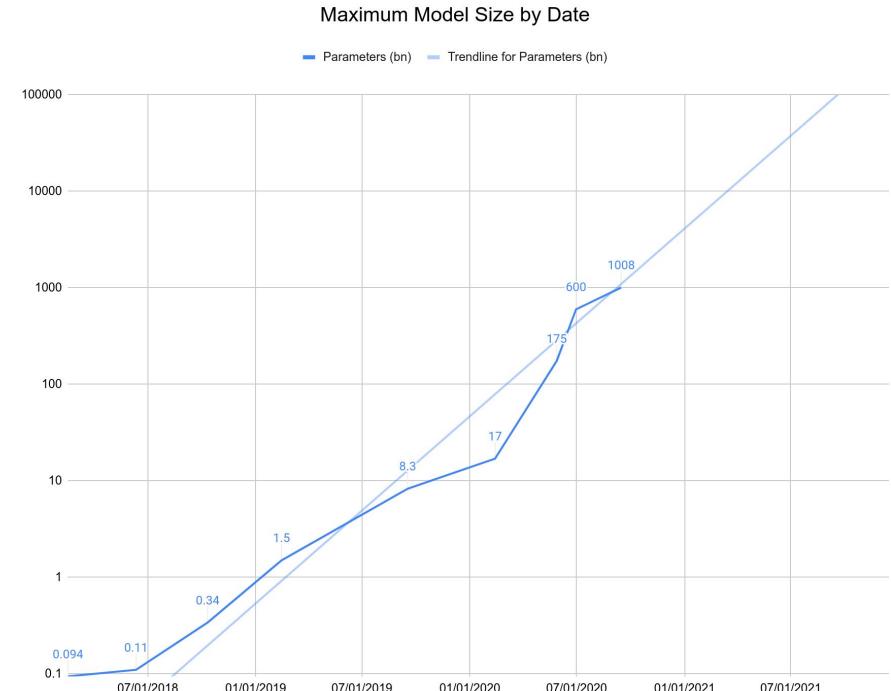
8 GPU server

21.6%

+16 GPU HPC cluster

16.5%

756 votes · Final results



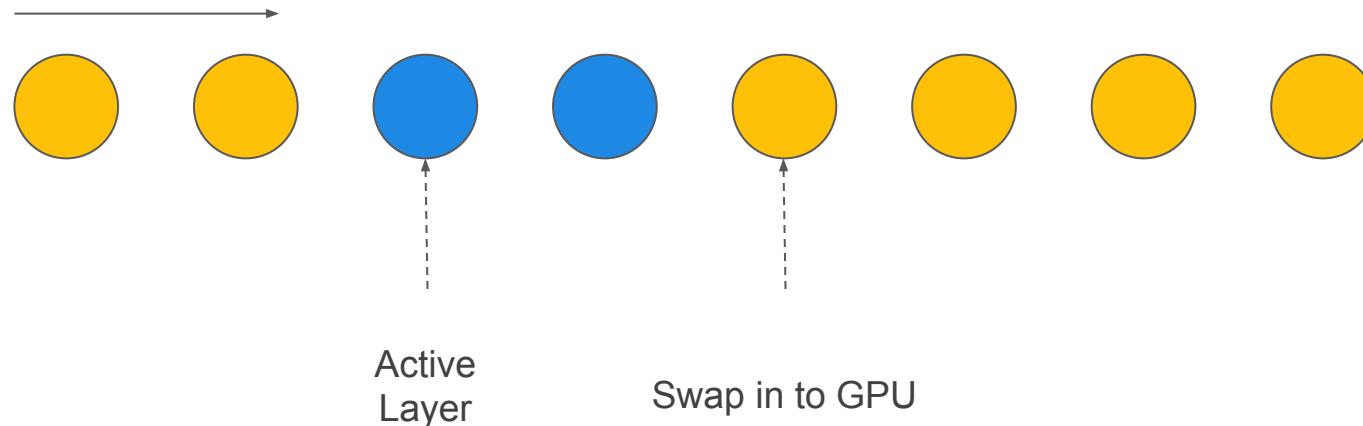
Memory Optimizations Overview

- Memory Swapping/Memory Paging
- FP16/BF16 training
- Gradient checkpointing
- Gradient accumulation
- Reversible residual connections

CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time

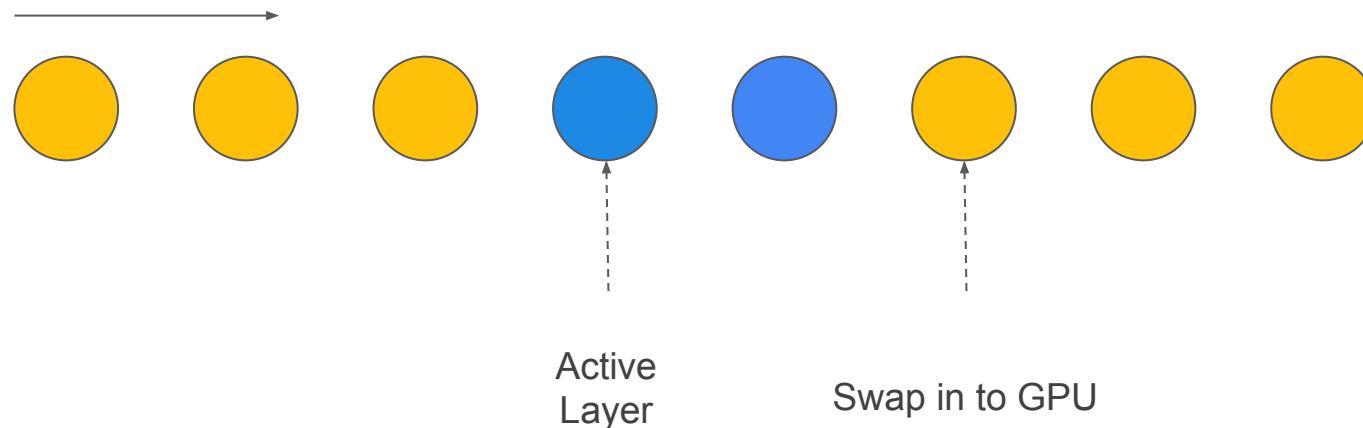
■ CPU ■ GPU



CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time

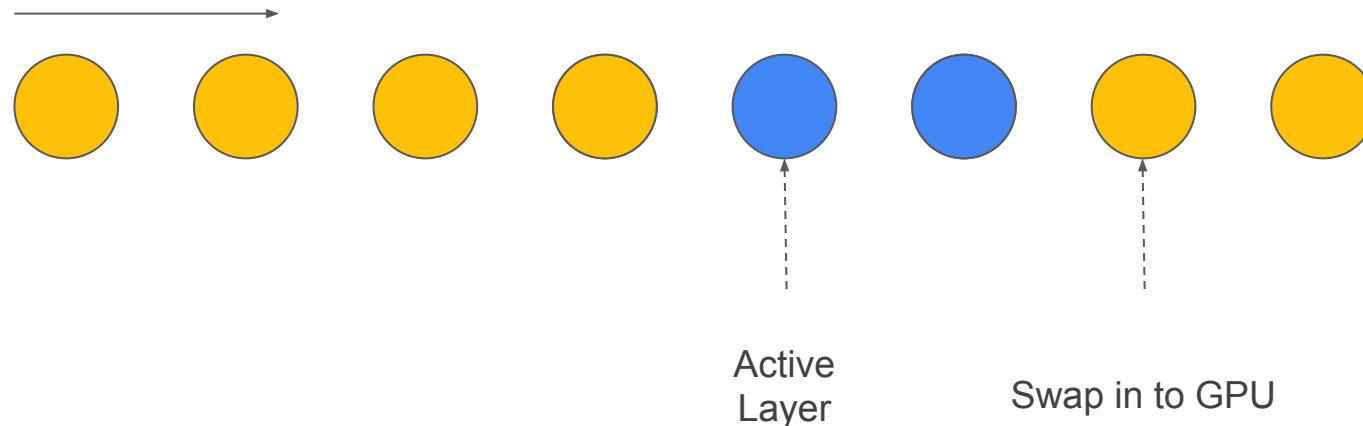
■ CPU ■ GPU



CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time

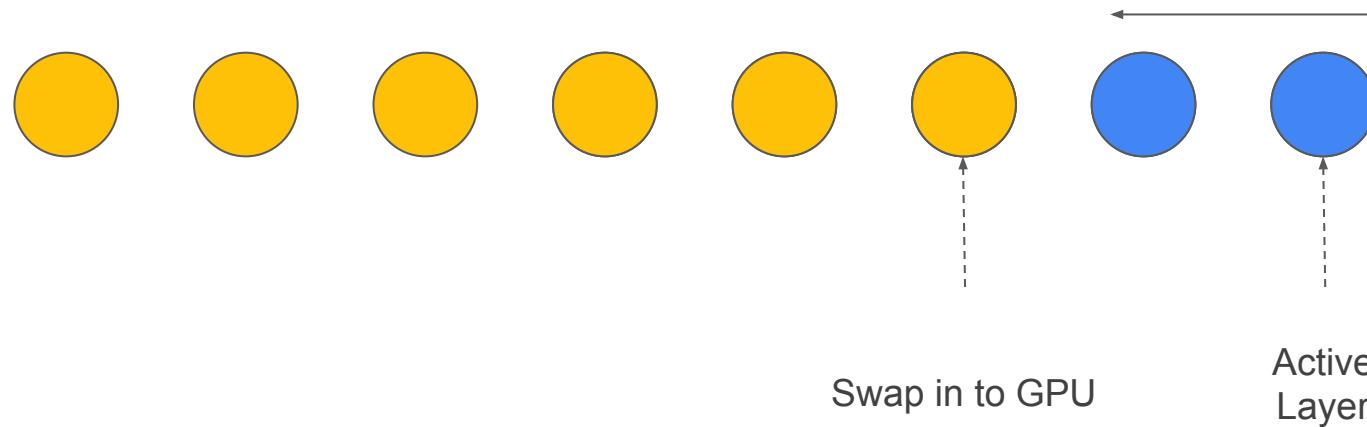
■ CPU ■ GPU



CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time

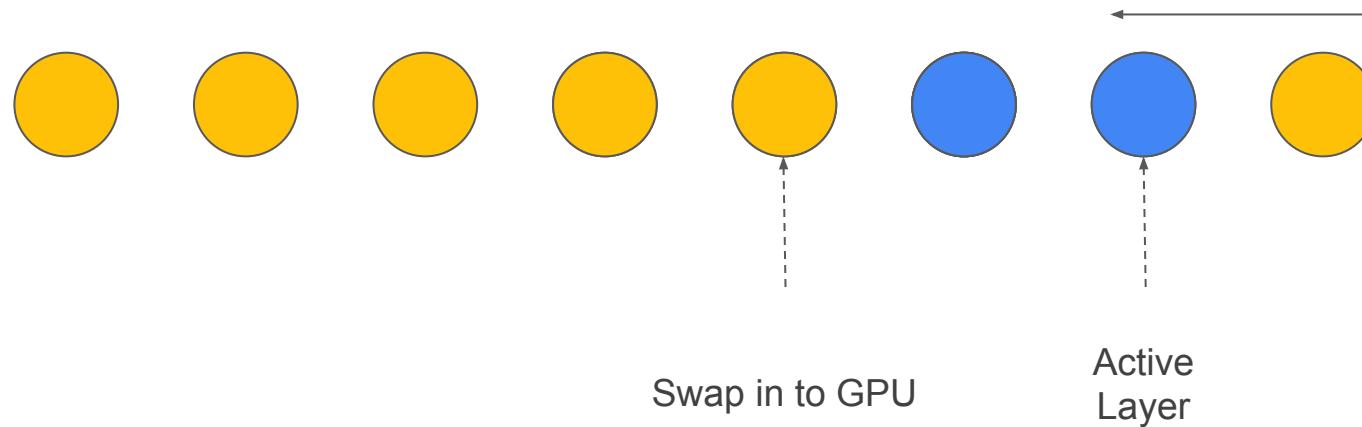
■ CPU ■ GPU



CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time

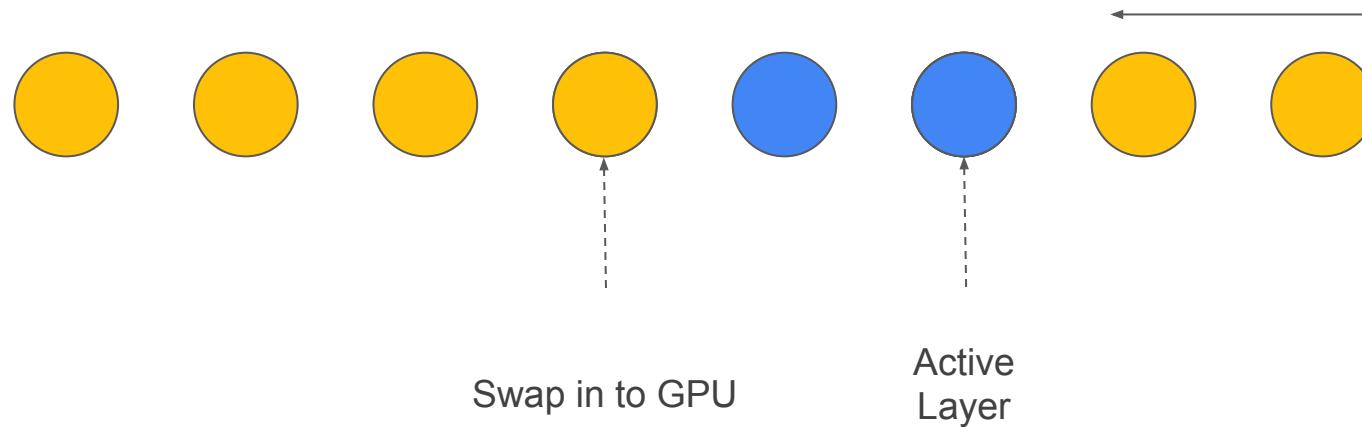
■ CPU ■ GPU



CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time

■ CPU ■ GPU



CPU<->GPU Memory Swapping / Paging

- Swap-out activations / weights to CPU once a layer is completed
- Swap-in activations /weights to GPU before a layer is started
- Exact timing of swap-in/swap-out depends on layers size and layer forward/backward time
- Benefits:
 - 60-80% memory reduction
 - Network usually not slower. If it is slower, swap-int layers earlier (less memory reduction)
 - Faster training due to larger batch size for very large models

Mixed Precision Training (FP16+FP32) / BF16 training

Mixed Precision Training:

- Keep 32-bit master weights
- Do forward pass with 16-bit
- Scale 16-bit loss to prevent under/overflow
- Compute gradients
- Update 32-bit weights; copy 32-bit weights to 16-bit buffers

BrainFloat-16 Training:

- Range: FP16 +-65504; BF16 & FP32 +-3e^38
- Cast everything to BF16
- Train normally (no under/overflow due to larger range)

Benefits:

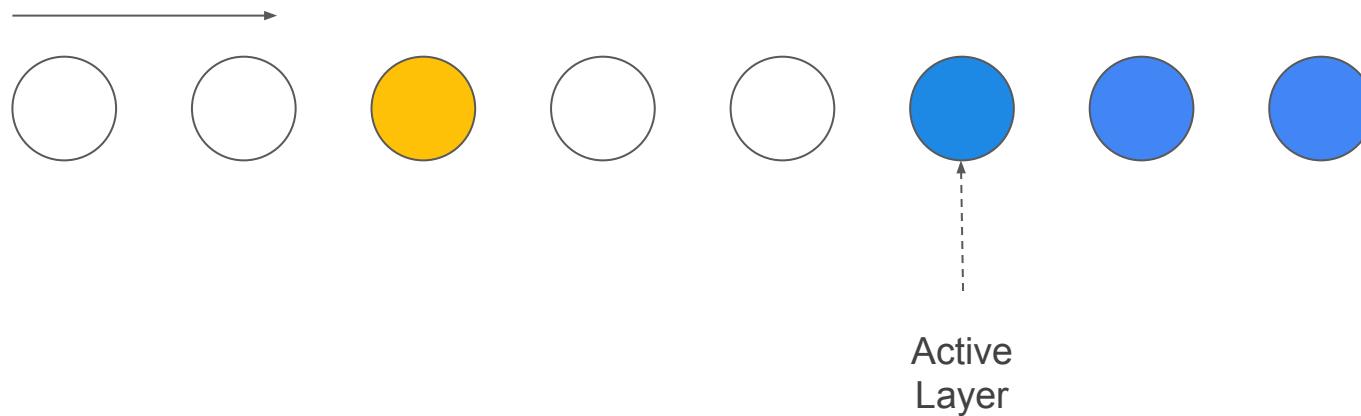
- Faster training, depending on network about 2x speedup
- Usually save some memory, especially if your activations are large

Micikevicius et al., 2018: [arXiv](#)

Gradient Checkpointing: Forward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

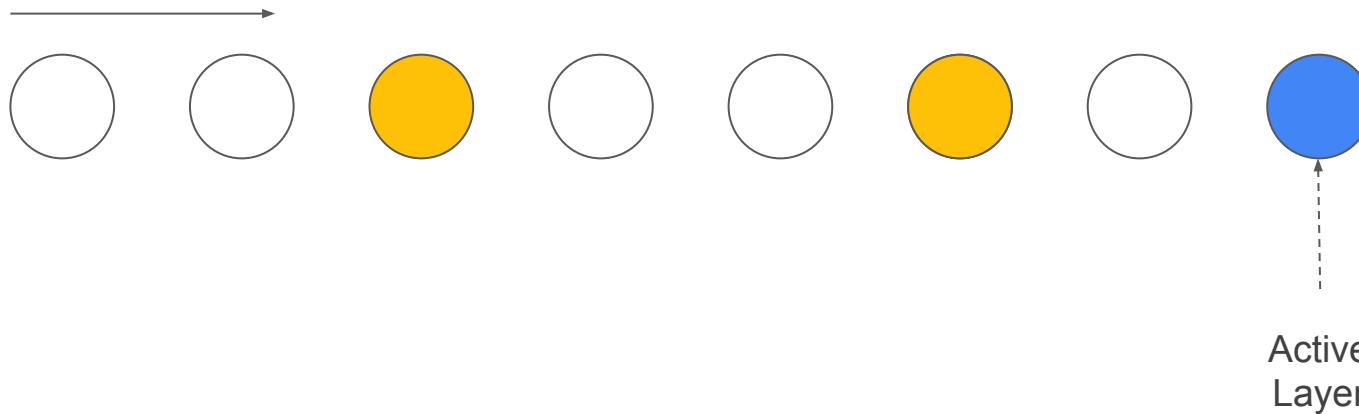
□ Dropped ■ Checkpointed ■ Not Computed Yet



Gradient Checkpointing: Forward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

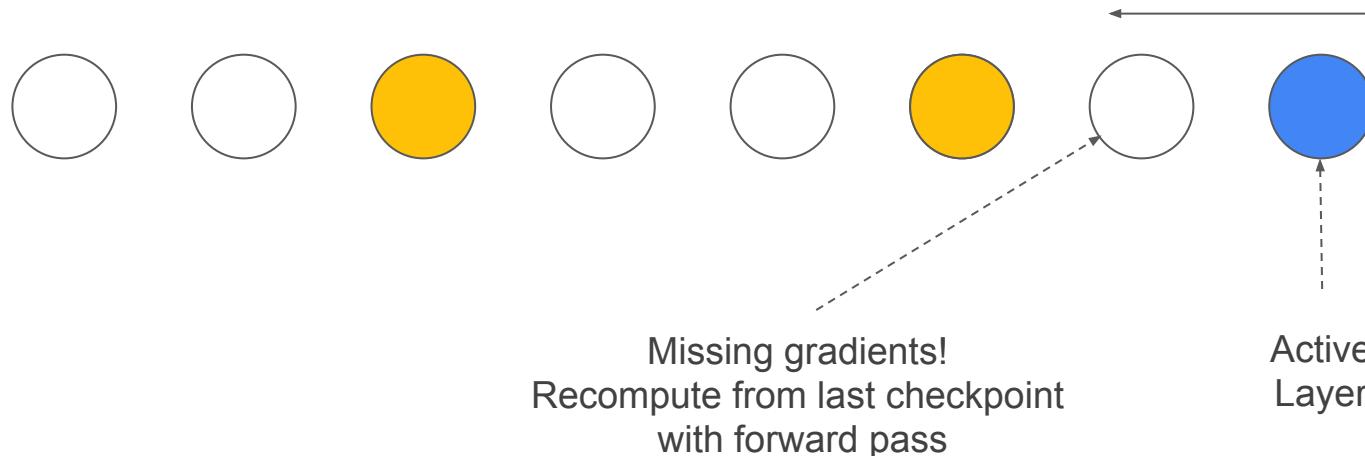
□ Dropped ■ Checkpointed ■ Not Computed Yet



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

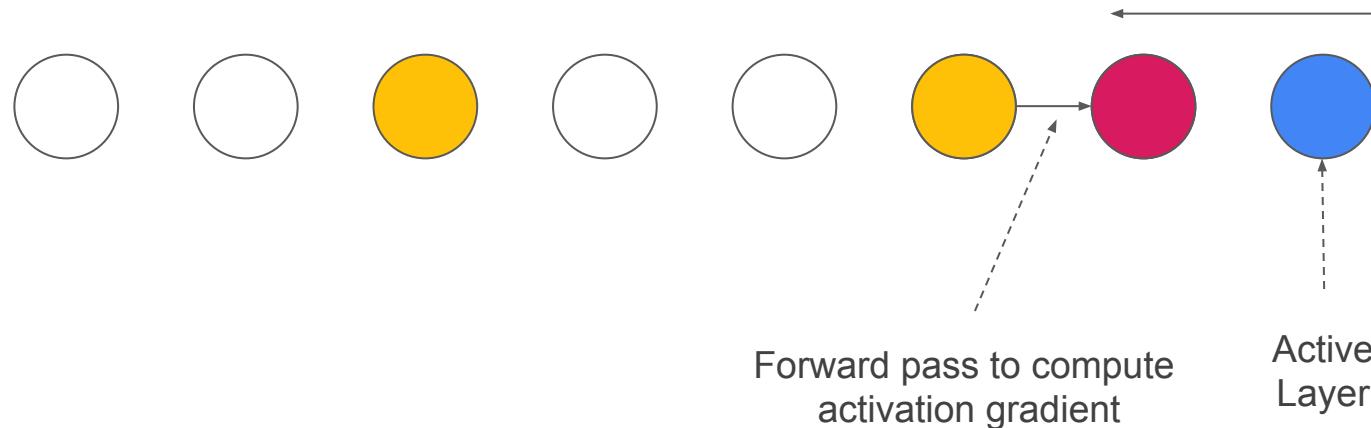
□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

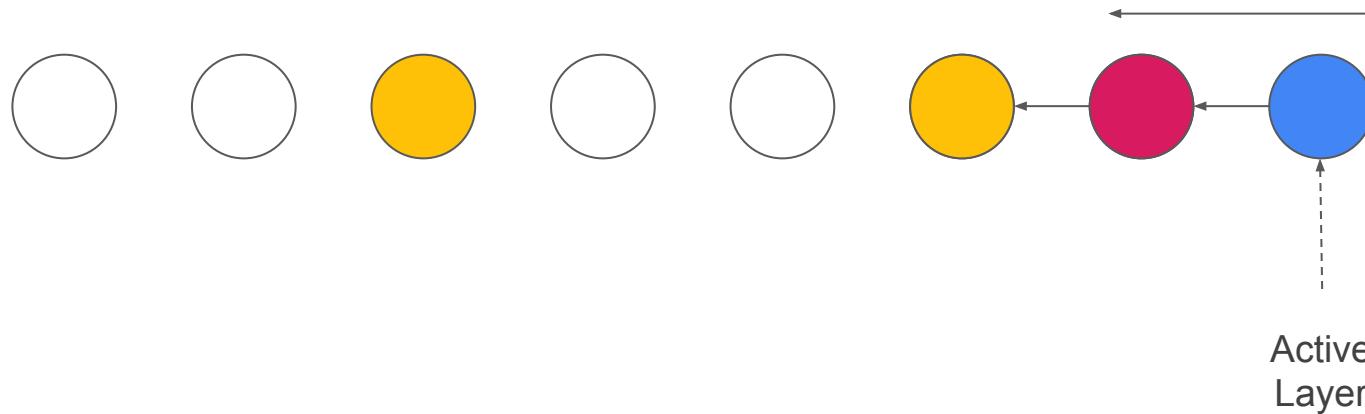
□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

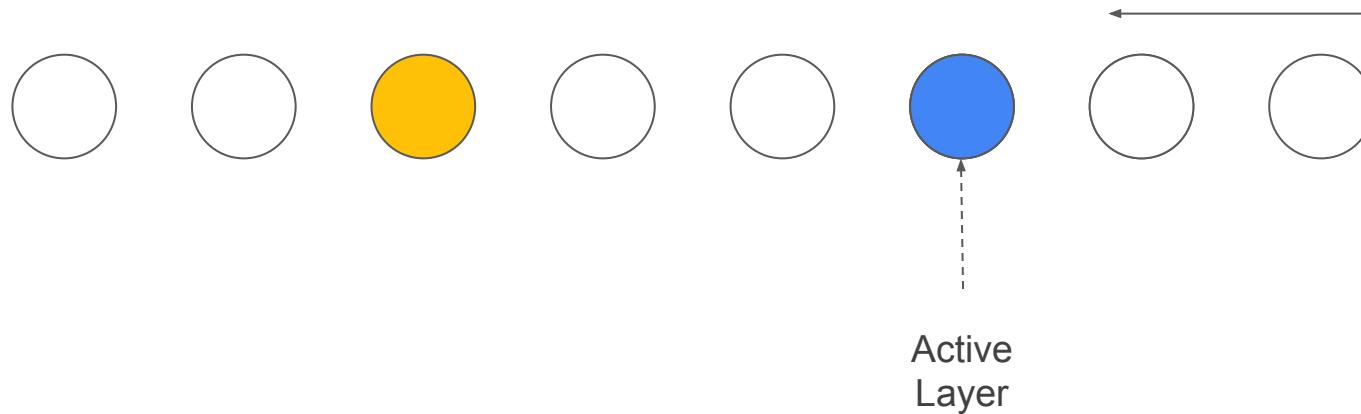
□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

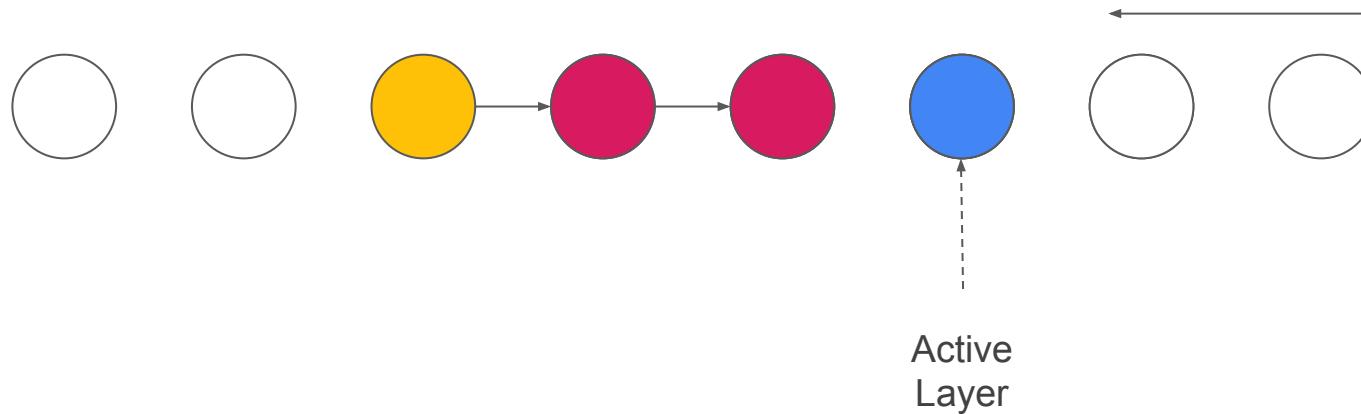
□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

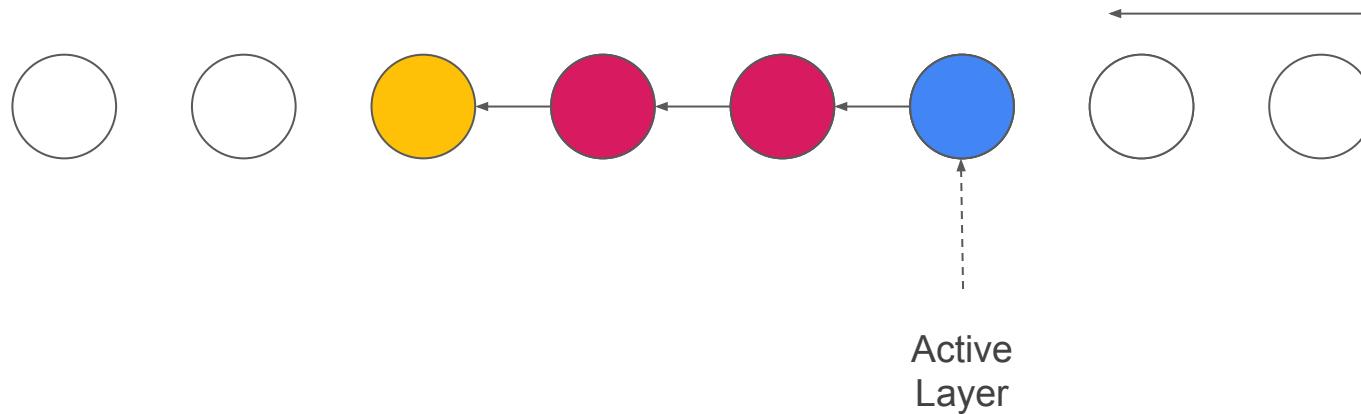
□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

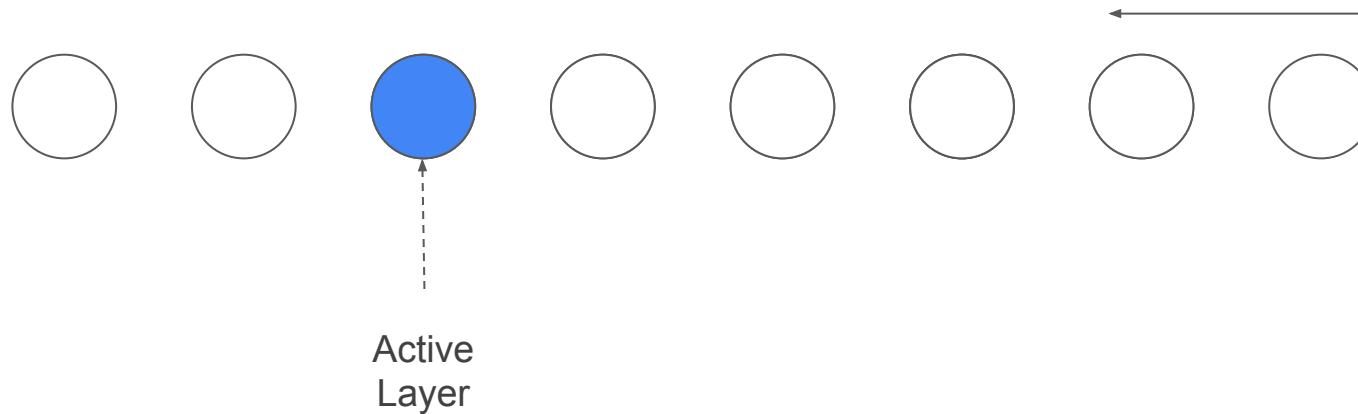
□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing: Backward

- Do not store activation gradients in the forward pass
- Recompute activation gradients in the backward pass by restarting a forward pass from a checkpoint node

□ Dropped ■ Checkpointed □ Not Computed Yet ■ Has Gradient



Gradient Checkpointing

Benefits:

- Trade computation to reduce memory footprint
- Best used for functions that are cheap to compute but have a large activation gradient (ReLU, layer norm, softmax)
- Very beneficial for nonlinear activation functions
- Easy to use in PyTorch ([torch.utils.checkpoint](#)) and TensorFlow 2.0 ([recompute_grad](#) (nightly))

Reversible Residual Connections

- Divide network output and residual connection into two halves. Compound them into a reversible structure:

Forward

$$y_1 = x_1 + \mathcal{F}(x_2)$$

$$y_2 = x_2 + \mathcal{G}(y_1)$$

Backward

$$x_2 = y_2 - \mathcal{G}(y_1)$$

$$x_1 = y_1 - \mathcal{F}(x_2)$$

Benefits:

- Saves some memory for free (if your framework supports it, e.g. JAX)
- Usually, gradient checkpointing should be preferred:
 - Can save more memory due to being more general.
 - Easy to implement. Supported by major frameworks.

Gradient Accumulation

- Split larger mini-batches into “micro-batches”
- Do standard forward/backward passes with micro-batches, but do not update the weights right away (and do not reset the gradient on the weights)
- Accumulate the gradient on the weights for all micro-batches
- Update the weights once enough micro-batches have been computed

Benefits / Tradeoffs:

- As long as your model runs with batch size 1 you can simulate any batch size
- Easy to implement and can reduce memory footprint significantly
- Slow if micro-batch size is very small
- Can improve data parallel performance significantly (speedups) especially for very large models

Parallelism

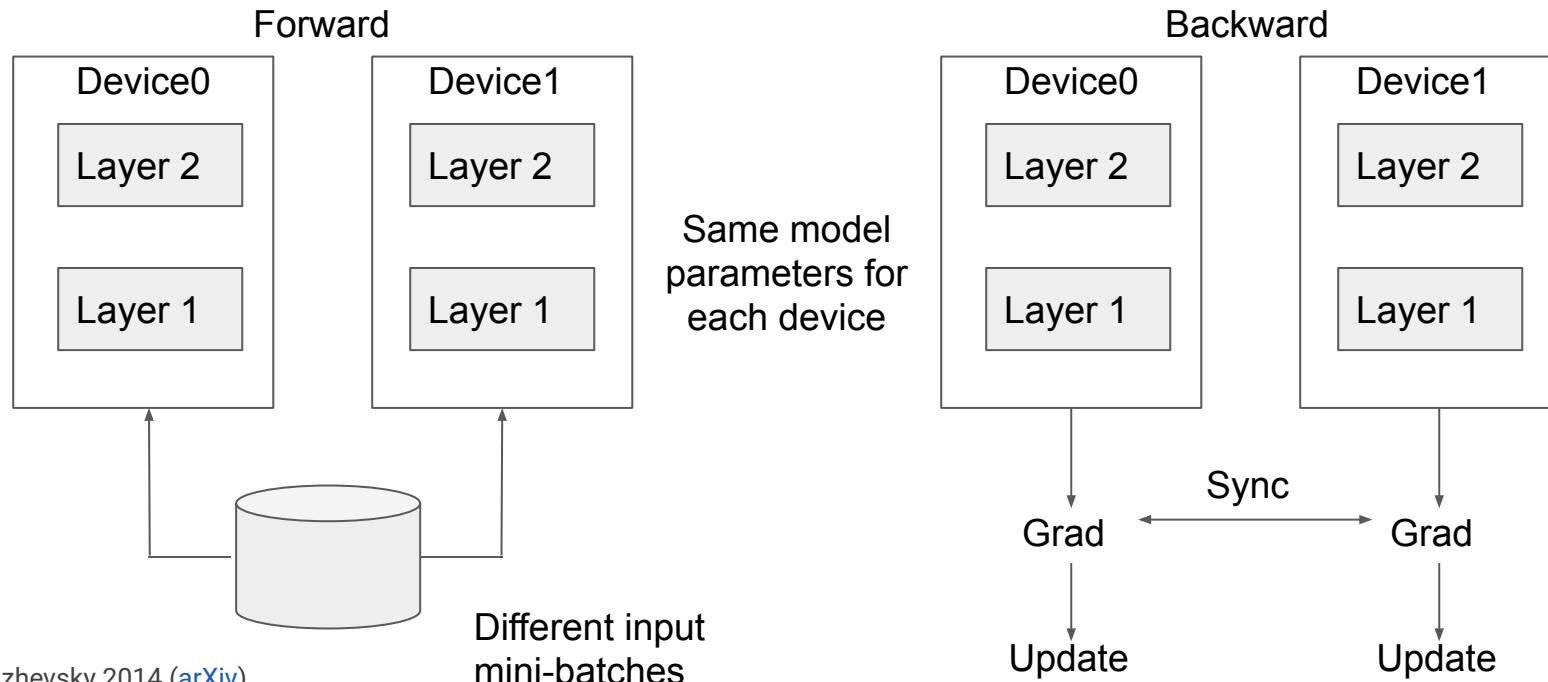


Parallelism Overview

- Data parallelism
- Model parallelism
- Pipeline parallelism
- ZeRo parallelism optimizations
- 3D parallelism

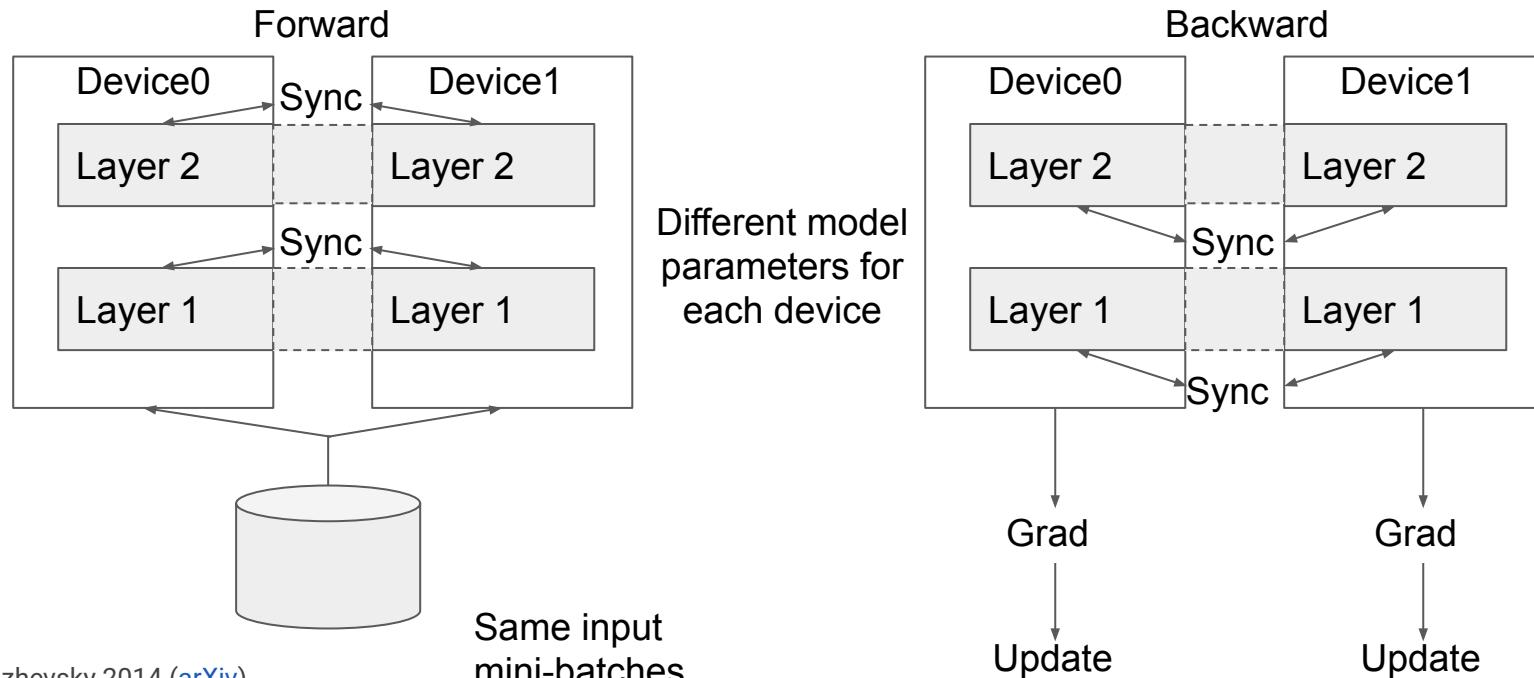
Data Parallelism

Idea: Keep the same model parameters across multiple accelerators. Feed them different mini-batches and average the gradient across accelerators.



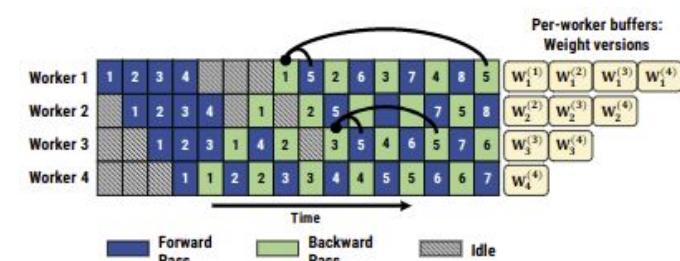
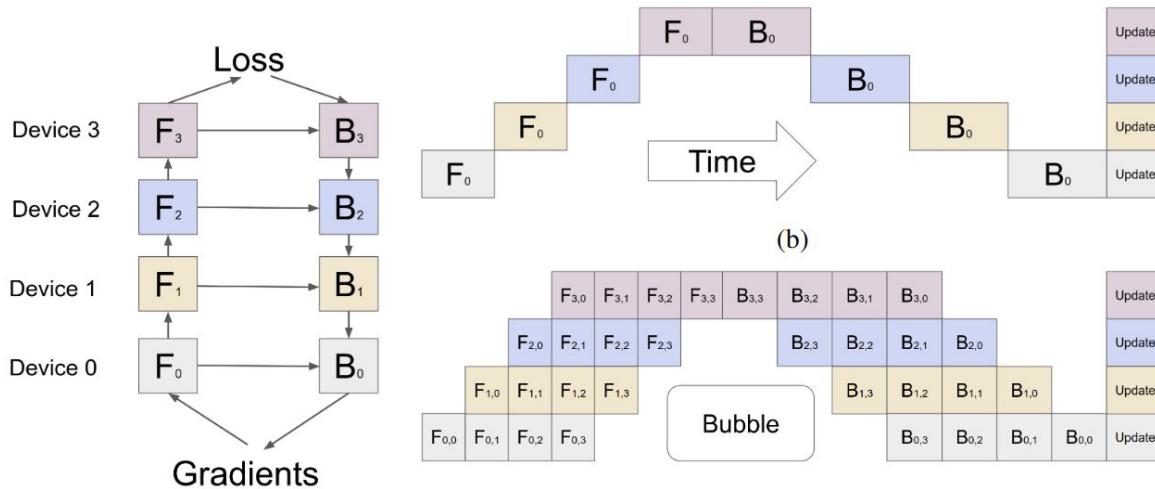
Model Parallelism

Idea: Keep the same mini-batch across multiple accelerators; split the layers parameters across all devices and synchronize layer outputs after each layer.



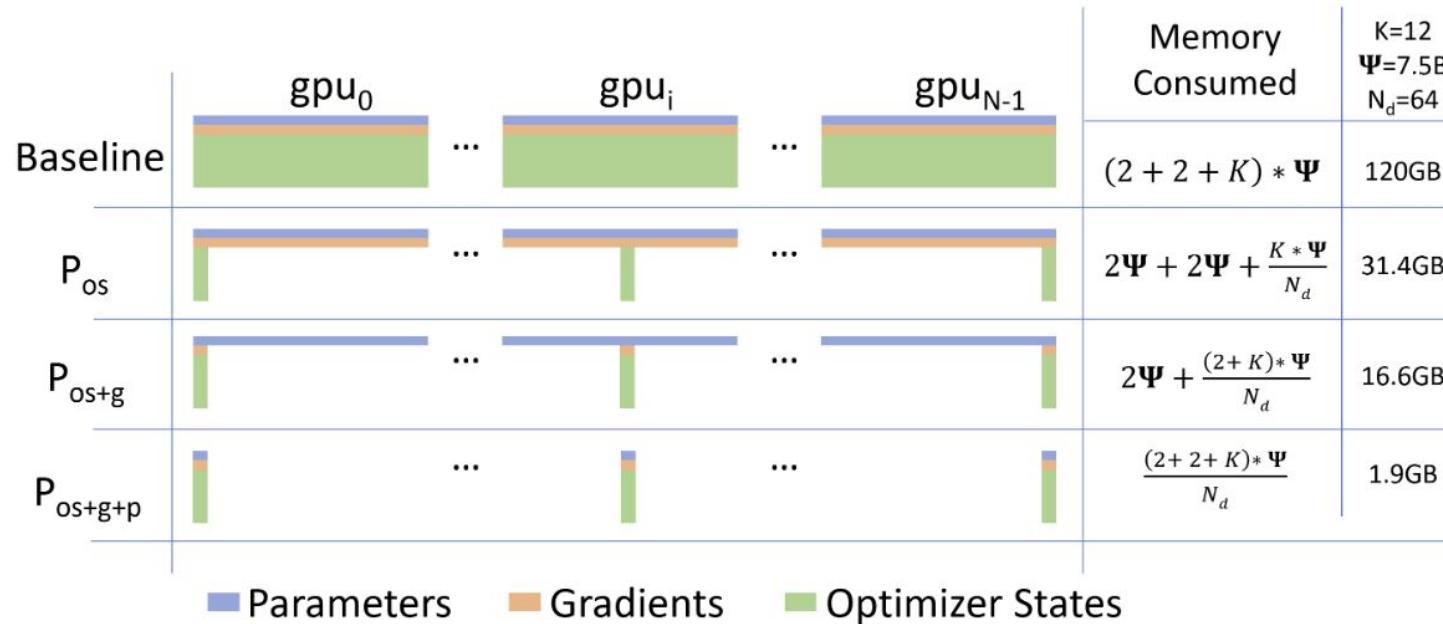
Pipeline Parallelism

Idea: Split network by depth into k pieces onto k accelerators. Each accelerator holds $1/k$ th of layers. Use micro-batches to overlap computation and communication.

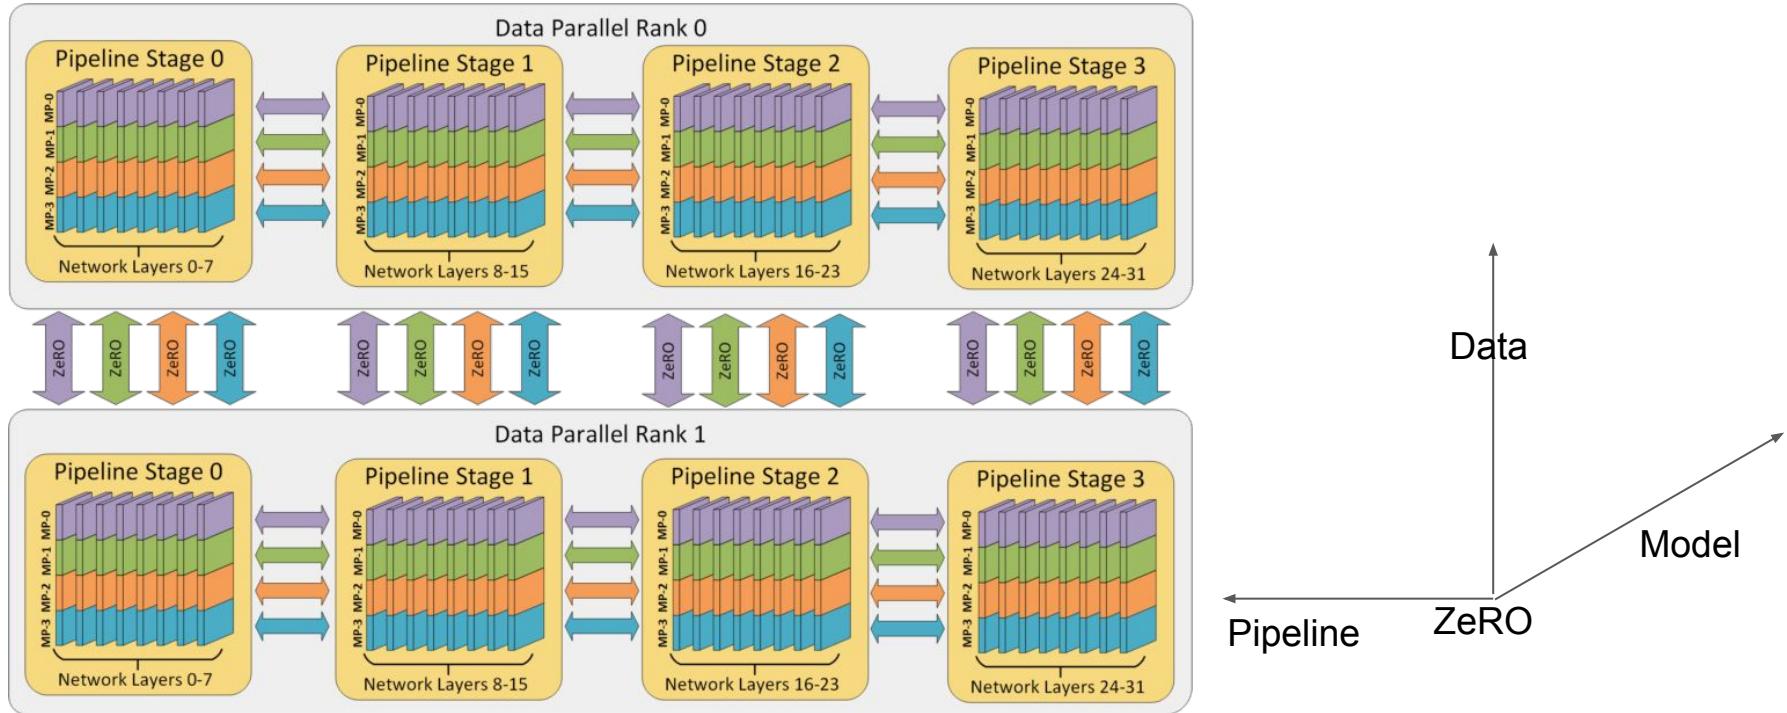


ZeRO Parallelism Optimizations

Idea: Gradients, parameters, and optimizer state only needed for active layer. We distribute the state across all GPUs and gather them together when we need them (when they become “active”).



3D Parallelism



Why 3D Parallelism?

- Model parallelism bad if batch size is too large. Communication cannot be overlapped with computation.
- Data parallelism bad if the batch size is too small.
- Pipeline parallelism decreases mini-batch size through micro-batches.
- Pipeline parallelism increase min-batch size through aggregation of micro-batches.
- Pipeline parallelism allows for simple overlap of communication and computation..

Efficiency Optimizations

Larger Batch Size

- GPUs are more efficient if fully utilized. That usually only happens if batch size is large
- GPUs run better if the mini-batch dimension is 32 or larger
- Often you can achieve *faster* training by using a memory efficiency technique which *slows* down training but enables training with larger batch size
- Larger batch sizes enables larger learning rates. While computation is slower, training might be faster.

Fused Kernels

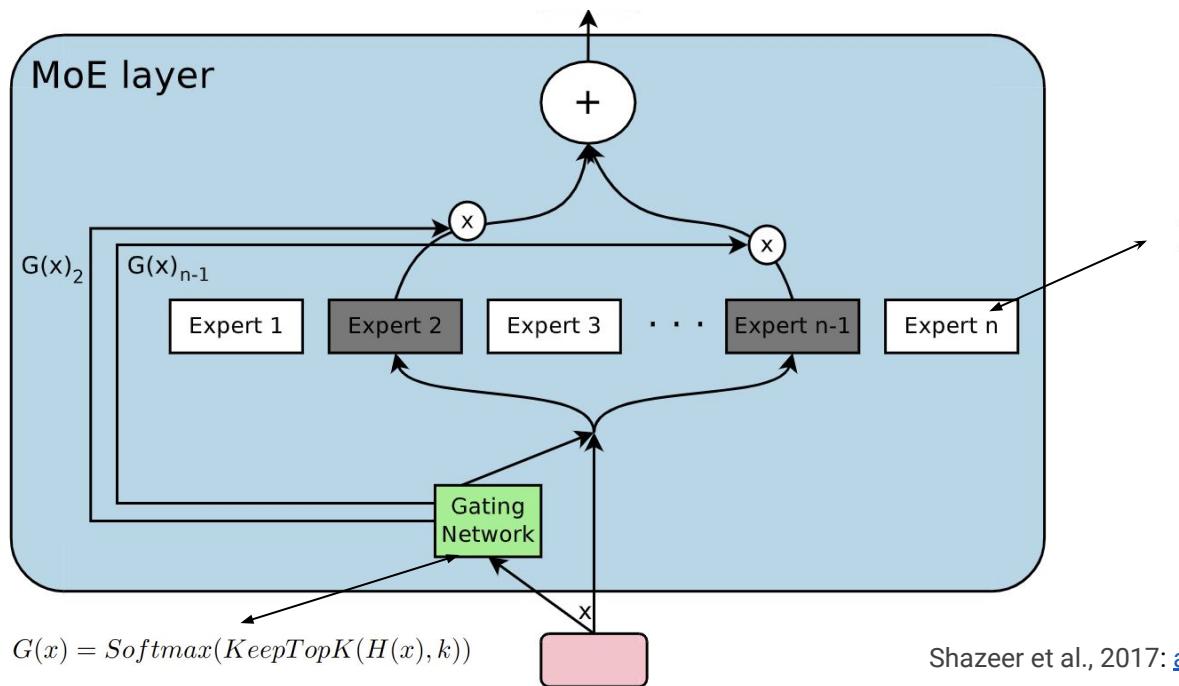
```
if group['weight_decay'] != 0:  
    grad.add_(group['weight_decay'], p.data)  
  
# Decay the first and second moment running average coefficient  
exp_avg.mul_(beta1).add_(1 - beta1, grad)  
exp_avg_sq.mul_(beta2).addcmul_(1 - beta2, grad, grad)  
if amsgrad:  
    # Maintains the maximum of all 2nd moment running avg. till now  
    torch.max(max_exp_avg_sq, exp_avg_sq, out=max_exp_avg_sq)  
    # Use the max. for normalizing running avg. of gradient  
    denom = (max_exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(group['eps'])  
else:  
    denom = (exp_avg_sq.sqrt() / math.sqrt(bias_correction2)).add_(group['eps'])  
  
step_size = group['lr'] / bias_correction1  
  
p.data.addcdiv_(-step_size, exp_avg, denom)
```

- Adam with 10^9 parameters:
 - 14 read/writes
 - 32-bit 10^9 parameters = 4 GB
 - Normal Adam: GPU with 600 GB/s $\rightarrow 14 \times 4 / 600 = 100$ ms
 - Fused Adam: 6ms

Mixture of Experts



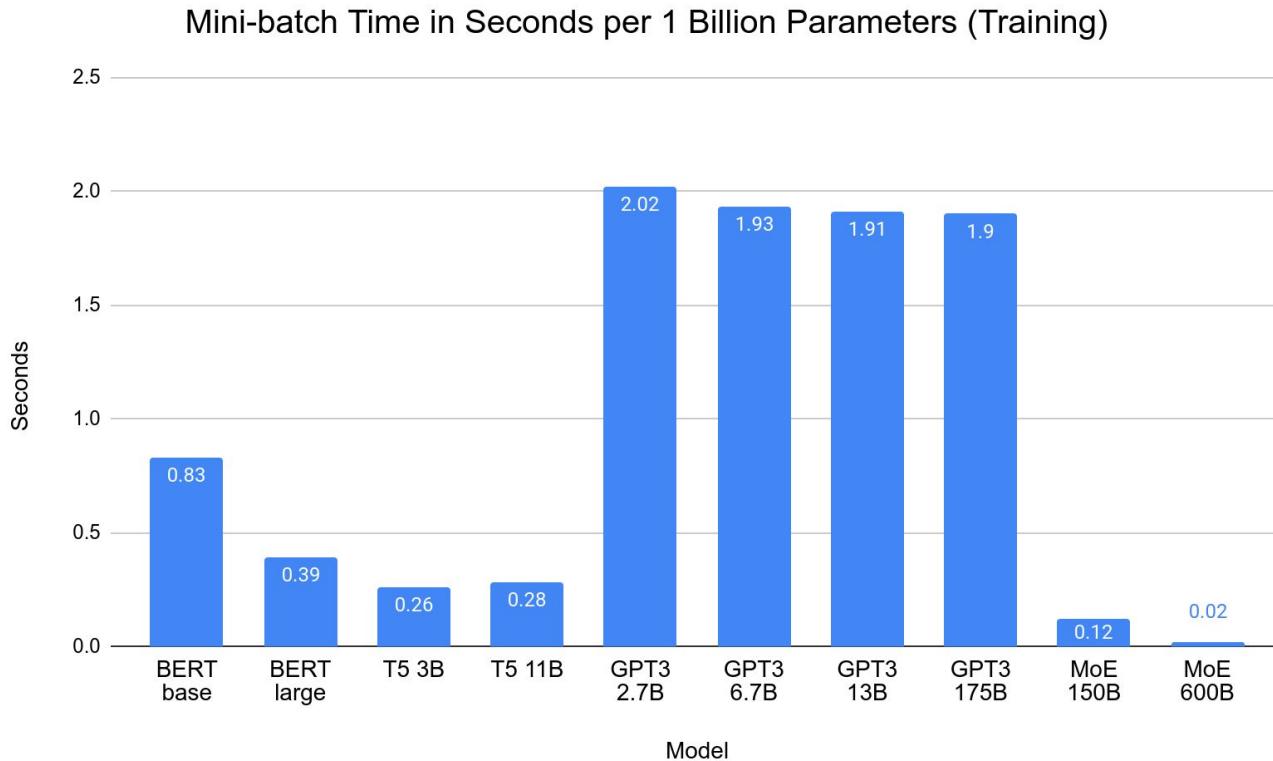
Mixture of Experts: Overview



Shazeer et al., 2017: [arXiv](#)

Lepikhin et al., 2020: [arXiv](#)

Transformers Mini-batch Time



Mixture of Experts: Balancing and Specialization v1

Version 1 ([Shazeer et al., 2017](#)):

Initialize W_g and W_{noise} with zeros, so outputs are driven by standard normal noise. This guarantees balancing across experts at the start of training.

The noise also helps to decrease early advantage of previously picked experts.

$$G(x) = \text{Softmax}(\text{KeepTopK}(H(x), k))$$

$$H(x)_i = (x \cdot W_g)_i + \text{StandardNormal()} \cdot \text{Softplus}((x \cdot W_{noise})_i)$$

Mixture of Experts: Balancing and Specialization v1

Version 1 ([Shazeer et al., 2017](#)):

An additional balancing loss assigns high loss to experts which have very high probability. This prevents failure cases where an expert is always picked with 100% probability.

$$Importance(X) = \sum_{x \in X} G(x)$$

$$L_{importance}(X) = w_{importance} \cdot CV(Importance(X))^2$$

Coefficient of variation: $CV(X) = \text{std}(X)/\text{mean}(X)$

Mixture of Experts: Balancing and Specialization v1

Version 1 ([Shazeer et al., 2017](#))

Importance loss can be satisfied by picking a subset of experts. To prevent this degeneration we want to pick *all* experts with roughly the same probability over time.

If we view the softplus term as something analogous to a standard deviation and the mean softmax as the expected value, we can express an approximate probability for this with a CDF of the normal distribution.

$$H(x)_i = (x \cdot W_g)_i + \text{StandardNormal}() \cdot \text{Softplus}((x \cdot W_{noise})_i)$$

$$P(x, i) = \Phi\left(\frac{(x \cdot W_g)_i - k\text{th_excluding}(H(x), k, i)}{\text{Softplus}((x \cdot W_{noise})_i)}\right)$$

$$\text{Load}(X)_i = \sum_{x \in X} P(x, i)$$

$$L_{load}(X) = w_{load} \cdot \text{CV}(\text{Load}(X))^2$$

Mixture of Experts: Balancing and Specialization v2

Version 2 ([Lepikhin et al., 2020](#)):

No noise. Initialize layers normally. Keep track of c_e , how many times each expert was used for the sequence S. With the mean gate probability of m_e we can now define a balancing auxiliary loss:

$$\ell_{aux} = m_e(c_e/S)$$

$$\mathcal{L} = \ell_{nll} + k * \ell_{aux}$$

Where k is a constant loss weight (a good value is 0.1; usually between 0.01 and 1.0)

Mixture of Experts: Balancing and Specialization v2

Version 2 ([Lepikhin et al., 2020](#)):

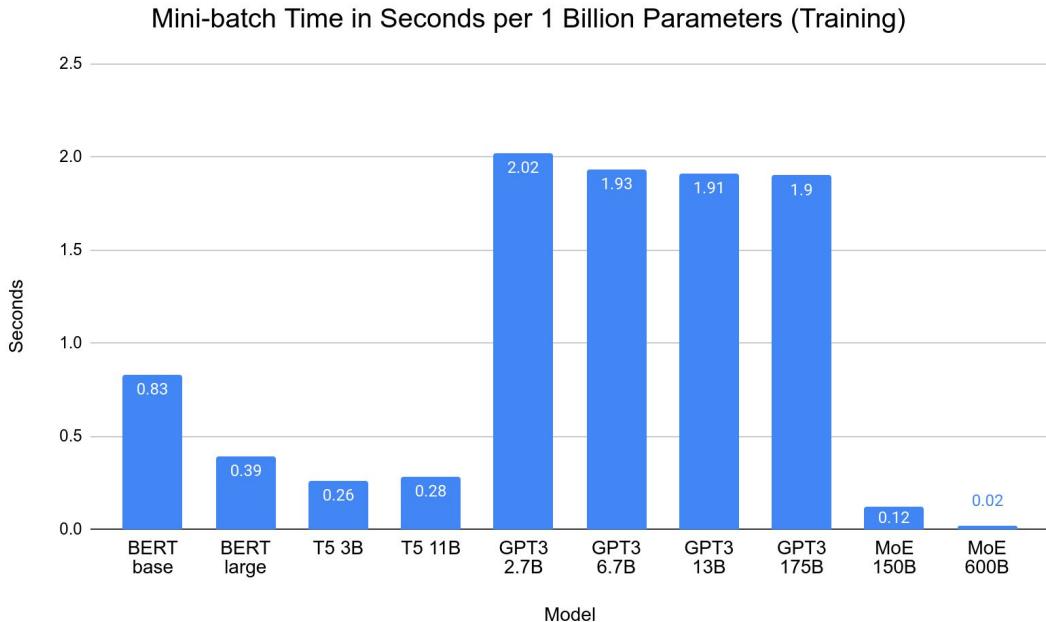
- Random dispatch: Use 2nd expert proportionally to the softmax gate probability.
- Have a frequency cutoff – a token budget – for each expert. If this budget is exceeded the expert degenerated to a zero matrix. This effectively reduces the output of the MoE layer to zero and thus only the residual connection output around the MoE layer is fed to the next layer.

Mixture of Experts: Balancing and Specialization

Many cases of expert degeneration:

1. **Overbalancing:** All experts are approximately equally used. However, gate probability approaches $1/\#Experts$. No expert is better than another expert.
2. **Underbalancing:** The same top-k experts are used for every token. This leads to two strong experts, but all other experts do not learn anything and are “wasted capacity”.
3. **Sequence-level degeneration:** Model balances experts by using each expert for a particular sequence index. For example, for indices 0, 1, 2, 3 always experts E3, E1, E2, E0. This leads to sequence experts, but not content experts.

Mixture of Experts: Benefits



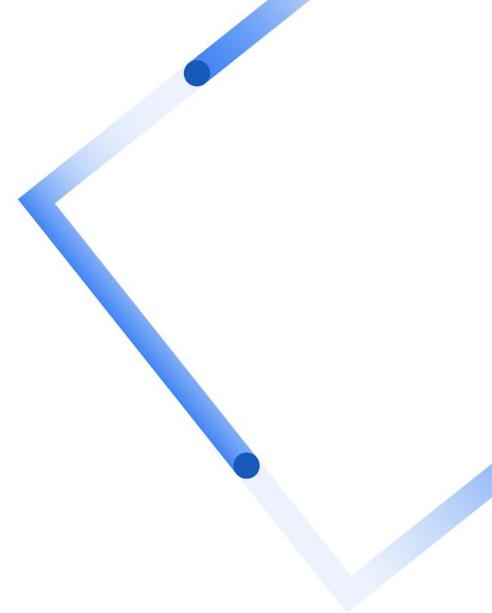
- Works well on diverse data like multilingual machine translation
- Can be difficult to train due to balancing/specialization issues
- Only faster than transformers if you can run it with a large enough batch size to saturate distributed experts
- If you scale the model across a cluster, you will need excellent interconnect performance
(TPU v4 Pod, NVIDIA SuperPod)

Shazeer et al., 2017: [arXiv](#)

Lepikhin et al., 2020: [arXiv](#)

07

Closing Notes



Why we should strive for efficiency

Our field has seen a dramatic increase in scale in the past 2 years.

Striving for efficiency means caring about:

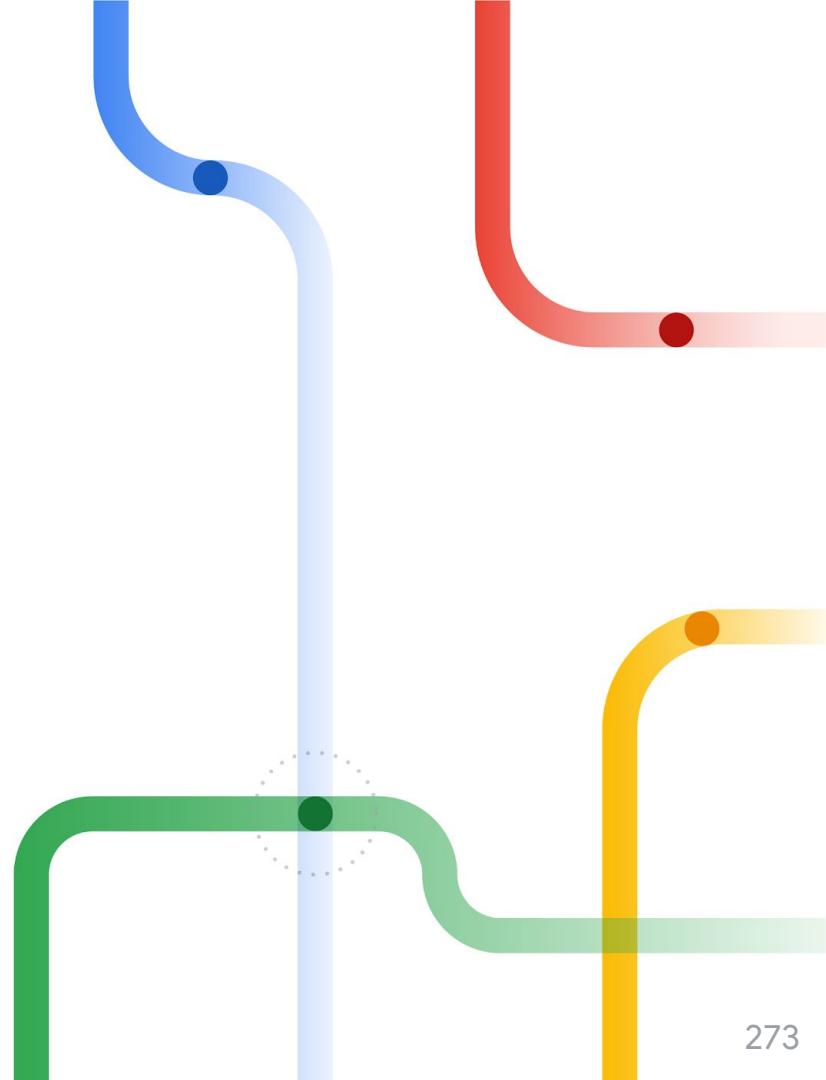
- 1) Costs
- 2) Accessibility
- 3) Production needs
- 4) The sustainability of this growth

Closing Notes

In this tutorial, we covered a wide range of ideas, applications and practical considerations that helps us build more efficient systems, including:

- 1) Core efficiency techniques
- 2) Efficiency improvements to attention mechanisms
- 3) Case studies of efficient models
- 4) Practical considerations for scaling models

We hope you
enjoyed it and
learned something
new!



Thank you!

Slides available at: bit.ly/2SmhKY7



Google Research

